

PROJECT: Diagnosing the latency–throughput tradeoff in vLLM serving under concurrent loads

- **Experiment 1:** Baseline setup and workload harness for mixed prompt-length serving
- **Experiment 1.1:** Fix max-tokens and context-budget failures by adding dynamic `max_tokens` to prevent request rejections
- **Experiment 2:** BAD (unfair) scheduling characterization with concurrency sweeps to expose head-of-line blocking and tail-latency spikes
- **Experiment 3:** Tail latency and throughput variation with concurrency sweep by holding workload and output length constant to isolate scheduling effects
- **Experiment 4:** Sweep of concurrency and `max-num-seqs` under a deliberately bad configuration to map scaling limits, queueing onset, and short request degradation from long-prefill contention
- **Experiment 5:** Find a strong FAIR scheduling configuration by first sweeping `max-num-batched-tokens` at lower concurrency ($c=16$), then validating at $c=32$, and finally running a `max-num-batched-tokens × max-num-seqs` grid at $c=32$ to pick a stable throughput vs tail-latency sweet spot

Installing libraries

```
In [ ]: !pip -q install datasets transformers
!pip -q install tqdm

!pip -q install "jedi>=0.16"
!pip -q install vllm
```

Importing Libraries

```
In [ ]: from google.colab import drive
import os

import json, random, statistics
from pathlib import Path
from datasets import load_dataset
from transformers import AutoTokenizer

import itertools

import numpy as np
import pandas as pd
```

```
from tqdm.notebook import tqdm
import random
import vllm
import torch
from vllm import LLM, SamplingParams
import requests, time

import concurrent.futures as cf
import pandas as pd

import shutil
import math
import subprocess
```

Mounting the Drive

In []: # Mounting Google Drive
drive.mount("/content/drive")

Mounted at /content/drive

In []: # Write locally first for faster write, then copy to Drive at the end
OUT_DIR = "/content/workloads"
os.makedirs(OUT_DIR, exist_ok=True)

Where to persist outputs in Drive
DRIVE_OUT_DIR = "/content/drive/MyDrive/llm_loadtest_workloads"
os.makedirs(DRIVE_OUT_DIR, exist_ok=True)

print("OUT_DIR:", OUT_DIR)
print("DRIVE_OUT_DIR:", DRIVE_OUT_DIR)

OUT_DIR: /content/workloads
DRIVE_OUT_DIR: /content/drive/MyDrive/llm_loadtest_workloads

In []: MODEL = "Qwen/Qwen2.5-3B-Instruct"
tok = AutoTokenizer.from_pretrained(MODEL, use_fast=True)

def tok_len(text: str) -> int:
 return len(tok.encode(text, add_special_tokens=False))

def in_range(n, r):
 return r[0] <= n <= r[1]

random.seed(42)

SHORT_RANGE = (8, 24) # NQ-friendly short prompts
MED_RANGE = (128, 256) # UltraChat-friendly medium prompts
LONG_RANGE = (2048, 4096) # long prompts via concatenation

How many prompts you want in each bucket
WANT_SHORT = 800
WANT_MEDIUM = 400
WANT_LONG = 600

Requests per workload file

```
TOTAL_REQS_PER_FILE = 400

print("Tokenizer loaded for:", MODEL)

tokenizer_config.json: 0.00B [00:00, ?B/s]
vocab.json: 0.00B [00:00, ?B/s]
merges.txt: 0.00B [00:00, ?B/s]
tokenizer.json: 0.00B [00:00, ?B/s]
Tokenizer loaded for: Qwen/Qwen2.5-3B-Instruct
```

Function to enable Dataset Streaming

```
In [ ]: def try_load_streaming(dataset_name, config_split_candidates):
    last_err = None

    for cfg, split in config_split_candidates:
        try:
            if cfg is None:
                ds = load_dataset(dataset_name, split=split, streaming=True)
            else:
                ds = load_dataset(dataset_name, cfg, split=split, streaming=True)
        print(f"Loaded {dataset_name} (config={cfg}, split={split})")

    return ds

except Exception as e:
    last_err = e
    print(f"Failed {dataset_name} (config={cfg}, split={split})")

raise RuntimeError(f"Could not load {dataset_name}. Last error: {last_err}")
```

```
In [ ]: # streaming both the datasets
ultra_stream = try_load_streaming(
    "HuggingFaceH4/ultrachat_200k",
    config_split_candidates=[
        (None, "train_sft"),
        ("default", "train_sft"),
        ("default", "train"),
    ],
)

nq_stream = try_load_streaming(
    "natural_questions",
    config_split_candidates=[
        (None, "train"),
        ("default", "train"),
    ],
)
```

```
README.md: 0.00B [00:00, ?B/s]
Loaded HuggingFaceH4/ultrachat_200k (config=None, split=train_sft)
README.md: 0.00B [00:00, ?B/s]
Resolving data files:  0%|          | 0/287 [00:00<?, ?it/s]
Resolving data files:  0%|          | 0/287 [00:00<?, ?it/s]
Loaded natural_questions (config=None, split=train)
```

Check if the Dataset loaded properly

```
In [ ]: ex = next(iter(nq_stream))
print("Keys:", ex.keys())
print("Sample:", {k: ex[k] for k in list(ex.keys())[:5]})

# clearing the output because it takes a lot of space in pdf version of n
```



```
In [ ]: ex = next(iter(ultra_stream))
print("Keys:", ex.keys())
print("Sample:", {k: ex[k] for k in list(ex.keys())[:5]})

# clearing the output because it takes a lot of space in pdf version of n
```

Functions to extract prompt from Dataset

```
In [ ]: def ultrachat_user_text(ex):
    prompt = ex.get("prompt")
    if isinstance(prompt, str) and prompt.strip():
        return prompt.strip()

    # This is a fallback: scan the "messages" list for the first user turn
    msgs = ex.get("messages")
    if isinstance(msgs, list):
        for m in msgs:
            if not isinstance(m, dict):
                continue

            role = (m.get("role") or "").lower()
            content = m.get("content") or ""

            if role == "user" and isinstance(content, str) and content.strip():
                return content.strip()

    return None


def nq_question_text(ex):
    q = ex.get("question")
    if isinstance(q, dict):
        t = q.get("text")
        return t.strip() if isinstance(t, str) and t.strip() else None
    if isinstance(q, str):
        return q.strip() if q.strip() else None
    return None
```

Create Small and Medium Buckets

```
In [ ]: from tqdm import tqdm

def stream_bucket(
    stream,
    text_fn,
    want_n,
```

```

    token_range,
    max_read,
    desc,
    print_every=500,
    also_keep_pool=True,
):
    """
    Stream from a HF streaming dataset, extract text, collect:
    - bucket: texts in token_range until want_n
    - pool: all extracted texts (optional)
    """

    bucket = []
    pool = [] if also_keep_pool else None

    for i, ex in tqdm(enumerate(stream), total=max_read, desc=desc):
        if i >= max_read:
            break

        s = text_fn(ex)
        if not s:
            continue

        if also_keep_pool:
            pool.append(s)

        n = tok_len(s)
        if len(bucket) < want_n and in_range(n, token_range):
            bucket.append(s)

        if (i % print_every) == 0:
            tqdm.write(f"read={i}, {desc}_bucket={len(bucket)}")

        if len(bucket) >= want_n:
            break

    return bucket, pool

```

In []: MAX_NQ_READ = 50000

```

short_bucket, nq_pool = stream_bucket(
    stream=nq_stream,
    text_fn=nq_question_text,
    want_n=WANT_SHORT,
    token_range=SHORT_RANGE,
    max_read=MAX_NQ_READ,
    desc="short",
    print_every=200, # keep your original cadence
    also_keep_pool=True,
)

print("short_bucket:", len(short_bucket))

```

```
short: 0%|          | 2/50000 [00:18<104:09:36, 7.50s/it]
read=0, short_bucket=1
```

```
short: 0%|          | 208/50000 [00:30<26:42, 31.07it/s]
read=200, short_bucket=201
```

```
short: 1%|          | 402/50000 [00:42<2:21:13, 5.85it/s]
read=400, short_bucket=401
```

```
short: 1%|          | 604/50000 [00:58<35:59, 22.88it/s]
```

```
read=600, short_bucket=601
short: 2%||          | 799/50000 [01:08<1:10:09, 11.69it/s]
short_bucket: 800
```

```
In [ ]: MAX_ULTRA_READ = 200000

medium_bucket, ultra_pool = stream_bucket(
    stream=ultra_stream,
    text_fn=ultrachat_user_text,
    want_n=WANT_MEDIUM,
    token_range=MED_RANGE,
    max_read=MAX_ULTRA_READ,
    desc="medium",
    print_every=500, # keep your original cadence
    also_keep_pool=True,
)

print("medium_bucket:", len(medium_bucket))
```

```
medium: 0%|          | 14/200000 [00:03<9:37:24, 5.77it/s]
read=0, medium_bucket=0
medium: 0%|          | 528/200000 [00:04<10:42, 310.50it/s]
read=500, medium_bucket=51
medium: 1%|          | 1040/200000 [00:06<09:46, 339.40it/s]
read=1,000, medium_bucket=106
medium: 1%|          | 1537/200000 [00:08<12:04, 274.11it/s]
read=1,500, medium_bucket=158
medium: 1%|          | 2042/200000 [00:09<09:52, 334.10it/s]
read=2,000, medium_bucket=198
medium: 1%||         | 2543/200000 [00:11<10:07, 325.01it/s]
read=2,500, medium_bucket=252
medium: 2%||         | 3004/200000 [00:13<23:16, 141.10it/s]
read=3,000, medium_bucket=305
medium: 2%||         | 3614/200000 [00:19<07:18, 448.23it/s]
read=3,500, medium_bucket=355
medium: 2%||         | 3990/200000 [00:20<17:07, 190.85it/s]
medium_bucket: 400
```

Create Large Bucket

```
In [ ]: def build_long_prompt(target_tokens, rng, snippets_pool, max_unique_before_parts = [],
                                used = set()):
    while True:
        idx = rng.randrange(len(snippets_pool))
        if idx in used:
            continue
        used.add(idx)

        parts.append(snippets_pool[idx].replace("\n", " ").strip())
        text = "\n\n".join(parts)

        if tok_len(text) >= target_tokens:
            break
```

```

        if len(used) > min(max_unique_before_reset, len(snippets_pool)) -
            used.clear()

        ids = tok.encode(text, add_special_tokens=False)[:target_tokens]
        return tok.decode(ids)

    def build_long_bucket(
        want_n,
        token_range,
        snippets_pool,
        seed=123,
        desc="long",
        print_every=10,
    ):
        assert len(snippets_pool) > 0, "snippets_pool must not be empty"

        rng = random.Random(seed)
        bucket = []

        for j in tqdm(range(want_n), desc=f"Building {desc}"):
            target = rng.randint(token_range[0], token_range[1])
            lp = build_long_prompt(target, rng, snippets_pool)
            bucket.append(lp)

            if ((j + 1) % print_every) == 0:
                tqdm.write(f"built={j+1}/{want_n}, last_len={tok_len(lp)} tokens")

    return bucket

```

In []:

```

snippets_pool = short_bucket + medium_bucket
long_bucket = build_long_bucket(
    want_n=WANT_LONG,
    token_range=LONG_RANGE,
    snippets_pool=snippets_pool,
    seed=123,
    desc="long",
    print_every=10,
)

print("long_bucket:", len(long_bucket))

```

```

Building long:  2%|          | 10/600 [00:02<03:26,  2.85it/s]
built=10/600, last_len=3366 tokens
Building long:  3%|          | 20/600 [00:06<03:07,  3.09it/s]
built=20/600, last_len=2955 tokens
Building long:  5%|          | 30/600 [00:11<05:48,  1.64it/s]
built=30/600, last_len=3380 tokens
Building long:  7%|          | 40/600 [00:17<07:19,  1.27it/s]
built=40/600, last_len=3161 tokens
Building long:  8%|          | 50/600 [00:25<07:18,  1.25it/s]
built=50/600, last_len=3393 tokens
Building long:  10%|         | 60/600 [00:29<03:09,  2.86it/s]
built=60/600, last_len=2295 tokens
Building long:  12%|         | 70/600 [00:32<02:16,  3.88it/s]
built=70/600, last_len=2631 tokens
Building long:  13%|         | 80/600 [00:37<05:57,  1.45it/s]
built=80/600, last_len=3856 tokens

```

Building long: 15% | [■] 90/600 [00:41<02:56, 2.89it/s]
built=90/600, last_len=2146 tokens

Building long: 17% | [■] 100/600 [00:46<03:35, 2.32it/s]
built=100/600, last_len=2627 tokens

Building long: 18% | [■] 110/600 [00:52<06:05, 1.34it/s]
built=110/600, last_len=4031 tokens

Building long: 20% | [■] 120/600 [00:57<03:45, 2.13it/s]
built=120/600, last_len=3992 tokens

Building long: 22% | [■] 130/600 [01:01<02:43, 2.87it/s]
built=130/600, last_len=2305 tokens

Building long: 23% | [■] 140/600 [01:07<04:23, 1.75it/s]
built=140/600, last_len=3250 tokens

Building long: 25% | [■] 150/600 [01:11<03:04, 2.45it/s]
built=150/600, last_len=3441 tokens

Building long: 27% | [■] 160/600 [01:14<02:37, 2.80it/s]
built=160/600, last_len=3468 tokens

Building long: 28% | [■] 170/600 [01:17<02:08, 3.36it/s]
built=170/600, last_len=2441 tokens

Building long: 30% | [■] 180/600 [01:22<03:39, 1.91it/s]
built=180/600, last_len=3415 tokens

Building long: 32% | [■] 190/600 [01:29<03:24, 2.01it/s]
built=190/600, last_len=3169 tokens

Building long: 33% | [■] 200/600 [01:33<02:45, 2.42it/s]
built=200/600, last_len=2460 tokens

Building long: 35% | [■] 210/600 [01:38<02:26, 2.65it/s]
built=210/600, last_len=2945 tokens

Building long: 37% | [■] 220/600 [01:42<02:06, 3.01it/s]
built=220/600, last_len=3628 tokens

Building long: 38% | [■] 230/600 [01:51<06:17, 1.02s/it]
built=230/600, last_len=3228 tokens

Building long: 40% | [■] 240/600 [01:57<04:22, 1.37it/s]
built=240/600, last_len=3886 tokens

Building long: 42% | [■] 250/600 [02:09<06:26, 1.10s/it]
built=250/600, last_len=3225 tokens

Building long: 43% | [■] 260/600 [02:14<02:30, 2.26it/s]
built=260/600, last_len=3722 tokens

Building long: 45% | [■] 270/600 [02:19<02:49, 1.95it/s]
built=270/600, last_len=2643 tokens

Building long: 47% | [■] 280/600 [02:23<03:09, 1.69it/s]
built=280/600, last_len=4031 tokens

Building long: 48% | [■] 290/600 [02:39<09:20, 1.81s/it]
built=290/600, last_len=2986 tokens

Building long: 50% | [■] 300/600 [02:45<03:48, 1.32it/s]
built=300/600, last_len=3837 tokens

Building long: 52% | [■] 310/600 [02:57<04:38, 1.04it/s]
built=310/600, last_len=3470 tokens

Building long: 53% | [■] 320/600 [03:08<04:19, 1.08it/s]
built=320/600, last_len=3117 tokens

Building long: 55% | [■] 330/600 [03:16<04:41, 1.04s/it]
built=330/600, last_len=3086 tokens

Building long: 57% | [■] 340/600 [03:23<02:40, 1.62it/s]
built=340/600, last_len=3688 tokens

```
Building long: 58%|██████| 350/600 [03:29<03:55, 1.06it/s]
built=350/600, last_len=3755 tokens
Building long: 60%|██████| 360/600 [03:39<03:56, 1.01it/s]
built=360/600, last_len=2462 tokens
Building long: 62%|██████| 370/600 [03:44<02:14, 1.71it/s]
built=370/600, last_len=3850 tokens
Building long: 63%|██████| 380/600 [03:57<04:01, 1.10s/it]
built=380/600, last_len=2184 tokens
Building long: 65%|██████| 390/600 [04:02<01:47, 1.95it/s]
built=390/600, last_len=3645 tokens
Building long: 67%|██████| 400/600 [04:09<02:21, 1.41it/s]
built=400/600, last_len=3122 tokens
Building long: 68%|██████| 410/600 [04:15<02:58, 1.07it/s]
built=410/600, last_len=3959 tokens
Building long: 70%|██████| 420/600 [04:27<02:44, 1.09it/s]
built=420/600, last_len=3146 tokens
Building long: 72%|██████| 430/600 [04:31<01:41, 1.68it/s]
built=430/600, last_len=3866 tokens
Building long: 73%|██████| 440/600 [04:38<01:38, 1.62it/s]
built=440/600, last_len=3164 tokens
Building long: 75%|██████| 450/600 [04:45<02:02, 1.22it/s]
built=450/600, last_len=3925 tokens
Building long: 77%|██████| 460/600 [04:51<01:14, 1.88it/s]
built=460/600, last_len=3569 tokens
Building long: 78%|██████| 470/600 [04:57<01:39, 1.31it/s]
built=470/600, last_len=3912 tokens
Building long: 80%|██████| 480/600 [05:03<01:08, 1.76it/s]
built=480/600, last_len=3028 tokens
Building long: 82%|██████| 490/600 [05:07<00:38, 2.87it/s]
built=490/600, last_len=2081 tokens
Building long: 83%|██████| 500/600 [05:10<00:38, 2.61it/s]
built=500/600, last_len=2345 tokens
Building long: 85%|██████| 510/600 [05:22<01:22, 1.09it/s]
built=510/600, last_len=2541 tokens
Building long: 87%|██████| 520/600 [05:34<01:55, 1.44s/it]
built=520/600, last_len=2754 tokens
Building long: 88%|██████| 530/600 [05:48<02:32, 2.17s/it]
built=530/600, last_len=3959 tokens
Building long: 90%|██████| 540/600 [05:58<01:14, 1.24s/it]
built=540/600, last_len=3576 tokens
Building long: 92%|██████| 550/600 [06:09<00:30, 1.62it/s]
built=550/600, last_len=2163 tokens
Building long: 93%|██████| 560/600 [06:21<00:43, 1.09s/it]
built=560/600, last_len=2606 tokens
Building long: 95%|██████| 570/600 [06:33<00:53, 1.78s/it]
built=570/600, last_len=3850 tokens
Building long: 97%|██████| 580/600 [06:39<00:09, 2.12it/s]
built=580/600, last_len=3674 tokens
Building long: 98%|██████| 590/600 [06:42<00:03, 2.52it/s]
built=590/600, last_len=3438 tokens
Building long: 100%|██████| 600/600 [06:48<00:00, 1.47it/s]
```

```
built=600/600, last_len=2664 tokens
long_bucket: 600
```

Create Mixed Buckets

```
In [ ]: from tqdm.auto import tqdm

TOTAL_REQS_PER_FILE = 400
SEED = 999

OVERWRITE_LOCAL = False
OVERWRITE_DRIVE = False

ENABLE_LONG_HEAVY = True
HEAVY_LONG_RANGE = (3500, 4096)
WANT_LONG_HEAVY = 600
USE_LONG_HEAVY_FOR = {"mix_60_40.jsonl", "mix_50_50.jsonl"}

MIX_SPECS = {
    "short_only.jsonl": {"short": 1.0},
    "medium_only.jsonl": {"medium": 1.0},
    "long_only.jsonl": {"long": 1.0},

    "mix_90_10.jsonl": {"short": 0.9, "long": 0.1},
    "mix_70_30.jsonl": {"short": 0.7, "long": 0.3},
    "mix_70_20_10.jsonl": {"short": 0.7, "medium": 0.2, "long": 0.1},

    # these use long heavy
    "mix_60_40.jsonl": {"short": 0.6, "long": 0.4},
    "mix_50_50.jsonl": {"short": 0.5, "long": 0.5},
}

assert "tok" in globals(), "Tokenizer 'tok' not found."
assert "short_bucket" in globals() and len(short_bucket) > 0, "short_bucket"
assert "medium_bucket" in globals() and len(medium_bucket) > 0, "medium_bucket"
assert "long_bucket" in globals() and len(long_bucket) > 0, "long_bucket"

def tok_len(text: str) -> int:
    return len(tok.encode(text, add_special_tokens=False))

def build_long_prompt(target_tokens, rng, snippets_pool):
    parts = []
    used = set()
    while True:
        idx = rng.randrange(len(snippets_pool))
        if idx in used:
            continue
        used.add(idx)

        parts.append(snippets_pool[idx].replace("\n", " ").strip())
        text = "\n\n".join(parts)

        if tok_len(text) >= target_tokens:
            break

    if len(used) > min(2000, len(snippets_pool) - 1):
```

```

        used.clear()

        ids = tok.encode(text, add_special_tokens=False) [:target_tokens]
        return tok.decode(ids)

def build_long_bucket(want_n, token_range, snippets_pool, seed=1234, desc=""):
    rng = random.Random(seed)
    bucket = []

    for _ in tqdm(range(want_n), desc=f"Building {desc} {token_range}"):
        target = rng.randint(token_range[0], token_range[1])
        bucket.append(build_long_prompt(target, rng, snippets_pool))

    return bucket

long_heavy_bucket = None
if ENABLE_LONG_HEAVY and len(USE_LONG_HEAVY_FOR) > 0:
    snippets_pool = short_bucket + medium_bucket
    long_heavy_bucket = build_long_bucket(
        want_n=WANT_LONG_HEAVY,
        token_range=HEAVY_LONG_RANGE,
        snippets_pool=snippets_pool,
        seed=1234,
        desc="long_heavy_bucket",
    )

    print(
        "long_heavy_bucket:",
        len(long_heavy_bucket),
        "example_len_tokens:",
        tok_len(long_heavy_bucket[0]),
    )

def write_jsonl(path, prompts, max_tokens=128, temperature=0.0):
    with open(path, "w", encoding="utf-8") as f:
        for p in prompts:
            row = {
                "messages": [{"role": "user", "content": p}],
                "max_tokens": max_tokens,
                "temperature": temperature,
            }
            f.write(json.dumps(row) + "\n")

def sample_prompts(bucket, k, rng):
    if k <= 0:
        return []
    if len(bucket) < k:
        raise ValueError(f"Bucket too small: need {k}, have {len(bucket)}")
    return rng.sample(bucket, k)

def normalize_ratios(ratios: dict) -> dict:
    r = dict(ratios)
    s = float(r.get("short", 0.0))
    m = float(r.get("medium", 0.0))
    l = float(r.get("long", 0.0))
    total = s + m + l
    if total == 0:
        return ratios
    else:
        return {key: value / total for key, value in ratios.items()}

```

```

if total > 1.000001:
    raise ValueError(f"Ratios sum > 1.0: {ratios}")

if "long" not in r:
    r["long"] = 1.0 - (s + m)

return r

def make_mix(total, ratios, short_b, medium_b, long_b):
    ratios = normalize_ratios(ratios)
    k_short = int(total * ratios.get("short", 0))
    k_medium = int(total * ratios.get("medium", 0))
    k_long = total - k_short - k_medium

    prompts = []
    prompts += sample_prompts(short_b, k_short, rng)
    prompts += sample_prompts(medium_b, k_medium, rng)
    prompts += sample_prompts(long_b, k_long, rng)

    rng.shuffle(prompts)

    return prompts

def should_write(path, overwrite=False):
    return overwrite or (not os.path.exists(path))

def safe_copy(local_path, drive_path, overwrite=False):
    if not overwrite and os.path.exists(drive_path):
        return False

    shutil.copy(local_path, drive_path)

    return True

rng = random.Random(SEED)

written = []
skipped = []

for fn, ratios in MIX_SPECS.items():
    local_path = os.path.join(OUT_DIR, fn)
    drive_path = os.path.join(DRIVE_OUT_DIR, fn)

    if (not OVERWRITE_LOCAL) and os.path.exists(local_path) and (not OVERWRITE_DRIVE):
        skipped.append(fn)
        continue

    use_heavy = (fn in USE_LONG_HEAVY_FOR) and (long_heavy_bucket is not None)
    long_source = long_heavy_bucket if use_heavy else long_bucket

    prompts = make_mix(
        total=TOTAL_REQS_PER_FILE,
        ratios=ratios,
        rng=rng,
        short_b=short_bucket,
        medium_b=medium_bucket,
        long_b=long_source,
    )

```

```

if should_write(local_path, overwrite=OVERWRITE_LOCAL):
    write_jsonl(local_path, prompts)
else:
    skipped.append(fn)
    continue

copied = safe_copy(local_path, drive_path, overwrite=OVERWRITE_DRIVE)
if not copied:
    print("Exists on Drive, skipping copy:", drive_path)

written.append(fn)

print("Done.")
print("Written:", written)
print("Skipped:", skipped)
print("Local dir:", OUT_DIR)
print("Drive dir:", DRIVE_OUT_DIR)

```

```

Building long_heavy_bucket (3500, 4096):  0%|          | 0/600 [00:00<?, ?it/s]
long_heavy_bucket: 600 example_len_tokens: 3951
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/short_only.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/medium_only.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/long_only.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/mix_90_10.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/mix_70_30.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/mix_70_20_10.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/mix_60_40.jsonl
Exists on Drive, skipping copy: /content/drive/MyDrive/llm_loadtest_workloads/mix_50_50.jsonl
Done.
Written: ['short_only.jsonl', 'medium_only.jsonl', 'long_only.jsonl', 'mix_90_10.jsonl', 'mix_70_30.jsonl', 'mix_70_20_10.jsonl', 'mix_60_40.jsonl', 'mix_50_50.jsonl']
Skipped: []
Local dir: /content/workloads
Drive dir: /content/drive/MyDrive/llm_loadtest_workloads

```

```

In [ ]: WORKLOAD_DIR = "/content/workloads"
FILES = sorted([f for f in os.listdir(WORKLOAD_DIR) if f.endswith(".jsonl")]

def tok_len(text: str) -> int:
    return len(tok.encode(text, add_special_tokens=False))

def read_jsonl(path: str):
    rows = []

    with open(path, "r", encoding="utf-8") as f:
        for line in f:
            if not line.strip():
                continue

```

```

        rows.append(json.loads(line))

    return rows

def extract_prompt(row: dict) -> str:
    msgs = row.get("messages", [])
    if isinstance(msgs, list) and len(msgs) > 0 and isinstance(msgs[0], dict):
        return (msgs[0].get("content") or "").strip()

    return ""

def stats(arr):
    arr = np.array(arr, dtype=float)

    if len(arr) == 0:
        return {}

    return {
        "n": int(len(arr)),
        "min": float(np.min(arr)),
        "p50": float(np.percentile(arr, 50)),
        "p90": float(np.percentile(arr, 90)),
        "p95": float(np.percentile(arr, 95)),
        "p99": float(np.percentile(arr, 99)),
        "max": float(np.max(arr)),
        "mean": float(np.mean(arr)),
    }

results = []
per_file_details = {}

for fn in tqdm(FILES, desc="Verifying workload JSONLs"):
    path = os.path.join(WORKLOAD_DIR, fn)
    rows = read_jsonl(path)

    prompts = [extract_prompt(r) for r in rows]
    prompts = [p for p in prompts if p] # drop empties

    tok_lens = [tok_len(p) for p in prompts]
    char_lens = [len(p) for p in prompts]

    s_tok = stats(tok_lens)
    s_chr = stats(char_lens)

    results.append({
        "file": fn,
        "n": s_tok.get("n", 0),
        "tok_min": s_tok.get("min"),
        "tok_p50": s_tok.get("p50"),
        "tok_p90": s_tok.get("p90"),
        "tok_p95": s_tok.get("p95"),
        "tok_p99": s_tok.get("p99"),
        "tok_max": s_tok.get("max"),
        "tok_mean": s_tok.get("mean"),
        "char_p50": s_chr.get("p50"),
        "char_p99": s_chr.get("p99"),
    })

```

```
# store some examples for sanity checks
if len(tok_lens) > 0:
    idx_sorted = np.argsort(tok_lens)
    smallest = idx_sorted[:2].tolist()
    largest = idx_sorted[-2:].tolist()
    per_file_details[fn] = {
        "smallest_examples": [(tok_lens[i], prompts[i][:200]) for i in range(2)],
        "largest_examples": [(tok_lens[i], prompts[i][:200]) for i in range(-2, 0)]}

df_stats = pd.DataFrame(results).sort_values("file")
display(df_stats)

print("\n--- Sanity samples (first 200 chars) ---")
for fn, det in per_file_details.items():
    print(f"\n### {fn}")
    print("Smallest:")
    for L, s in det["smallest_examples"]:
        print(f"  tok={L} | {s[:200]}")
    print("Largest:")
    for L, s in det["largest_examples"]:
        print(f"  tok={L} | {s[:200]})")
```

Verifying workload JSONNLs: 0% | 0/8 [00:00<?, ?it/s]

	file	n	tok_min	tok_p50	tok_p90	tok_p95	tok_p99	tok_max
0	long_only.jsonl	400	2051.0	3068.5	3931.1	4028.10	4083.01	4092.0
1	medium_only.jsonl	400	128.0	169.0	231.1	240.05	253.00	256.0
2	mix_50_50.jsonl	400	8.0	1761.0	4018.1	4062.10	4090.02	4096.0
3	mix_60_40.jsonl	400	8.0	12.5	3983.1	4042.10	4089.01	4095.0
4	mix_70_20_10.jsonl	400	8.0	11.0	432.9	3189.05	3898.31	4090.0
5	mix_70_30.jsonl	400	8.0	11.0	3525.5	3852.10	4070.13	4092.0
6	mix_90_10.jsonl	400	8.0	10.0	239.1	3332.80	3887.60	4090.0
7	short_only.jsonl	400	8.0	10.0	14.0	15.00	19.01	23.0

--- Sanity samples (first 200 chars) ---

long_only.jsonl

Smallest:

tok=2051 | 'who wrote the song come thou fount of every blessing\n\nCan you summarize the setting and atmosphere at Sue Ann Staff Estate Winery Wddings? Answer according to: Surrounded by lush vineyards and perche'

tok=2058 | 'Write a fast-paced, plot-driven thriller novel that chronicles the story of a young woman abducted by a notoriously insidious serial killer. The narrative style should be highly tense and suspenseful,'

Largest:

tok=4090 | "elvis presley daddy don't you walk so fast\n\nhow old do you have to be to get a tattoo in utah\n\nbruce lee hall of mirrors enter the dragon\n\nWrite a 75,000-word historical romance novel set during the H"

tok=4092 | 'when was the last time a total eclipse happened\n\nWrite a detailed report that analyzes and summarizes the various ways in which technology has influenced and transformed the manufacturing industry. Yo'

medium_only.jsonl

Smallest:

tok=128 | 'Great way to get your message out there, "Happy Birthday", "Will You Marry Me" and More!\nLet them know who you are Now!\nGreat way to Brand and be able to change your tagline!\nWe Can Customize to Almost'

tok=128 | 'Implement a queue data structure in Java that can enqueue and dequeue elements. The queue should have a maximum size of n and should be implemented using an array. The enqueue operation should add ele'

Largest:

tok=255 | 'such as. The gruffelo) or poems for children. Do you have any tips? Recommendations? Or would you like to write poems yourself with me? I think poems are an important part of German literature and esp'

tok=256 | '1. Roast and powder coriander seeds, cumin seeds and red chillies.\n2. Grind together ginger-garlic. Grind coconut and poppy seeds separately.\n3. Chop the onions finely.\n4. Put the minced meat into a v'

mix_50_50.jsonl

Smallest:

tok=8 | 'most commonly used punishment in the united states'

tok=8 | 'i m in the love with the coco'

Largest:

tok=4095 | 'who does the voice of shaggy on scooby doo\n\nntom and jerry movie the fast and the furry\n\nIs it possible to have a combination of drawers, cupboards, and wine racks in the same sideboard?: Hand Made any'

tok=4096 | 'deeper than the holler what is a holler\n\nwho played tj on head of the class\nnis the song fight song in a movie\n\nthe hollywood ten who refused to answer questions\n\nfools rush in where ghosts fear to t r'

mix_60_40.jsonl

Smallest:

tok=8 | 'does president have to be born in usa'

tok=8 | 'who wants to marry a prince tv show'

Largest:

tok=4095 | 'Can you provide a brief summary of the central problem facing Dell management according to the HBS Case DELL: NEW HORIZONS?: (2) The following questions are based on the HBS Case DELL: NEW HORIZONS. ('

tok=4095 | 'who does the voice of shaggy on scooby doo\n\nntom and jerry movie the fast and the furry\n\nIs it possible to have a combination of drawers, cupboards, and wine racks in the same sideboard?: Hand Made any'

mix_70_20_10.jsonl

```

Smallest:
tok=8 | 'properties of red black tree in data structure'
tok=8 | 'what are the parts of the finger called'
Largest:
tok=4084 | 'Write a free verse poem that captures the beauty and essence
of each season, drawing inspiration from the natural elements around you.
Incorporate vivid descriptions of the sights, sounds, smells, and'
tok=4090 | "elvis presley daddy don't you walk so fast\n\nhow old do you
have to be to get a tattoo in utah\n\nbruce lee hall of mirrors enter the
dragon\n\nWrite a 75,000-word historical romance novel set during the H"

### mix_70_30.jsonl
Smallest:
tok=8 | 'what number book is mark in the bible'
tok=8 | 'the legend of the monkey king tv series'
Largest:
tok=4090 | "elvis presley daddy don't you walk so fast\n\nhow old do you
have to be to get a tattoo in utah\n\nbruce lee hall of mirrors enter the
dragon\n\nWrite a 75,000-word historical romance novel set during the H"
tok=4092 | 'when was the last time a total eclipse happened\n\nWrite a d
etailed report that analyzes and summarizes the various ways in which tech
nology has influenced and transformed the manufacturing industry. Yo'

### mix_90_10.jsonl
Smallest:
tok=8 | 'why do ice cubes get smaller over time'
tok=8 | 'early blues musicians used which combination of instruments'
Largest:
tok=4032 | "what is the meaning of yanny and laurel\n\nearly blues music
ians used which combination of instruments\n\nCan you please summarize M0
O's services and what makes them unique in the industry? Answer accordi"
tok=4090 | "elvis presley daddy don't you walk so fast\n\nhow old do you
have to be to get a tattoo in utah\n\nbruce lee hall of mirrors enter the
dragon\n\nWrite a 75,000-word historical romance novel set during the H"

### short_only.jsonl
Smallest:
tok=8 | 'who sang take that look off your face'
tok=8 | 'where is the us military base in japan'
Largest:
tok=21 | 'nasa conducted a poll with 9 names to choose the name of the r
over which landed on mars recently'
tok=23 | 'erupted on august 3 1965 as the result of a routine arrest of
a drunk driver'

```

Setting up vLLM server and Dataset loading

```
In [ ]: print("vLLM:", vllm.__version__)
print("torch:", torch.__version__, "cuda:", torch.version.cuda)
print("GPU:", torch.cuda.get_device_name(0))
```

```
vLLM: 0.13.0
torch: 2.9.0+cu126 cuda: 12.6
GPU: NVIDIA L4
```

```
In [ ]: # Use this cell, when restarting the Server again in the same runtime
# it clears out the previous instances prevents issues in server setup
```

```
!pkill -f "vllm.entrypoints.openai.api_server" || true
!pkill -f "EngineCore" || true
!pkill -f "vllm" || true

^C
^C
^C
```

```
In [ ]: MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"
PORT = 8000

MAX_MODEL_LEN = 4096
GPU_UTIL = 0.70
DTYPE = "half" # fp16 weights/compute

# Launch server
# run it in the background and log to a file.
cmd = f"""
nohup python3 -m vllm.entrypoints.openai.api_server \
--model {MODEL_ID} \
--host 127.0.0.1 \
--port {PORT} \
--max-model-len {MAX_MODEL_LEN} \
--gpu-memory-utilization {GPU_UTIL} \
--dtype {DTYPE} \
--enforce-eager \
--disable-log-stats \
> vllm_server.log 2>&1 &
"""

print("Launching vLLM server...")
!bash -lc "$cmd"

print("Server process launched. Showing last 30 log lines:")
!tail -n 30 vllm_server.log
```

Launching vLLM server...
 Server process launched. Showing last 30 log lines:

Function to check the successful startup of the server

```
In [ ]: import time, requests

BASE = "http://127.0.0.1:8000"

def wait_for_server(timeout_s=600):
    t0 = time.time()
    while time.time() - t0 < timeout_s:
        try:
            r = requests.get(f"{BASE}/v1/models", timeout=2)
            if r.status_code == 200:
                return True, r.json()
        except Exception:
            pass
        time.sleep(2)
    return False, None
```

```
ok, payload = wait_for_server()
print("Ready:", ok)
if ok:
    print("Models:", [m["id"] for m in payload.get("data", [])])
else:
    print("Not ready yet. Check logs:")
    !tail -n 60 vllm_server.log
```

Ready: True
Models: ['Qwen/Qwen2.5-3B-Instruct']

```
In [ ]: # Server take almost a minute to start,
# keep checking the logs to confirm successful startup
# before running any experiments

!tail -n 200 vllm_server.log
```

```
2026-01-11 19:14:50.463943: E external/local_xla/xla/stream_executor/cuda/
cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register
factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are writ-
ten to STDERR
E0000 00:00:1768158890.485170    3519 cuda_dnn.cc:8579] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one ha-
s already been registered
E0000 00:00:1768158890.491587    3519 cuda_blas.cc:1407] Unable to registe-
r cuBLAS factory: Attempting to register factory for plugin cuBLAS when on-
e has already been registered
W0000 00:00:1768158890.507836    3519 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
ame target more than once.
W0000 00:00:1768158890.507865    3519 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
ame target more than once.
W0000 00:00:1768158890.507868    3519 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
ame target more than once.
W0000 00:00:1768158890.507870    3519 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
ame target more than once.
(APIServer pid=3519) INFO 01-11 19:14:58 [api_server.py:1351] vLLM API ser-
ver version 0.13.0
(APIServer pid=3519) INFO 01-11 19:14:58 [utils.py:253] non-default args:
{'host': '127.0.0.1', 'model': 'Qwen/Qwen2.5-3B-Instruct', 'dtype': 'hal-
f', 'max_model_len': 4096, 'enforce_eager': True, 'gpu_memory_utilizatio-
n': 0.7, 'disable_log_stats': True}
(APIServer pid=3519) INFO 01-11 19:14:59 [model.py:514] Resolved architect-
ure: Qwen2ForCausalLM
(APIServer pid=3519) WARNING 01-11 19:14:59 [model.py:2005] Casting torch.
bfloating16 to torch.float16.
(APIServer pid=3519) INFO 01-11 19:14:59 [model.py:1661] Using max model l-
en 4096
(APIServer pid=3519) INFO 01-11 19:15:00 [scheduler.py:230] Chunked prefil-
l is enabled with max_num_batched_tokens=2048.
(APIServer pid=3519) WARNING 01-11 19:15:00 [vllm.py:622] Enforce eager se-
t, overriding optimization level to -00
(APIServer pid=3519) INFO 01-11 19:15:00 [vllm.py:722] Cudagraph is disabl-
ed under eager mode
2026-01-11 19:15:08.298361: E external/local_xla/xla/stream_executor/cuda/
cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register
factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are writ-
ten to STDERR
E0000 00:00:1768158908.320069    3649 cuda_dnn.cc:8579] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one ha-
s already been registered
E0000 00:00:1768158908.326565    3649 cuda_blas.cc:1407] Unable to registe-
r cuBLAS factory: Attempting to register factory for plugin cuBLAS when on-
e has already been registered
W0000 00:00:1768158908.342417    3649 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
ame target more than once.
W0000 00:00:1768158908.342450    3649 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
ame target more than once.
W0000 00:00:1768158908.342454    3649 computation_placer.cc:177] computati-
on placer already registered. Please check linkage and avoid linking the s-
```

ame target more than once.

W0000 00:00:1768158908.342457 3649 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.

(EngineCore_DP0 pid=3649) INFO 01-11 19:15:16 [core.py:93] Initializing a V1 LLM engine (v0.13.0) with config: model='Qwen/Qwen2.5-3B-Instruct', speculative_config=None, tokenizer='Qwen/Qwen2.5-3B-Instruct', skip_tokenizer_init=False, tokenizer_mode=auto, revision=None, tokenizer_revision=None, trust_remote_code=False, dtype=torch.float16, max_seq_len=4096, download_dir=None, load_format=auto, tensor_parallel_size=1, pipeline_parallel_size=1, data_parallel_size=1, disable_custom_all_reduce=False, quantization=None, enforce_eager=True, kv_cache_dtype=auto, device_config=cuda, structured_outputs_config=StructuredOutputsConfig(backend='auto', disable_fallback=False, disable_any_whitespace=False, disable_additional_properties=False, reasoning_parser='', reasoning_parser_plugin='', enable_in_reasoning=False), observability_config=ObservabilityConfig(show_hidden_metrics_for_version=None, otlp_traces_endpoint=None, collect_detailed_traces=None, kv_cache_metrics=False, kv_cache_metrics_sample=0.01, cudagraph_metrics=False, enable_layerwise_nvtx_tracing=False), seed=0, served_model_name=Qwen/Qwen2.5-3B-Instruct, enable_prefix_caching=True, enable_chunked_prefill=True, pooler_config=None, compilation_config={'level': None, 'mode': <CompilationMode.NONE: 0>, 'debug_dump_path': None, 'cache_dir': '', 'compile_cache_save_format': 'binary', 'backend': 'inductor', 'custom_ops': ['all'], 'splitting_ops': [], 'compile_mm_encoder': False, 'compile_sizes': [], 'compile_ranges_split_points': [2048], 'inductor_compile_config': {'enable_auto_functionalized_v2': False, 'combo_kernels': True, 'benchmark_combo_kernel': True}, 'inductor_passes': {}}, 'cudagraph_mode': <CUDAGraphMode.NONE: 0>, 'cudagraph_num_of_warmups': 0, 'cudagraph_capture_sizes': [], 'cudagraph_copy_inputs': False, 'cudagraph_specialize_lora': True, 'use_inductor_graph_partition': False, 'pass_config': {'fuse_norm_quant': False, 'fuse_act_quant': False, 'fuse_attn_quant': False, 'eliminate_noops': False, 'enable_sp': False, 'fuse_gemm_comms': False, 'fuse_allreduce_rms': False}, 'max_cudagraph_capture_size': 0, 'dynamic_shapes_config': {'type': <DynamicShapesType.BACKED: 'backed'>, 'evaluate_guards': False}, 'local_cache_dir': None}

(EngineCore_DP0 pid=3649) INFO 01-11 19:15:16 [parallel_state.py:1203] world_size=1 rank=0 local_rank=0 distributed_init_method=tcp://172.28.0.12:58367 backend=nccl

(EngineCore_DP0 pid=3649) INFO 01-11 19:15:16 [parallel_state.py:1411] rank 0 in world size 1 is assigned as DP rank 0, PP rank 0, PCP rank 0, TP rank 0, EP rank 0

(EngineCore_DP0 pid=3649) INFO 01-11 19:15:17 [gpu_model_runner.py:3562] Starting to load model Qwen/Qwen2.5-3B-Instruct...

(EngineCore_DP0 pid=3649) INFO 01-11 19:15:18 [cuda.py:351] Using FLASH_ATTN attention backend out of potential backends: ('FLASH_ATTN', 'FLASHINFER', 'TRITON_ATTN', 'FLEX_ATTENTION')

(EngineCore_DP0 pid=3649)

Loading safetensors checkpoint shards: 0% Completed | 0/2 [00:00<?, ?its]

(EngineCore_DP0 pid=3649)

Loading safetensors checkpoint shards: 50% Completed | 1/2 [00:01<00:01, 1.16s/it]

(EngineCore_DP0 pid=3649)

Loading safetensors checkpoint shards: 100% Completed | 2/2 [00:03<00:00, 1.67s/it]

(EngineCore_DP0 pid=3649)

Loading safetensors checkpoint shards: 100% Completed | 2/2 [00:03<00:00, 1.59s/it]

(EngineCore_DP0 pid=3649)

(EngineCore_DP0 pid=3649) INFO 01-11 19:15:21 [default_loader.py:308] Load

```
ing weights took 3.27 seconds
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:22 [gpu_model_runner.py:3659] Model loading took 5.7916 GiB memory and 4.039437 seconds
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:24 [gpu_worker.py:375] Available KV cache memory: 8.29 GiB
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:25 [kv_cache_utils.py:1291] GPU KV cache size: 241,456 tokens
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:25 [kv_cache_utils.py:1296] Maximum concurrency for 4,096 tokens per request: 58.95x
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:25 [core.py:259] init engine (profile, create kv cache, warmup model) took 2.45 seconds
(EngineCore_DP0 pid=3649) WARNING 01-11 19:15:26 [vllm.py:629] Inductor compilation was disabled by user settings, Optimizations settings that are only active during Inductor compilation will be ignored.
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:26 [vllm.py:722] Cudagraph is disabled under eager mode
(APIServer pid=3519) INFO 01-11 19:15:26 [api_server.py:1099] Supported tasks: ['generate']
(APIServer pid=3519) WARNING 01-11 19:15:26 [model.py:1487] Default sampling parameters have been overridden by the model's Hugging Face generation config recommended from the model creator. If this is not intended, please relaunch vLLM instance with `--generation-config vllm`.
(APIServer pid=3519) INFO 01-11 19:15:26 [serving_responses.py:201] Using default chat sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=3519) INFO 01-11 19:15:26 [serving_chat.py:137] Using default chat sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=3519) INFO 01-11 19:15:27 [serving_completion.py:77] Using default completion sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=3519) INFO 01-11 19:15:27 [serving_chat.py:137] Using default chat sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=3519) INFO 01-11 19:15:27 [api_server.py:1425] Starting vLLM API server 0 on http://127.0.0.1:8000
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:38] Available routes are:
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /openapi.json, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /docs, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /docs/oauth2-redirect, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /redoc, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /scale_elastic_ep, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /is_scaling_elastic_ep, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /tokenize, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /detokenize, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /inference/v1/generate, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /pause, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /resume, Methods: POST
```

```
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /is_pause
d, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /metrics,
Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /health,
Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /load, Me
thods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/model
s, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /version,
Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/respo
nses, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/respo
nses/{response_id}, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/respo
nses/{response_id}/cancel, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/messa
ges, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/chat/
completions, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/compl
etions, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/audi
o/transcriptions, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/audi
o/translations, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /ping, Me
thods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /ping, Me
thods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /invocati
ons, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /classif
y, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/embed
dings, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /score, M
ethods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /score,
Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /rerank,
Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/reran
k, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v2/reran
k, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /pooling,
Methods: POST
(APIServer pid=3519) INFO:      Started server process [3519]
(APIServer pid=3519) INFO:      Waiting for application startup.
(APIServer pid=3519) INFO:      Application startup complete.
```

Sanity check after Server start

```
In [ ]: import requests, json
        BASE="http://127.0.0.1:8000"
```

```

payload = {
    "model": "Qwen/Qwen2.5-3B-Instruct",
    "messages": [{"role": "user", "content": "Say hi in one sentence."}],
    "max_tokens": 32,
    "temperature": 0.0
}

r = requests.post(f"{BASE}/v1/chat/completions", json=payload, timeout=12)
print(r.status_code)
print(r.json()["choices"][0]["message"]["content"])

```

200

Hi there! How can I assist you today?

In []: # check the GPU consumption once the vLLM server is up
! nvidia-smi

```

Sun Jan 11 19:55:23 2026
+-----
-----+
| NVIDIA-SMI 550.54.15           Driver Version: 550.54.15      CUDA Ve
rsion: 12.4          |
|-----+-----+-----+
| GPU  Name                  Persistence-M | Bus-Id        Disp.A | Volat
ile Uncorr. ECC |
| Fan  Temp     Perf          Pwr:Usage/Cap |             Memory-Usage | GPU-U
til  Compute M. |
|          |                               |               |
MIG M. |
|=====+=====+=====+=====+=====+=====+=====|
|     0  NVIDIA L4              Off  | 00000000:00:03.0 Off |
0 |
| N/A   56C     P0            30W / 72W | 16225MiB / 23034MiB |
0%     Default |
|          |                               |
N/A |
|-----+-----+-----+
+-----+
-----+
| Processes:
|
| GPU  GI  CI          PID  Type  Process name
GPU Memory |
|          ID  ID
Usage       |
|=====+=====+=====+=====+=====+=====+=====|
|-----+-----+-----+
+-----+
-----+

```

Load the dataset and run tests

In []: from google.colab import drive
drive.mount("/content/drive")

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Load the dataset from Drive

```
In [ ]: import json, os, random
from pathlib import Path

PROMPT_DIR = Path("/content/drive/MyDrive/llm_loadtest_workloads")

FILES = [
    "short_only.jsonl",
    "medium_only.jsonl",
    "long_only.jsonl",
    "mix_90_10.jsonl",
    "mix_70_30.jsonl",
    "mix_70_20_10.jsonl",
    "mix_60_40.jsonl",
    "mix_50_50.jsonl"
]

def load_jsonl_prompts(path: Path, limit=None, seed=123):
    prompts = []
    with path.open("r", encoding="utf-8") as f:
        for line in f:
            obj = json.loads(line)

            if "prompt" in obj:
                prompts.append(obj["prompt"])

            elif "messages" in obj:
                users = [m["content"] for m in obj["messages"] if m.get("user")]
                if users:
                    prompts.append(users[-1])

            if limit and len(prompts) >= limit:
                break

    random.Random(seed).shuffle(prompts)
    return prompts

prompt_sets = {}
for fn in FILES:
    p = PROMPT_DIR / fn
    if p.exists():
        prompt_sets[fn] = load_jsonl_prompts(p)
        print(f"Loaded {fn}: {len(prompt_sets[fn])} prompts")
    else:
        print(f"Missing: {p}")
```

```
Loaded short_only.jsonl: 400 prompts
Loaded medium_only.jsonl: 400 prompts
Loaded long_only.jsonl: 400 prompts
Loaded mix_90_10.jsonl: 400 prompts
Loaded mix_70_30.jsonl: 400 prompts
Loaded mix_70_20_10.jsonl: 400 prompts
Loaded mix_60_40.jsonl: 400 prompts
Loaded mix_50_50.jsonl: 400 prompts
```

Checking if the dataset is properly read

```
In [ ]: from transformers import AutoTokenizer
import numpy as np, random

TOK_MODEL = "Qwen/Qwen2.5-3B-Instruct"
tok = AutoTokenizer.from_pretrained(TOK_MODEL, use_fast=True)

def tok_len(s: str) -> int:
    return len(tok.encode(s, add_special_tokens=False))

def summarize_lengths(name, prompts, n=200):
    sample = random.sample(prompts, min(n, len(prompts)))
    lens = np.array([tok_len(p) for p in sample])
    print(f"{name:18s} n={len(sample):4d} min={lens.min():4d} p50={int(lens[lens.argsort()][n//2]):4d} max={lens.max():4d}")

for fn, prompts in prompt_sets.items():
    summarize_lengths(fn, prompts)

short_only.jsonl      n= 200  min=   8  p50=   10  p90=   13  p99=   16  max=
16
medium_only.jsonl    n= 200  min= 128  p50= 166  p90= 230  p99= 246  max= 2
55
long_only.jsonl       n= 200  min=2051  p50=3196  p90=3949  p99=4062  max=40
93
mix_90_10.jsonl       n= 200  min=   8  p50=   10  p90=2319  p99=3944  max=40
02
mix_70_30.jsonl       n= 200  min=   8  p50=   11  p90=3312  p99=3924  max=40
31
mix_70_20_10.jsonl    n= 200  min=   8  p50=   11  p90= 231  p99=3660  max=38
33
mix_60_40.jsonl       n= 200  min=   8  p50=   12  p90=3970  p99=4082  max=40
90
mix_50_50.jsonl       n= 200  min=   8  p50=3510  p90=4002  p99=4084  max=40
96
```

Experiment 1

In this experiment, we stress-test vLLM serving under mixed prompt-length workloads to reproduce two real-world issues:

- (1) context-budget failures where long prompts combined with a fixed max_tokens cause requests to error out (HTTP 400 max token limit), and
- (2) degraded serving performance as the fraction of long prompts increases. We run controlled workloads (short-only and mixed ratios like 90/10 and 70/30), then compare success rate, throughput (RPS), and latency metrics (p50/p90/p99) to quantify how long prompts amplify tail latency and reduce effective throughput.

This sets up the next experiments where we mitigate the failure mode (dynamic max_tokens / context budgeting) and run experiments to see how can we configure our servers to perform better under mixed loads.

Single-Request Client Helper (Latency + Token Usage)

```
In [ ]: import time, requests

BASE = "http://127.0.0.1:8000"
MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

def one_request(prompt, max_tokens=128, temperature=0.0, timeout=120):
    t0 = time.perf_counter()
    payload = {
        "model": MODEL_ID,
        "messages": [{"role": "user", "content": prompt}],
        "max_tokens": max_tokens,
        "temperature": temperature,
    }
    r = requests.post(f"{BASE}/v1/chat/completions", json=payload, timeout=timeout)
    latency = time.perf_counter() - t0

    # Hard fail with useful error
    if r.status_code != 200:
        txt = r.text[:500]
        raise RuntimeError(f"HTTP {r.status_code}: {txt}")

    j = r.json()
    usage = j.get("usage", {})
    return {
        "latency_s": latency,
        "prompt_tokens": usage.get("prompt_tokens"),
        "completion_tokens": usage.get("completion_tokens"),
        "total_tokens": usage.get("total_tokens"),
        "status_code": r.status_code,
    }
```

Some checks before starting the experiment

```
In [ ]: # Is the server process actually running?  
!ps aux | grep -E "vllm\.entrypoints\.openai\.api_server|api_server" | gr  
  
# Is anything listening on port 8000, and which process owns it?  
!ss -ltnp | grep ":8000" || true  
  
# What did the server log say most recently?  
!tail -n 80 vllm_server.log || true  
  
print(requests.get("http://127.0.0.1:8000/v1/models").status_code)  
print(requests.get("http://127.0.0.1:8000/v1/models").json())
```

```

root      3519  0.6  3.1 14040592 1743404 ?    Sl  19:14  0:24 python3
-m vllm.entrypoints.openai.api_server --model Qwen/Qwen2.5-3B-Instruct --host 127.0.0.1 --port 8000 --max-model-len 4096 --gpu-memory-utilization 0.7 --dtype half --enforce-eager --disable-log-stats
LISTEN 0      2048          127.0.0.1:8000        0.0.0.0:*    users:(("python
3",pid=3519,fd=33))
(APIServer pid=3519) INFO 01-11 19:14:59 [model.py:1661] Using max model len 4096
(APIServer pid=3519) INFO 01-11 19:15:00 [scheduler.py:230] Chunked prefill is enabled with max_num_batched_tokens=2048.
(APIServer pid=3519) WARNING 01-11 19:15:00 [vllm.py:622] Enforce eager set, overriding optimization level to -00
(APIServer pid=3519) INFO 01-11 19:15:00 [vllm.py:722] Cudagraph is disabled under eager mode
2026-01-11 19:15:08.298361: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are written to STDERR
E0000 00:00:1768158908.320069    3649 cuda_dnn.cc:8579] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered
E0000 00:00:1768158908.326565    3649 cuda_blas.cc:1407] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered
W0000 00:00:1768158908.342417    3649 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1768158908.342450    3649 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1768158908.342454    3649 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
W0000 00:00:1768158908.342457    3649 computation_placer.cc:177] computation placer already registered. Please check linkage and avoid linking the same target more than once.
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:16 [core.py:93] Initializing a V1 LLM engine (v0.13.0) with config: model='Qwen/Qwen2.5-3B-Instruct', speculative_config=None, tokenizer='Qwen/Qwen2.5-3B-Instruct', skip_tokenizer_init=False, tokenizer_mode=auto, revision=None, tokenizer_revision=None, trust_remote_code=False, dtype=torch.float16, max_seq_len=4096, download_dir=None, load_format=auto, tensor_parallel_size=1, pipeline_parallel_size=1, data_parallel_size=1, disable_custom_all_reduce=False, quantization=None, enforce_eager=True, kv_cache_dtype=auto, device_config=cuda, structured_outputs_config=StructuredOutputsConfig(backend='auto', disable_fallback=False, disable_any_whitespace=False, disable_additional_properties=False, reasoning_parser='', reasoning_parser_plugin='', enable_in_reasoning=False), observability_config=ObservabilityConfig(show_hidden_metrics_for_version=None, otlp_traces_endpoint=None, collect_detailed_traces=None, kv_cache_metrics=False, kv_cache_metrics_sample=0.01, cudagraph_metrics=False, enable_layerwise_nvtx_tracing=False), seed=0, served_model_name=Qwen/Qwen2.5-3B-Instruct, enable_prefix_caching=True, enable_chunked_prefill=True, pooler_config=None, compilation_config={'level': None, 'mode': <CompilationMode.NONE: 0>, 'debug_dump_path': None, 'cache_dir': '', 'compile_cache_save_format': 'binary', 'backend': 'inductor', 'custom_ops': ['all'], 'splitting_ops': [], 'compile_mm_encoder': False, 'compile_sizes': [], 'compile_ranges_split_points': [2048], 'inductor_compile_config': {'enable_auto_functionalized_v2': False, 'combo_kernels': True, 'benchmark_combo_kernel': True}, 'inductor_passes': {}}, 'cudagraph_mode': <CUADAGraphMode.NONE: 0>, 'cud

```

```
agraph_num_of_warmups': 0, 'cudagraph_capture_sizes': [], 'cudagraph_copy_inputs': False, 'cudagraph_specialize_lora': True, 'use_inductor_graph_partition': False, 'pass_config': {'fuse_norm_quant': False, 'fuse_act_quant': False, 'fuse_attn_quant': False, 'eliminate_noops': False, 'enable_sp': False, 'fuse_gemm_comms': False, 'fuse_allreduce_rms': False}, 'max_cudagraph_capture_size': 0, 'dynamic_shapes_config': {'type': <DynamicShapes Type.BACKED: 'backed'>, 'evaluate_guards': False}, 'local_cache_dir': None}  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:16 [parallel_state.py:1203] world_size=1 rank=0 local_rank=0 distributed_init_method=tcp://172.28.0.12:58367 backend=nccl  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:16 [parallel_state.py:1411] rank 0 in world size 1 is assigned as DP rank 0, PP rank 0, PCP rank 0, TP rank 0, EP rank 0  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:17 [gpu_model_runner.py:3562] Starting to load model Qwen/Qwen2.5-3B-Instruct...  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:18 [cuda.py:351] Using FLASH_ATTN attention backend out of potential backends: ('FLASH_ATTN', 'FLASHINFER', 'TRITON_ATTN', 'FLEX_ATTENTION')  
Loading safetensors checkpoint shards: 0% Completed | 0/2 [00:00<?, ?its]  
Loading safetensors checkpoint shards: 50% Completed | 1/2 [00:01<00:01, 1.16s/it]  
Loading safetensors checkpoint shards: 100% Completed | 2/2 [00:03<00:00, 1.67s/it]  
Loading safetensors checkpoint shards: 100% Completed | 2/2 [00:03<00:00, 1.59s/it]  
(EngineCore_DP0 pid=3649)  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:21 [default_loader.py:308] Loading weights took 3.27 seconds  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:22 [gpu_model_runner.py:3659] Model loading took 5.7916 GiB memory and 4.039437 seconds  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:24 [gpu_worker.py:375] Available KV cache memory: 8.29 GiB  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:25 [kv_cache_utils.py:1291] GPU KV cache size: 241,456 tokens  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:25 [kv_cache_utils.py:1296] Maximum concurrency for 4,096 tokens per request: 58.95x  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:25 [core.py:259] init engine (profile, create kv cache, warmup model) took 2.45 seconds  
(EngineCore_DP0 pid=3649) WARNING 01-11 19:15:26 [vllm.py:629] Inductor compilation was disabled by user settings, Optimizations settings that are only active during Inductor compilation will be ignored.  
(EngineCore_DP0 pid=3649) INFO 01-11 19:15:26 [vllm.py:722] Cudagraph is disabled under eager mode  
(APIServer pid=3519) INFO 01-11 19:15:26 [api_server.py:1099] Supported tasks: ['generate']  
(APIServer pid=3519) WARNING 01-11 19:15:26 [model.py:1487] Default sampling parameters have been overridden by the model's Hugging Face generation config recommended from the model creator. If this is not intended, please relaunch vLLM instance with `--generation-config vllm`.  
(APIServer pid=3519) INFO 01-11 19:15:26 [serving_responses.py:201] Using default chat sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}  
(APIServer pid=3519) INFO 01-11 19:15:26 [serving_chat.py:137] Using default chat sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}  
(APIServer pid=3519) INFO 01-11 19:15:27 [serving_completion.py:77] Using default completion sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
```

```
(APIServer pid=3519) INFO 01-11 19:15:27 [serving_chat.py:137] Using default chat sampling params from model: {'repetition_penalty': 1.05, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=3519) INFO 01-11 19:15:27 [api_server.py:1425] Starting vLLM API server 0 on http://127.0.0.1:8000
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:38] Available routes are:
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /openapi.json, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /docs, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /docs/oauth2-redirect, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /redoc, Methods: GET, HEAD
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /scale_elastic_ep, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /is_scaling_elastic_ep, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /tokenize, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /detokenize, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /inference/v1/generate, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /pause, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /resume, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /is_pause, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /metrics, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /health, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /load, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/models, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /version, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/responses, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/responses/{response_id}, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/responses/{response_id}/cancel, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/messages, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/chat/completions, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/completions, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/audio/transcriptions, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/audio/translations, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /ping, Methods: GET
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /ping, Me
```

```

thods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /invocations, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /classify, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/embeddings, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /score, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/score, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /rerank, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v1/rerank, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /v2/rerank, Methods: POST
(APIServer pid=3519) INFO 01-11 19:15:27 [launcher.py:46] Route: /pooling, Methods: POST
(APIServer pid=3519) INFO:      Started server process [3519]
(APIServer pid=3519) INFO:      Waiting for application startup.
(APIServer pid=3519) INFO:      Application startup complete.
(APIServer pid=3519) INFO 01-11 19:37:27 [chat_utils.py:590] Detected the chat template content format to be 'string'. You can set `--chat-template-content-format` to override this.
(APIServer pid=3519) INFO:      127.0.0.1:48736 - "POST /v1/chat/completions HTTP/1.1" 200 OK
(APIServer pid=3519) INFO:      127.0.0.1:45104 - "GET /v1/models HTTP/1.1" 200 OK
200
{
  'object': 'list',
  'data': [
    {
      'id': 'Qwen/Qwen2.5-3B-Instruct',
      'object': 'model',
      'created': 1768162789,
      'owned_by': 'vllm',
      'root': 'Qwen/Qwen2.5-3B-Instruct',
      'parent': None,
      'max_model_len': 4096,
      'permission': [
        {
          'id': 'modelperm-ba9025123a14f02f',
          'object': 'model_permission',
          'created': 1768162789,
          'allow_create_engine': False,
          'allow_sampling': True,
          'allow_logprobs': True,
          'allow_search_indices': False,
          'allow_view': True,
          'allow_fine_tuning': False,
          'organization': '*',
          'group': None,
          'is_blocking': False
        }
      ]
    }
  ]
}

```

Server Health Check + Verify vLLM OpenAI Endpoint Is Working

```

In [ ]: import requests, json

BASE = "http://127.0.0.1:8000"
MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

payload = {
    "model": MODEL_ID,
    "messages": [{"role": "user", "content": "Say hello in one sentence."},
    {"max_tokens": 32,
     "temperature": 0.0,
    }]
}

r = requests.post(f"{BASE}/v1/chat/completions", json=payload, timeout=12)
print("status:", r.status_code)
print("text:", r.text[:400])

```

```

status: 200
text: {"id":"chatcmpl-bc2f4e8695ca898b","object":"chat.completion","create
d":1768163326,"model":"Qwen/Qwen2.5-3B-Instruct","choices":[{"index":0,"me
ssage":{"role":"assistant","content":"Hello! How can I assist you toda
y?"}, "refusal":null, "annotations":null, "audio":null, "function_call":null, "t
ool_calls":[],"reasoning":null, "reasoning_content":null}, "logprobs":nul
l, "finish_reason":"stop", "stop_reason":null

```

Single-Request Sanity Check (Real Workload Prompt)

```
In [ ]: test_prompt = prompt_sets["short_only.jsonl"][0]
print("prompt tokens (rough chars):", len(test_prompt))

out = one_request(test_prompt, max_tokens=32)
print(out)

prompt tokens (rough chars): 49
{'latency_s': 0.8831302409998898, 'prompt_tokens': 39, 'completion_tokens': 32, 'total_tokens': 71, 'status_code': 200}
```

Client-Side Load Generator (ThreadPool) + Latency/RPS Summary

```
In [ ]: def run_load(prompts, concurrency=8, total_requests=200, max_tokens=128,
    rng = random.Random(123)
    picked = [rng.choice(prompts) for _ in range(total_requests)]

    results = []
    t_start = time.perf_counter()

    def task(i):
        p = picked[i]
        try:
            res = one_request(p, max_tokens=max_tokens)
            res["ok"] = True
        except Exception as e:
            res = {"ok": False, "error": str(e), "latency_s": None,
                   "prompt_tokens": None, "completion_tokens": None, "tot
    return res

    ok_ct = 0
    fail_ct = 0

    with cf.ThreadPoolExecutor(max_workers=concurrency) as ex:
        futs = [ex.submit(task, i) for i in range(total_requests)]

        pbar = tqdm(total=total_requests, desc=desc)
        for f in cf.as_completed(futs):
            r = f.result()
            results.append(r)
            if r.get("ok"):
                ok_ct += 1
            else:
                fail_ct += 1
            pbar.update(1)
            pbar.set_postfix(ok=ok_ct, fail=fail_ct)
```

```

    pbar.close()

    t_end = time.perf_counter()

    df = pd.DataFrame(results)
    ok_df = df[df["ok"] == True].copy()

    summary = {
        "Concurrency": concurrency,
        "Total Requests": total_requests,
        "Ok": int(ok_df.shape[0]),
        "Fail": int(df.shape[0] - ok_df.shape[0]),
        "Wall Time (s)": t_end - t_start,
        "RPS": ok_df.shape[0] / (t_end - t_start) if (t_end - t_start) > 0 else None,
        "P50 (s)": float(ok_df["Latency (s)"].quantile(0.50)) if len(ok_df) > 0 else None,
        "P90 (s)": float(ok_df["Latency (s)"].quantile(0.90)) if len(ok_df) > 0 else None,
        "P99 (s)": float(ok_df["Latency (s)"].quantile(0.99)) if len(ok_df) > 0 else None,
        "Mean (s)": float(ok_df["Latency (s)"].mean()) if len(ok_df) > 0 else None
    }
    return df, summary

```

Reproducing Context-Length Failures Under Load (Max Token Limit)

```

In [ ]: workloads = [
    ("short_only", "short_only.jsonl"),
    ("mix_90_10", "mix_90_10.jsonl"),
    ("mix_70_30", "mix_70_30.jsonl"),
]

all_summaries = []
all_dfs = {}

for name, fn in tqdm(workloads, desc="Workloads"):
    print("\nRunning:", name)
    df, summ = run_load(
        prompt_sets[fn],
        concurrency=8,
        total_requests=200,
        max_tokens=128,
        desc=f"{name} (c=8)"
    )
    all_dfs[name] = df
    all_summaries.append({"workload": name, **summ})
    print(summ)

summary_df = pd.DataFrame(all_summaries)
summary_df

```

Workloads: 0% | 0/3 [00:00<?, ?it/s]
 Running: short_only
 short_only (c=8): 0% | 0/200 [00:00<?, ?it/s]
 {'Concurrency': 8, 'Total Requests': 200, 'Ok': 200, 'Fail': 0, 'Wall Time (s)': 79.18587238500186, 'RPS': 2.5257030575807717, 'P50 (s)': 3.553790949499671, 'P90 (s)': 3.5761803588007752, 'P99 (s)': 3.5867652782488584, 'Mean (s)': 3.104981383839822}

Running: mix_90_10

```
mix_90_10 (c=8): 0% | 0/200 [00:00<?, ?it/s]
{'concurrency': 8, 'total_requests': 200, 'ok': 198, 'fail': 2, 'wall_time_s': 86.622785784999, 'rps': 2.2857727121757945, 'p50_s': 3.6156531919987174, 'p90_s': 4.1372710210980586, 'p99_s': 4.528918918627678, 'mean_s': 3.43375558661609}
```

```
Running: mix_70_30
mix_70_30 (c=8): 0% | 0/200 [00:00<?, ?it/s]
{'concurrency': 8, 'total_requests': 200, 'ok': 195, 'fail': 5, 'wall_time_s': 99.59230680399924, 'rps': 1.9579825616828626, 'p50_s': 4.101179446999595, 'p90_s': 4.894588489399757, 'p99_s': 5.37605518930228, 'mean_s': 4.054637723287175}
```

Out[]:	workload	concurrency	total_requests	ok	fail	wall_time_s	rps	p5
	0 short_only	8	200	200	0	79.185872	2.525703	3.5553
	1 mix_90_10	8	200	198	2	86.622786	2.285773	3.615
	2 mix_70_30	8	200	195	5	99.592307	1.957983	4.101

```
In [ ]: for name, df in all_dfs.items():
    print("\n====", name, "====")
    print("rows:", len(df))
    print("cols:", list(df.columns))

    if "ok" in df.columns:
        fails = df[df["ok"] == False]
    elif "status" in df.columns:
        fails = df[df["status"] != 200]
    else:
        print("I don't see ok/status columns - tell me df.columns and I'll continue")

    print("fail count:", len(fails))
    if len(fails):
        cols_to_show = [c for c in ["status", "error", "latency_s", "prompt_tokens"]]
        display(fails[cols_to_show].head(10))
        if "error" in fails.columns:
            print("\nTop error reasons:")
            display(fails["error"].value_counts().head(10))
```

```
==== short_only ====
rows: 200
cols: ['latency_s', 'prompt_tokens', 'completion_tokens', 'total_tokens', 'status_code', 'ok']
fail count: 0

==== mix_90_10 ====
rows: 200
cols: ['latency_s', 'prompt_tokens', 'completion_tokens', 'total_tokens', 'status_code', 'ok', 'error']
fail count: 2
```

		error	latency_s	prompt_tokens
159	HTTP 400: {"error": {"message": "'max_tokens' or...}}	NaN	NaN	NaN
168	HTTP 400: {"error": {"message": "'max_tokens' or...}}	NaN	NaN	NaN

Top error reasons:

	count
	error
HTTP 400: {"error":{"message":"'max_tokens' or 'max_completion_tokens' is too large: 128. This model's maximum context length is 4096 tokens and your request has 4030 input tokens (128 > 4096 - 4030). None","type":"BadRequestError","param":null,"code":400}}	1
HTTP 400: {"error":{"message":"'max_tokens' or 'max_completion_tokens' is too large: 128. This model's maximum context length is 4096 tokens and your request has 4040 input tokens (128 > 4096 - 4040). None","type":"BadRequestError","param":null,"code":400}}	1

dtype: int64
 === mix_70_30 ===
 rows: 200
 cols: ['latency_s', 'prompt_tokens', 'completion_tokens', 'total_tokens', 'status_code', 'ok', 'error']
 fail count: 5

		error	latency_s	prompt_tokens
85	HTTP 400: {"error":{"message":"'max_tokens' or...	NaN	NaN	NaN
90	HTTP 400: {"error":{"message":"'max_tokens' or...	NaN	NaN	NaN
118	HTTP 400: {"error":{"message":"'max_tokens' or...	NaN	NaN	NaN
178	HTTP 400: {"error":{"message":"'max_tokens' or...	NaN	NaN	NaN
189	HTTP 400: {"error":{"message":"'max_tokens' or...	NaN	NaN	NaN

Top error reasons:

	count
	error
HTTP 400: {"error":{"message":"'max_tokens' or 'max_completion_tokens' is too large: 128. This model's maximum context length is 4096 tokens and your request has 4040 input tokens (128 > 4096 - 4040). None","type":"BadRequestError","param":null,"code":400}}	2
HTTP 400: {"error":{"message":"'max_tokens' or 'max_completion_tokens' is too large: 128. This model's maximum context length is 4096 tokens and your request has 3984 input tokens (128 > 4096 - 3984). None","type":"BadRequestError","param":null,"code":400}}	1
HTTP 400: {"error":{"message":"'max_tokens' or 'max_completion_tokens' is too large: 128. This model's maximum context length is 4096 tokens and your request has 4015 input tokens (128 > 4096 - 4015). None","type":"BadRequestError","param":null,"code":400}}	1
HTTP 400: {"error":{"message":"'max_tokens' or 'max_completion_tokens' is too large: 128. This model's maximum context length is 4096 tokens and your request has 4060 input tokens (128 > 4096 - 4060). None","type":"BadRequestError","param":null,"code":400}}	1

dtype: int64

Conclusion

In the mixed workloads (mix_90_10 and mix_70_30), we start seeing HTTP 400 failures even though the same code works fine for short-only prompts. The error message shows why: the model has a fixed maximum context length of 4096 tokens, and some of our long prompts already consume ~3980–4060 input tokens.

When we blindly request a fixed max_tokens=128 for completion, the request becomes impossible because the remaining context budget (4096 - input_tokens) is smaller than 128. vLLM therefore rejects those requests with a BadRequest error (e.g., "128 > 4096 – 4040").

This confirms a real serving issue under mixed prompt lengths: static max_tokens can cause hard failures whenever inputs approach the context window, motivating the next experiment where we dynamically cap max_tokens based on the available token budget per request.

```
In [ ]: import matplotlib.pyplot as plt
import numpy as np

labels = summary_df["workload"].tolist()

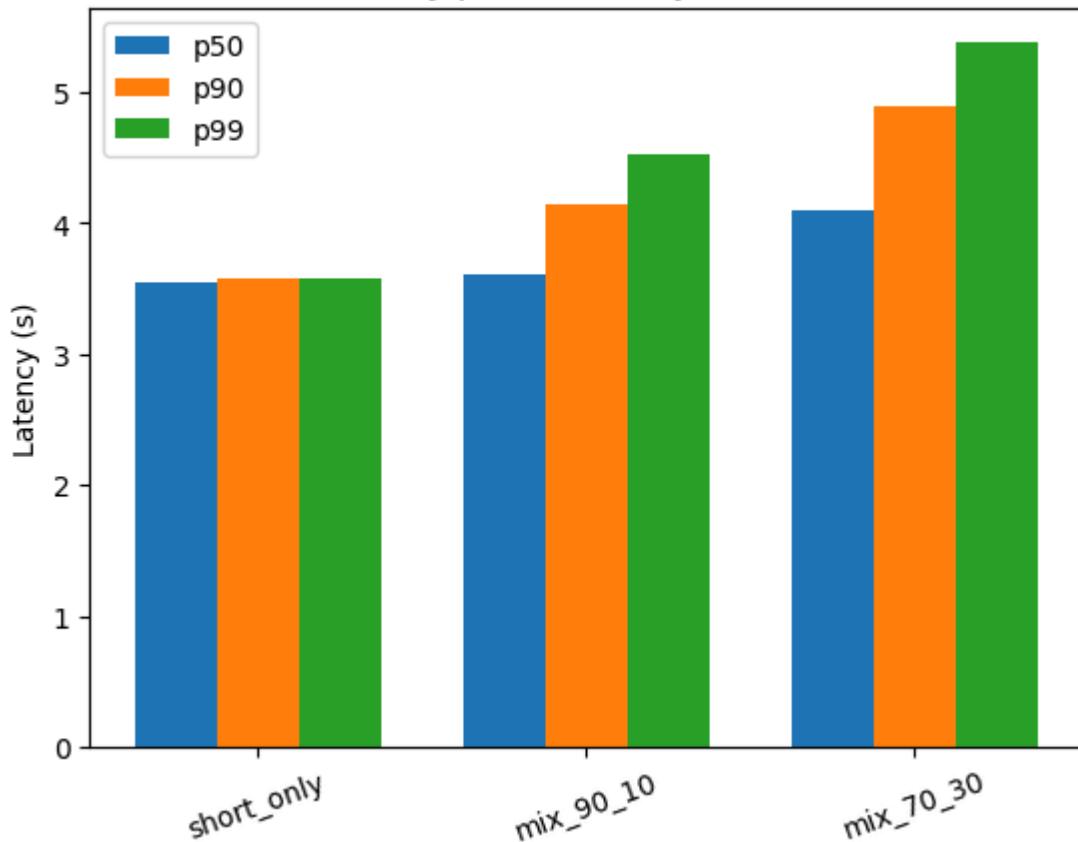
# p50/p90/p99
p50 = summary_df["p50_s"].values
p90 = summary_df["p90_s"].values
p99 = summary_df["p99_s"].values

x = np.arange(len(labels))
w = 0.25

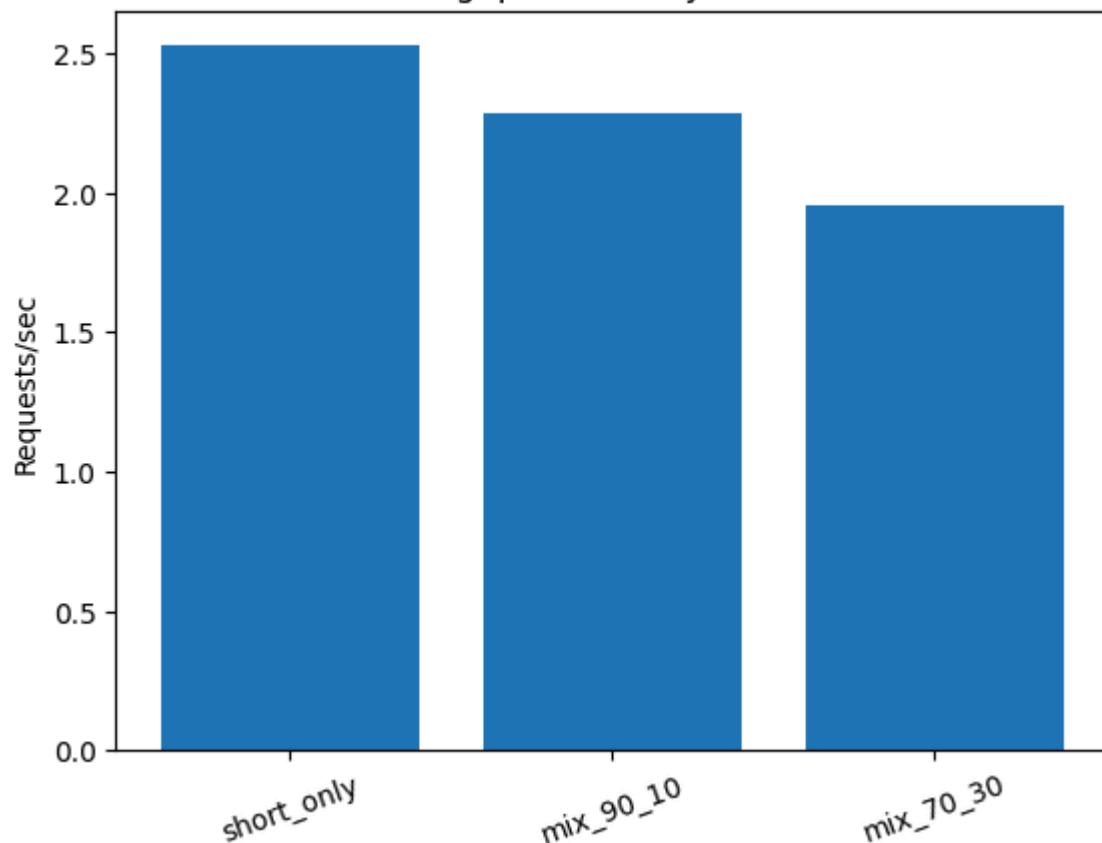
plt.figure()
plt.bar(x - w, p50, width=w, label="p50")
plt.bar(x, p90, width=w, label="p90")
plt.bar(x + w, p99, width=w, label="p99")
plt.xticks(x, labels, rotation=20)
plt.ylabel("Latency (s)")
plt.title("Latency percentiles by workload")
plt.legend()
plt.show()

# RPS
plt.figure()
plt.bar(labels, summary_df["rps"].values)
plt.ylabel("Requests/sec")
plt.title("Throughput (RPS) by workload")
plt.xticks(rotation=20)
plt.show()
```

Latency percentiles by workload



Throughput (RPS) by workload



Plot Interpretation

These two plots summarize the impact of prompt-length mix on both latency tail and throughput at a fixed concurrency ($c=8$, 200 total requests each).

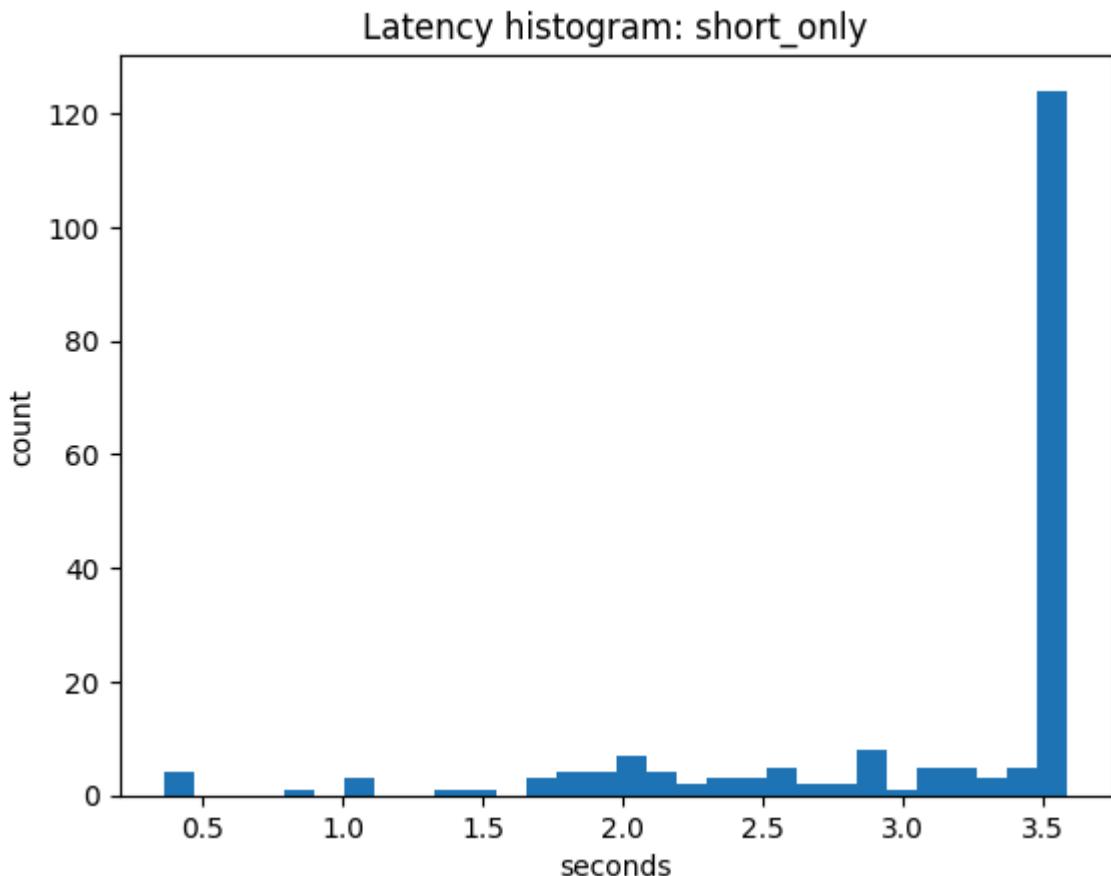
As we move from short_only \rightarrow mix_90_10 \rightarrow mix_70_30, the proportion of long prompts increases, which pushes up p90/p99 latency and reduces overall RPS. This is the classic “mixed-length serving” effect.

Long prompts spend much more time in prefill and consume more KV-cache and batching budget, so they slow down batch completion and increase queueing for everyone, which shows up most clearly in tail latency.

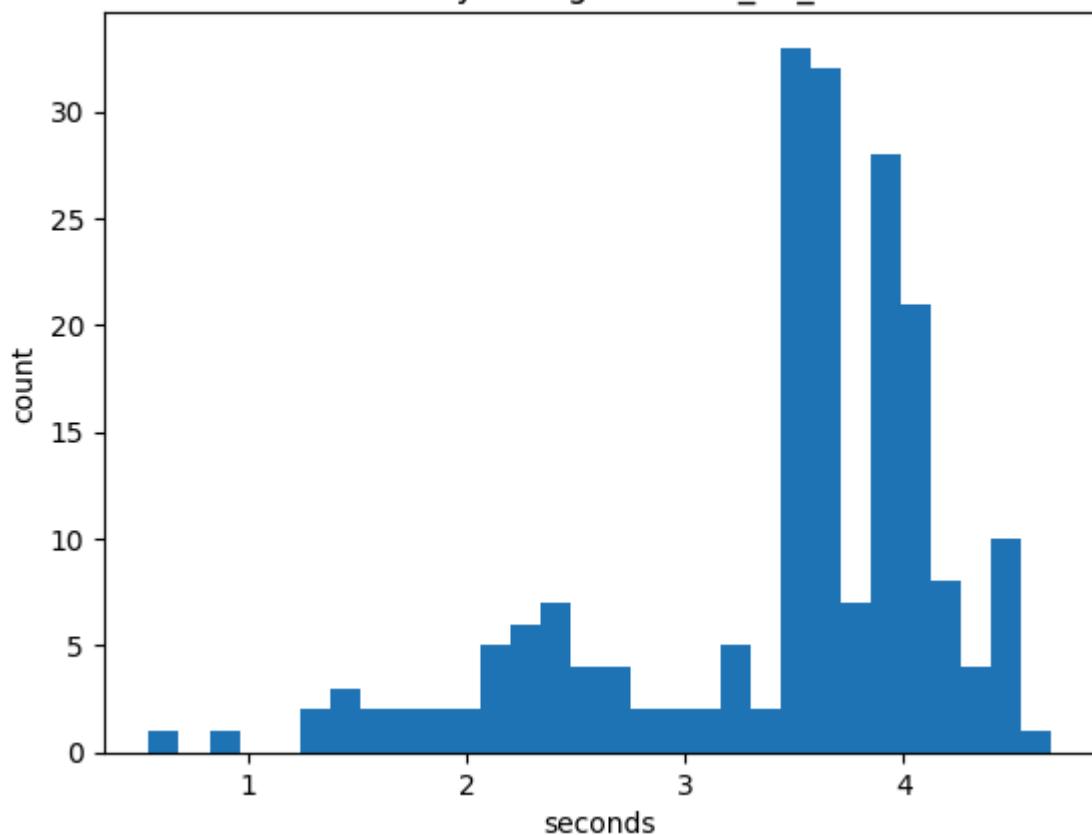
The throughput chart reinforces this: more long requests means fewer requests completed per second. Also note that the latency percentiles here are computed over successful (ok) requests.

The few failed requests from long prompts are not included in the percentile bars, and those failures are exactly what we address next with the max_tokens safety and dynamic max_tokens fix.

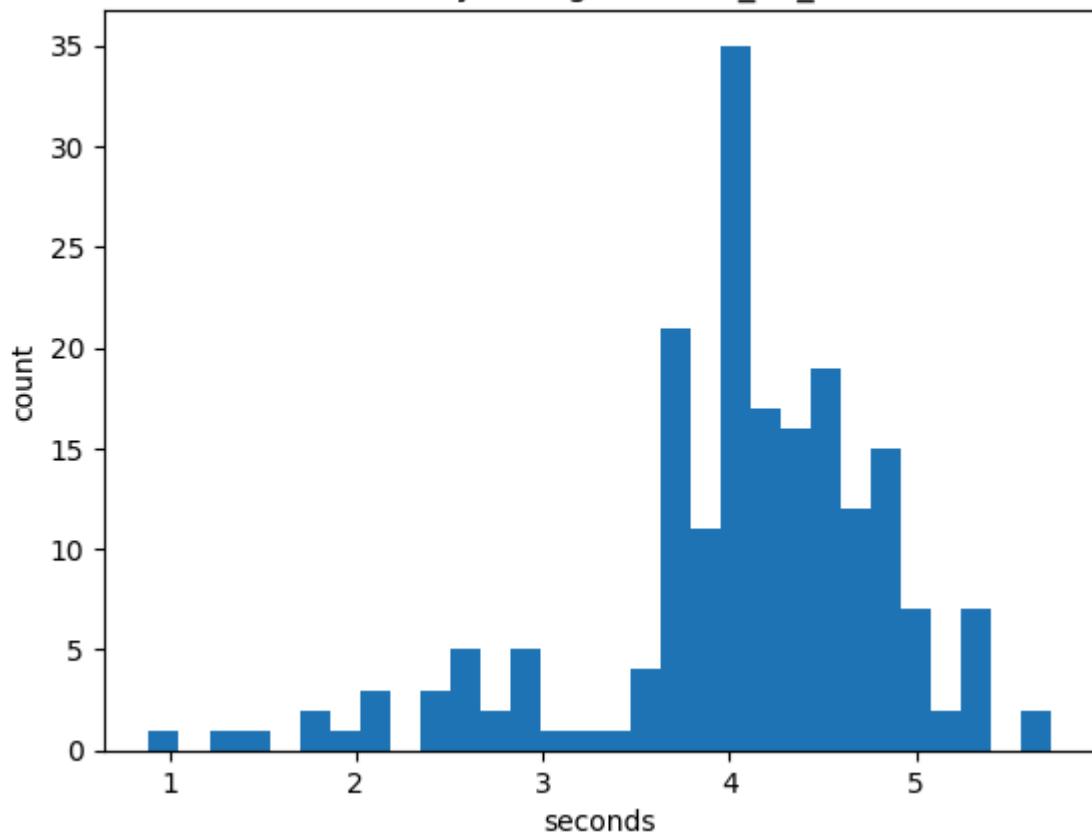
```
In [ ]: for name, df in all_dfs.items():
    if "latency_s" in df.columns:
        vals = df[df["latency_s"].notna()]["latency_s"].values
        plt.figure()
        plt.hist(vals, bins=30)
        plt.title(f"Latency histogram: {name}")
        plt.xlabel("seconds")
        plt.ylabel("count")
        plt.show()
```



Latency histogram: mix_90_10



Latency histogram: mix_70_30



Plot Interpretation

These histograms show the full latency distribution for each workload, not just the percentile summaries. The short_only workload is tightly clustered, which indicates stable, predictable latency when all prompts are small.

As we introduce longer prompts in mix_90_10 and mix_70_30, the distribution shifts to the right and spreads out, meaning higher average latency and greater variability. You can also see a heavier "right tail" in the mixed workloads, which is exactly what drives the p90/p99 increases we saw earlier.

This supports the core point of the experiment: mixed prompt lengths increase tail latency and make serving behavior less predictable.

In the next step we addressing context-budget failures.

Experiment 1.1: Eliminating max_token_limit Failures with Context-Budgeting

Why this matters ??

In real LLM services, requests do not arrive with uniform prompt sizes. You typically see a mix of short questions, medium chat turns, and occasional very long inputs (documents, logs, transcripts, RAG context, tool traces). If the system sometimes returns hard failures (for example, HTTP 400 context overflow), the service becomes unreliable, breaks downstream pipelines (agents, batch jobs, retries), and makes it harder to meet latency and availability targets.

A robust serving layer should enforce a clear context budget policy that keeps every request within the model's context window, while degrading gracefully when the input is too large.

Why we cannot "just cap prompt length"??

Simply limiting the prompt size is not a complete solution:

- **Capability loss:** Long inputs are often the whole point of the application (summarization, document Q&A, log diagnosis).
- **Prompt + completion share the same budget:** Even a "legal-sized" prompt can overflow if `max_tokens` is too large for the remaining space.
- **Naive truncation hurts quality:** Blindly cutting the prompt can drop critical context and silently degrade answer quality.
- **Chat formatting overhead:** System messages and chat templates add tokens that a simple character cap does not account for.

What we implement in Experiment 1.1

In this experiment, we implement a **context-budgeting policy** that guarantees requests fit the model's context window:

- First, shrink the output (`max_tokens`) to fit the remaining budget.
- If there still is not enough space for a minimum output, trim the input (head+tail) to create room.
- Log metadata (estimated prompt tokens, chosen `max_tokens`, whether trimming occurred) so the behavior is measurable and reproducible.

```
In [ ]: # clean out previous instances of vLLM if you want to start a new server
```

```
!pkill -9 -f "vllm.entrypoints.openai.api_server" || true
!pkill -9 -f "EngineCore" || true
!pkill -9 -f "vllm" || true
!sleep 2
!nvidia-smi
```

```
^C
^C
^C
Sun Jan 11 22:01:00 2026
+-----
+-----+
| NVIDIA-SMI 550.54.15           Driver Version: 550.54.15      CUDA Ve
rsion: 12.4          |
+-----+
| GPU  Name                  Persistence-M | Bus-Id          Disp.A | Volat
ile Uncorr. ECC |
| Fan  Temp     Perf          Pwr:Usage/Cap |          Memory-Usage | GPU-U
til  Compute M. |
|          |                               |          |
MIG M.  |
+-----+
| 0  NVIDIA L4                Off | 00000000:00:03.0 Off |
| 0 |
| N/A  54C     P8          18W / 72W | 3MiB / 23034MiB |
| 0%    Default   |
|          |                               |
N/A  |
+-----+
+-----+
| Processes:
|
| GPU  GI  CI          PID  Type  Process name
GPU Memory |
|          ID  ID
Usage       |
+-----+
| No running processes found
|
+-----+
```

```
In [ ]: !lsof -i :8000 || true
COMMAND PID USER FD TYPE DEVICE SIZE/OFF NODE NAME
python3 1144 root 79u IPv4 606432      0t0  TCP localhost:49490->localhost:8000 (CLOSE_WAIT)

In [ ]: MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"
PORT = 8000

MAX_MODEL_LEN = 4096
GPU_UTIL = 0.70
DTYPE = "half"

cmd = f"""
nohup python -m vllm.entrypoints.openai.api_server \
--model {MODEL_ID} \
--host 127.0.0.1 \
--port {PORT} \
--max-model-len {MAX_MODEL_LEN} \
--gpu-memory-utilization {GPU_UTIL} \
--dtype {DTYPE} \
--disable-log-stats \
> vllm_server.log 2>&1 &
"""

print("Launching vLLM server...")
!bash -lc "$cmd"
!tail -n 30 vllm_server.log
```

Launching vLLM server...

Function to check the successful startup of the server.

```
In [ ]: BASE = "http://127.0.0.1:8000"

def wait_for_server(timeout_s=600):
    t0 = time.time()
    while time.time() - t0 < timeout_s:
        try:
            r = requests.get(f"{BASE}/v1/models", timeout=2)
            if r.status_code == 200:
                return True, r.json()
        except Exception:
            pass
        time.sleep(2)
    return False, None

ok, payload = wait_for_server()
print("Ready:", ok)
if ok:
    print("Models:", [m["id"] for m in payload.get("data", [])])
else:
    print("Not ready yet. Check logs:")
    !tail -n 60 vllm_server.log
```

Ready: True
Models: ['Qwen/Qwen2.5-3B-Instruct']

Dynamic Output Budgeting and Prompt Trimming to Avoid Context Overflow

```
In [ ]: SAFETY_MARGIN = 16          # keep small headroom to avoid borderline fail
DESIRED_MAX_TOKENS = 128        # same as Experiment 1
MIN_OUTPUT_TOKENS = 32          # minimum output you want to guarantee

def count_prompt_tokens(messages):
    """
    Counts tokens for the full chat-formatted prompt that will be sent to
    Owen tokenizer supports apply_chat_template().
    """
    ids = tok.apply_chat_template(
        messages,
        tokenize=True,
        add_generation_prompt=True,
        return_tensors=None
    )
    return len(ids)

def head_tail_trim_text_by_tokens(text, target_tokens, head_frac=0.35):
    """
    Keep some head + some tail tokens, drop middle.
    """
    ids = tok.encode(text, add_special_tokens=False)
    if len(ids) <= target_tokens:
        return text, False

    head_tokens = int(target_tokens * head_frac)
    tail_tokens = target_tokens - head_tokens

    trimmed_ids = ids[:head_tokens] + ids[-tail_tokens:]
    return tok.decode(trimmed_ids), True

def build_payload_with_budget(prompt_text,
                               desired_max_tokens=DESIRED_MAX_TOKENS,
                               min_output_tokens=MIN_OUTPUT_TOKENS,
                               safety_margin=SAFETY_MARGIN,
                               temperature=0.0):
    """
    Returns (payload, meta).
    - Ensures no context overflow.
    - If needed, trims user prompt (head+tail) to make room for min output
    """
    messages = [{"role": "user", "content": prompt_text}]
    ptoks = count_prompt_tokens(messages)

    # First attempt: just shrink output tokens if needed
    max_out = min(desired_max_tokens, MAX_MODEL_LEN - safety_margin - ptoks)

    trimmed = False
    original_ptoks = ptoks
    original_prompt = prompt_text

    # If not enough room even for min output -> trim input to create room
    if max_out < min_output_tokens:
```

```

target_prompt_budget = MAX_MODEL_LEN - safety_margin - min_output
if target_prompt_budget <= 0:
    raise RuntimeError("Model context too small after safety margin")

# We need to reduce prompt tokens from ptoks down to target_prompt_budget.
# We'll trim only the USER CONTENT tokens.
#
# Use a simple search on how many user-content tokens we can keep
# (Because chat template adds overhead tokens.)
lo, hi = 0, len(tok.encode(prompt_text, add_special_tokens=False))
best_text = ""
best_ptoks = None

while lo <= hi:
    mid = (lo + hi) // 2
    candidate_text, _ = head_tail_trim_text_by_tokens(prompt_text)
    candidate_msgs = [{"role": "user", "content": candidate_text}]
    candidate_ptoks = count_prompt_tokens(candidate_msgs)

    if candidate_ptoks <= target_prompt_budget:
        best_text = candidate_text
        best_ptoks = candidate_ptoks
        lo = mid + 1
    else:
        hi = mid - 1

if best_ptoks is None:
    raise RuntimeError(
        f"Prompt too long to fit even after trimming. "
        f"prompt_tokens={ptoks}, budget={target_prompt_budget}"
    )

prompt_text = best_text
trimmed = True
messages = [{"role": "user", "content": prompt_text}]
ptoks = best_ptoks

max_out = min(desired_max_tokens, MAX_MODEL_LEN - safety_margin - min_output)

# Final guard
if max_out <= 0:
    raise RuntimeError(
        f"No room for output tokens. prompt_tokens={ptoks}, "
        f"max_len={MAX_MODEL_LEN}, safety_margin={safety_margin}"
    )

payload = {
    "model": MODEL_ID,
    "messages": messages,
    "max_tokens": int(max_out),
    "temperature": temperature,
}

meta = {
    "prompt_tokens_est": ptoks,
    "max_tokens_used": int(max_out),
    "trimmed": trimmed,
    "prompt_tokens_est_original": original_ptoks,
    "original_prompt_len_chars": len(original_prompt),
    "final_prompt_len_chars": len(prompt_text),
}

```

```

    }
    return payload, meta

```

Request Execution with Dynamic Token Budgeting

```
In [ ]: def one_request(prompt, desired_max_tokens=128, temperature=0.0, timeout=t0 = time.perf_counter():

    payload, meta = build_payload_with_budget(
        prompt_text=prompt,
        desired_max_tokens=desired_max_tokens,
        min_output_tokens=MIN_OUTPUT_TOKENS,
        safety_margin=SAFETY_MARGIN,
        temperature=temperature,
    )

    r = requests.post(f"{BASE}/v1/chat/completions", json=payload, timeout=timeout)
    latency = time.perf_counter() - t0

    if r.status_code != 200:
        raise RuntimeError(f"HTTP {r.status_code}: {r.text[:800]}")

    j = r.json()
    usage = j.get("usage", {})

    return {
        "latency_s": latency,
        "prompt_tokens": usage.get("prompt_tokens"),
        "completion_tokens": usage.get("completion_tokens"),
        "total_tokens": usage.get("total_tokens"),
        "status_code": r.status_code,
        # extra debugging info (useful for your writeup)
        **meta
    }
```

```
In [ ]: # smoke test of one_request function

test_prompt = "Explain in 2 sentences what chunked prefill does in vLLM."
out = one_request(test_prompt, desired_max_tokens=128, temperature=0.0, timeout=10)
```

```
Out[ ]: {'latency_s': 1.3368042790007166,
         'prompt_tokens': 46,
         'completion_tokens': 49,
         'total_tokens': 95,
         'status_code': 200,
         'prompt_tokens_est': 46,
         'max_tokens_used': 128,
         'trimmed': False,
         'prompt_tokens_est_original': 46,
         'original_prompt_len_chars': 57,
         'final_prompt_len_chars': 57}
```

Load Test Harness Using Budget-Aware Requests (Dynamic max_tokens + Optional Prompt Trimming)

```
In [ ]: def run_load(prompts, concurrency=8, total_requests=200, desired_max_toke
rng = random.Random(123)
picked = [rng.choice(prompts) for _ in range(total_requests)]

results = []
t_start = time.perf_counter()

def task(i):
    p = picked[i]
    try:
        res = one_request(p, desired_max_tokens=desired_max_tokens, t
        res["ok"] = True
    except Exception as e:
        res = {"ok": False, "error": str(e), "latency_s": None,
               "prompt_tokens": None, "completion_tokens": None, "tot
    return res

ok_ct = 0
fail_ct = 0

with cf.ThreadPoolExecutor(max_workers=concurrency) as ex:
    futs = [ex.submit(task, i) for i in range(total_requests)]
    pbar = tqdm(total=total_requests, desc=desc)
    for f in cf.as_completed(futs):
        r = f.result()
        results.append(r)
        if r.get("ok"):
            ok_ct += 1
        else:
            fail_ct += 1
        pbar.update(1)
        pbar.set_postfix(ok=ok_ct, fail=fail_ct)
    pbar.close()

t_end = time.perf_counter()

df = pd.DataFrame(results)
ok_df = df[df["ok"] == True].copy()

summary = {
    "concurrency": concurrency,
    "total_requests": total_requests,
    "ok": int(ok_df.shape[0]),
    "fail": int(df.shape[0] - ok_df.shape[0]),
    "wall_time_s": t_end - t_start,
    "rps": ok_df.shape[0] / (t_end - t_start) if (t_end - t_start) >
    "p50_s": float(ok_df["latency_s"].quantile(0.50)) if len(ok_df) e
    "p90_s": float(ok_df["latency_s"].quantile(0.90)) if len(ok_df) e
    "p99_s": float(ok_df["latency_s"].quantile(0.99)) if len(ok_df) e
    "mean_s": float(ok_df["latency_s"].mean()) if len(ok_df) else Non
}
return df, summary
```

Run Safe (Budget-Aware) Load Tests Across Workloads

```
In [ ]: workloads = [
    ("short_only", "short_only.jsonl"),
    ("mix_90_10", "mix_90_10.jsonl"),
    ("mix_70_30", "mix_70_30.jsonl"),
]

all_summaries = []
all_dfs = {}

for name, fn in tqdm(workloads, desc="Workloads (safe)"):
    print("\nRunning:", name)
    df, summ = run_load(
        prompt_sets[fn],
        concurrency=8,
        total_requests=200,
        desired_max_tokens=128,
        desc=f"{name} (c=8, safe)"
    )
    all_dfs[name] = df
    all_summaries.append({"workload": name, **summ})
    print(summ)

summary_df_safe = pd.DataFrame(all_summaries)
summary_df_safe
```

Workloads (safe): 0% | 0/3 [00:00<?, ?it/s]
Running: short_only
short_only (c=8, safe): 0% | 0/200 [00:00<?, ?it/s]
{'concurrency': 8, 'total_requests': 200, 'ok': 200, 'fail': 0, 'wall_time_s': 79.25956869500078, 'rps': 2.523354634563067, 'p50_s': 3.556057800500639, 'p90_s': 3.5734452239999883, 'p99_s': 3.58166045346994, 'mean_s': 3.106944176310044}

Running: mix_90_10
mix_90_10 (c=8, safe): 0% | 0/200 [00:00<?, ?it/s]
{'concurrency': 8, 'total_requests': 200, 'ok': 200, 'fail': 0, 'wall_time_s': 88.17387503600003, 'rps': 2.268245553667037, 'p50_s': 3.640730647500277, 'p90_s': 4.24218545370004, 'p99_s': 4.5321530194697335, 'mean_s': 3.4690503159849686}

Running: mix_70_30
mix_70_30 (c=8, safe): 0% | 0/200 [00:00<?, ?it/s]
{'concurrency': 8, 'total_requests': 200, 'ok': 200, 'fail': 0, 'wall_time_s': 102.89451057299993, 'rps': 1.943738289693377, 'p50_s': 4.175408424499437, 'p90_s': 4.893263480699716, 'p99_s': 5.3932446032813255, 'mean_s': 4.073671482370019}

Out []:	workload	concurrency	total_requests	ok	fail	wall_time_s	rps	p5
0	short_only	8	200	200	0	79.259569	2.523355	3.556
1	mix_90_10	8	200	200	0	88.173875	2.268246	3.640
2	mix_70_30	8	200	200	0	102.894511	1.943738	4.175

Conclusion

In Experiment 1.1 (the budget-aware request builder), the key improvement is reliability: all three workloads now complete with 0 failures (200/200 OK), including mixes that previously hit HTTP 400 context overflow when long prompts left insufficient room for the requested completion (max_tokens).

This confirms the fix works as intended, because each request now dynamically adjusts max_tokens to fit the remaining context window and trims the prompt only when necessary to guarantee a minimum output budget.

Performance trends remain consistent with Experiment 1: as the proportion of long prompts increases (short_only -> mix_90_10 -> mix_70_30), latency rises and throughput (RPS) drops due to heavier prefill and decoding work, but crucially this happens without any request-level hard failures, making the system robust under mixed prompt lengths.

Experiment 2: Diagnosing Tail Latency with a “Bad” vLLM Batching Config

```
In [ ]: !pkill -f "vllm.entrypoints.openai.api_server" || true
!pkill -f "VLLM::EngineCore" || true
!pkill -f "EngineCore" || true
time.sleep(2)

!rm -f vllm_bad.log
print("killed + log removed")
```

```
^C
^C
^C
killed + log removed
```

Creating a Bad Config for LLM Server

To replicate a “**chunked prefill effectively OFF**” situation

Why this is a “bad” config: We set `--max-num-batched-tokens 8192` (very large) while leaving `--max-num-seqs` unconstrained (defaults can be high). This encourages vLLM to pack big prefills into a batch, which can let long prompts dominate compute and delay short requests.

Chunked prefetch caveat: Newer vLLM versions have chunked prefetch **enabled by default**, but its *practical* benefit depends on admission and batching knobs. If `max-num-batched-tokens` is too permissive (and/or `max-num-seqs` allows over-admission), long prefills can still monopolize batch capacity, effectively **diminishing chunked prefetch's ability to interleave work and protect short requests.**

Head-of-line blocking risk: With mixed workloads (50/50 short and long), long-prefill requests can occupy most of the batch token budget, causing short prompts to wait behind long prefills, inflating **p90/p99** latency.

```
In [ ]: !nohup python3 -m vllm.entrypoints.openai.api_server \
--model Qwen/Qwen2.5-3B-Instruct \
--host 127.0.0.1 --port 8000 \
--dtype half \
--max-model-len 4096 \
--gpu-memory-utilization 0.70 \
--disable-log-stats \
--enforce-eager \
--max-num-batched-tokens 8192 \
--long-prefill-token-threshold 1024 \
> vllm_bad.log 2>&1 </dev/null &
```

```
In [ ]: BASE = "http://127.0.0.1:8000"

def wait_for_server(timeout_s=600, poll_s=2):
    t0 = time.time()
    last_err = None
    while time.time() - t0 < timeout_s:
        try:
            r = requests.get(f"{BASE}/v1/models", timeout=2)
            if r.status_code == 200:
                return True, r.json()
            last_err = f"HTTP {r.status_code}: {r.text[:200]}"
        except Exception as e:
            last_err = f"{type(e).__name__}: {e}"
        time.sleep(poll_s)
    return False, last_err

ok, models_or_err = wait_for_server()
print("Ready:", ok)

if ok:
    print("Models:", [m.get("id") for m in models_or_err.get("data", [])])
else:
    print("Not ready yet. Last error:", models_or_err)
    print("Check logs:")
    !tail -n 80 vllm_bad.log
```

Ready: True
 Models: ['Qwen/Qwen2.5-3B-Instruct']

```
In [ ]: !tail -n 60 vllm_bad.log
```

```
(EngineCore_DP0 pid=28652)
Loading safetensors checkpoint shards: 0% Completed | 0/2 [00:00<?, ?it/
s]
(EngineCore_DP0 pid=28652)
Loading safetensors checkpoint shards: 50% Completed | 1/2 [00:01<00:01,
1.06s/it]
(EngineCore_DP0 pid=28652)
Loading safetensors checkpoint shards: 100% Completed | 2/2 [00:02<00:00,
1.50s/it]
(EngineCore_DP0 pid=28652)
Loading safetensors checkpoint shards: 100% Completed | 2/2 [00:02<00:00,
1.44s/it]
(EngineCore_DP0 pid=28652)
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:05 [default_loader.py:308] Loa
ding weights took 2.95 seconds
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:05 [gpu_model_runner.py:3659]
Model loading took 5.7916 GiB memory and 4.417226 seconds
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:08 [gpu_worker.py:375] Availab
le KV cache memory: 8.27 GiB
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:08 [kv_cache_utils.py:1291] GP
U KV cache size: 240,768 tokens
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:08 [kv_cache_utils.py:1296] Ma
ximum concurrency for 4,096 tokens per request: 58.78x
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:08 [core.py:259] init engine
(profile, create kv cache, warmup model) took 2.94 seconds
(EngineCore_DP0 pid=28652) WARNING 01-12 03:47:11 [vllm.py:629] Inductor c
ompilation was disabled by user settings, Optimizations settings that are o
nly active duringInductor compilation will be ignored.
(EngineCore_DP0 pid=28652) INFO 01-12 03:47:11 [vllm.py:722] Cudagraph is
disabled under eager mode
(APIServer pid=28521) INFO 01-12 03:47:11 [api_server.py:1099] Supported t
asks: ['generate']
(APIServer pid=28521) WARNING 01-12 03:47:11 [model.py:1487] Default sampl
ing parameters have been overridden by the model's Hugging Face generation
config recommended from the model creator. If this is not intended, please
relaunch vLLM instance with `--generation-config vllm`.
(APIServer pid=28521) INFO 01-12 03:47:11 [serving_responses.py:201] Using
default chat sampling params from model: {'repetition_penalty': 1.05, 'tem
perature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=28521) INFO 01-12 03:47:12 [serving_chat.py:137] Using defa
ult chat sampling params from model: {'repetition_penalty': 1.05, 'tempera
ture': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=28521) INFO 01-12 03:47:12 [serving_completion.py:77] Using
default completion sampling params from model: {'repetition_penalty': 1.0
5, 'temperature': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=28521) INFO 01-12 03:47:12 [serving_chat.py:137] Using defa
ult chat sampling params from model: {'repetition_penalty': 1.05, 'tempera
ture': 0.7, 'top_k': 20, 'top_p': 0.8}
(APIServer pid=28521) INFO 01-12 03:47:12 [api_server.py:1425] Starting vL
LM API server 0 on http://127.0.0.1:8000
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:38] Available route
s are:
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /openap
i.json, Methods: GET, HEAD
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /docs, M
ethods: GET, HEAD
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /docs/oa
uth2-redirect, Methods: GET, HEAD
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /redoc,
Methods: GET, HEAD
```

```
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /scale_elastic_ep, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /is_scaling_elastic_ep, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /tokenize, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /detokenize, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /inference/v1/generate, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /pause, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /resume, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /is_paused, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /metrics, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /health, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /load, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/models, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /version, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/responses, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/responses/{response_id}, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/responses/{response_id}/cancel, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/messages, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/chat/completions, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/competitions, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/audio/transcriptions, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/audio/translations, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /ping, Methods: GET
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /ping, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /invocations, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /classify, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /embeddings, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /score, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /score, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /rerank, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v1/rerank, Methods: POST
```

```
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /v2/rera
nk, Methods: POST
(APIServer pid=28521) INFO 01-12 03:47:12 [launcher.py:46] Route: /poolin
g, Methods: POST
(APIServer pid=28521) INFO:      Started server process [28521]
(APIServer pid=28521) INFO:      Waiting for application startup.
(APIServer pid=28521) INFO:      Application startup complete.
```

Load Tokenizer

```
In [ ]: import requests
from transformers import AutoTokenizer

BASE = "http://127.0.0.1:8000"
MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

tok = AutoTokenizer.from_pretrained(MODEL_ID, use_fast=True)

models = requests.get(f"{BASE}/v1/models").json()
model_info = next(m for m in models["data"] if m["id"] == MODEL_ID)
MAX_MODEL_LEN = int(model_info.get("max_model_len", 4096))

print("tok loaded")
print("MAX_MODEL_LEN from server:", MAX_MODEL_LEN)
```

tok loaded
MAX_MODEL_LEN from server: 4096

Run Load Test on Bad Server Config

```
In [ ]: import pandas as pd

WL_FILE = "mix_50_50.jsonl"
CONCURRENCY = 32
TOTAL_REQ = 200
DESIRED_MAX_TOKENS = 16

df_bad, summ_bad = run_load(
    prompt_sets[WL_FILE],
    concurrency=CONCURRENCY,
    total_requests=TOTAL_REQ,
    desired_max_tokens=DESIRED_MAX_TOKENS,
    desc=f"BAD_unfair | {WL_FILE} | c={CONCURRENCY}"
)

print("BAD summary:", summ_bad)
df_bad.head()
```

BAD_unfair | mix_50_50.jsonl | c=32: 0% | 0/200 [00:00<?, ?it/s]
BAD summary: {'concurrency': 32, 'total_requests': 200, 'ok': 200, 'fail': 0, 'wall_time_s': 42.19113462899986, 'rps': 4.740332341347631, 'p50_s': 6.181310366500156, 'p90_s': 10.99196271639994, 'p99_s': 11.637151490029915, 'mean_s': 6.493085269189996}

Out []:	latency_s	prompt_tokens	completion_tokens	total_tokens	status_code	prompt
0	8.678932	37	16	53	200	
1	8.703615	40	16	56	200	
2	8.683890	42	16	58	200	
3	8.700286	37	16	53	200	
4	8.697547	37	16	53	200	

Compute Overall and Short-Prompt Tail Latency Metrics (p50/p90/p99)

```
In [ ]: def tail_metrics(df, label, short_cutoff_tokens=64):
    ok = df[df["ok"] == True].copy()

    # overall
    overall = {
        "label": label,
        "subset": "overall",
        "n": len(ok),
        "p50": float(ok["latency_s"].quantile(0.50)),
        "p90": float(ok["latency_s"].quantile(0.90)),
        "p99": float(ok["latency_s"].quantile(0.99)),
        "mean": float(ok["latency_s"].mean()),
    }

    # short-only subset (inside the mixed workload)
    short = ok[ok["prompt_tokens"] <= short_cutoff_tokens].copy()
    if len(short) == 0:
        short_stats = {"label": label, "subset": f"short<= {short_cutoff_tokens}"}
        "p50": None, "p90": None, "p99": None, "mean": None
    else:
        short_stats = {
            "label": label,
            "subset": f"short<= {short_cutoff_tokens}",
            "n": len(short),
            "p50": float(short["latency_s"].quantile(0.50)),
            "p90": float(short["latency_s"].quantile(0.90)),
            "p99": float(short["latency_s"].quantile(0.99)),
            "mean": float(short["latency_s"].mean()),
        }

    return overall, short_stats
```

```
In [ ]: overall_bad, short_bad = tail_metrics(df_bad, "BAD_unfair", short_cutoff_tokens=64)
overall_bad, short_bad
```

```
Out[ ]: ({'label': 'BAD_unfair',
  'subset': 'overall',
  'n': 200,
  'p50': 6.181310366500156,
  'p90': 10.99196271639994,
  'p99': 11.637151490029915,
  'mean': 6.493085269189996},
{'label': 'BAD_unfair',
  'subset': 'short<= 64',
  'n': 108,
  'p50': 5.480035559000271,
  'p90': 8.717383960199822,
  'p99': 10.618033962269847,
  'mean': 5.465330376453692})
```

Experiment 3: Tail latency and Throughput variation with Concurrency Sweep

This sweep isolates scheduling effects by holding workload and max batched token length constant, then measuring how tail latency and throughput evolve with concurrency.

```
In [ ]: from google.colab import drive
drive.mount("/content/drive")

import os, json, shutil, time
import pandas as pd

WL_FILE = "mix_50_50.jsonl"
TOTAL_REQ = 200
DESIRED_MAX_TOKENS = 16
CONCURRENCY_LIST = [1, 2, 4, 8, 16, 32, 48, 64]

SHORT_CUTOFF = 64
LONG_CUTOFF = 2048

LABEL = "BAD_unfair"

RUN_DIR = "/content/drive/MyDrive/vllm_sweeps/exp2_bad_unfair"
os.makedirs(RUN_DIR, exist_ok=True)

CKPT_CSV      = os.path.join(RUN_DIR, "concurrency_checkpoint.csv")
MANIFEST_JSON = os.path.join(RUN_DIR, "concurrency_manifest.json")

SAVE_RAW_DF = False
RAW_DIR = os.path.join(RUN_DIR, "raw_dfs")
if SAVE_RAW_DF:
    os.makedirs(RAW_DIR, exist_ok=True)

manifest = {
    "label": LABEL,
    "workload_file": WL_FILE,
    "total_requests": TOTAL_REQ,
    "desired_max_tokens": DESIRED_MAX_TOKENS,
    "concurrency_list": CONCURRENCY_LIST,
    "short_cutoff": SHORT_CUTOFF,
```

```

        "long_cutoff": LONG_CUTOFF,
    }

    if os.path.exists(MANIFEST_JSON):
        try:
            old = json.loads(open(MANIFEST_JSON, "r").read())
            if old != manifest:
                print("Manifest differs from existing checkpoint. You may be")
                print("Existing:", old)
                print("New:", manifest)
        except Exception:
            print("Could not read existing manifest. Proceeding.")
    else:
        with open(MANIFEST_JSON, "w") as f:
            json.dump(manifest, f, indent=2)

    def summarize_run(df, summ, label, workload, concurrency):
        ok = df[df["ok"] == True].copy()

        row = {
            "label": label,
            "workload": workload,
            "concurrency": int(concurrency),

            "ok": summ.get("ok"),
            "fail": summ.get("fail"),
            "rps": summ.get("rps"),
            "wall_time_s": summ.get("wall_time_s"),
            "p50_s": summ.get("p50_s"),
            "p90_s": summ.get("p90_s"),
            "p99_s": summ.get("p99_s"),
            "mean_s": summ.get("mean_s"),
        }

        short = ok[ok["prompt_tokens"] <= SHORT_CUTOFF]
        row["short_n"] = int(len(short))
        row["short_p50_s"] = float(short["latency_s"].quantile(0.50)) if len(
            short) > 0 else None
        row["short_p90_s"] = float(short["latency_s"].quantile(0.90)) if len(
            short) > 0 else None
        row["short_p99_s"] = float(short["latency_s"].quantile(0.99)) if len(
            short) > 0 else None

        longg = ok[ok["prompt_tokens"] >= LONG_CUTOFF]
        row["long_n"] = int(len(longg))
        row["long_p50_s"] = float(longg["latency_s"].quantile(0.50)) if len(
            longg) > 0 else None
        row["long_p90_s"] = float(longg["latency_s"].quantile(0.90)) if len(
            longg) > 0 else None
        row["long_p99_s"] = float(longg["latency_s"].quantile(0.99)) if len(
            longg) > 0 else None

    return row

    def load_ckpt_df(path):
        if not os.path.exists(path):
            return pd.DataFrame()
        df = pd.read_csv(path)
        if "concurrency" in df.columns:
            df["concurrency"] = pd.to_numeric(df["concurrency"], errors="coer")
        return df

    def done_concurrenties(df_ckpt):
        if df_ckpt.empty:
            return set()
        good = df_ckpt[df_ckpt["concurrency"].notna()].copy()

```

```

if "status" in good.columns:
    good = good[good["status"].isna()].copy()
return set(int(x) for x in good["concurrency"].dropna().tolist())

def append_ckpt_row(row, path):
    df_row = pd.DataFrame([row])
    header = not os.path.exists(path)
    df_row.to_csv(path, mode="a", header=header, index=False)

def maybe_save_raw_df(df, concurrency):
    if not SAVE_RAW_DF:
        return
    outp = os.path.join(RAW_DIR, f"raw_c{int(concurrency)}.parquet")
    try:
        df.to_parquet(outp, index=False)
    except Exception:
        outp = os.path.join(RAW_DIR, f"raw_c{int(concurrency)}.csv")
        df.to_csv(outp, index=False)

ckpt_df = load_ckpt_df(CKPT_CSV)
done = done_concurrenties(ckpt_df)
print("Loaded checkpoint rows:", len(ckpt_df))
print("Completed concurrenties:", sorted(done))

for c in CONCURRENCY_LIST:
    c = int(c)
    if c in done:
        print("Skipping:", c)
        continue

    print("Running:", LABEL, WL_FILE, "c=", c)

    try:
        df, summ = run_load(
            prompt_sets[WL_FILE],
            concurrency=c,
            total_requests=TOTAL_REQ,
            desired_max_tokens=DESIRED_MAX_TOKENS,
            desc=f"{LABEL} | {WL_FILE} | c={c}",
        )

        row = summarize_run(df, summ, label=LABEL, workload=WL_FILE, concurrency=c)
        append_ckpt_row(row, CKPT_CSV)
        maybe_save_raw_df(df, concurrency=c)

        ckpt_df = load_ckpt_df(CKPT_CSV)
        done = done_concurrenties(ckpt_df)

    except Exception as e:
        fail_row = {
            "label": LABEL,
            "workload": WL_FILE,
            "concurrency": c,
            "ok": None,
            "fail": None,
            "rps": None,
            "wall_time_s": None,
            "p50_s": None,
            "p90_s": None,
            "p99_s": None,
        }

```

```

        "mean_s": None,
        "short_n": None,
        "short_p50_s": None,
        "short_p90_s": None,
        "short_p99_s": None,
        "long_n": None,
        "long_p50_s": None,
        "long_p90_s": None,
        "long_p99_s": None,
        "status": "run_failed",
        "error": type(e).__name__,
    }
append_ckpt_row(fail_row, CKPT_CSV)
print("run_load failed:", type(e).__name__, str(e)[:200])

summary_df = load_ckpt_df(CKPT_CSV).copy()
if "status" in summary_df.columns:
    summary_df = summary_df[summary_df["status"].isna()].copy()

summary_df = summary_df.sort_values("concurrency").reset_index(drop=True)

print("Concurrency sweep done/resumed. Summary DF loaded from Drive:")
display(summary_df)

FINAL_CSV = os.path.join(RUN_DIR, "summary_df_final.csv")
summary_df.to_csv(FINAL_CSV, index=False)
print("Final snapshot saved:", FINAL_CSV)

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Loaded checkpoint rows: 0

Completed concurrencies: []

Running: BAD_unfair mix_50_50.jsonl c= 1	
BAD_unfair mix_50_50.jsonl c=1: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 2	
BAD_unfair mix_50_50.jsonl c=2: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 4	
BAD_unfair mix_50_50.jsonl c=4: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 8	
BAD_unfair mix_50_50.jsonl c=8: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 16	
BAD_unfair mix_50_50.jsonl c=16: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 32	
BAD_unfair mix_50_50.jsonl c=32: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 48	
BAD_unfair mix_50_50.jsonl c=48: 0%	0/200 [00:00<?, ?it/s]
Running: BAD_unfair mix_50_50.jsonl c= 64	
BAD_unfair mix_50_50.jsonl c=64: 0%	0/200 [00:00<?, ?it/s]

Concurrency sweep done/resumed. Summary DF loaded from Drive:

label	workload	concurrency	ok	fail	rps	wall_time_s	p50
0	BAD_unfair	mix_50_50.jsonl	1	200	0	1.310728	152.586979
1	BAD_unfair	mix_50_50.jsonl	2	200	0	2.350512	85.087858
2	BAD_unfair	mix_50_50.jsonl	4	200	0	3.588961	55.726434
3	BAD_unfair	mix_50_50.jsonl	8	200	0	4.897939	40.833506
4	BAD_unfair	mix_50_50.jsonl	16	200	0	5.902047	33.886548
5	BAD_unfair	mix_50_50.jsonl	32	200	0	5.912537	33.826429
6	BAD_unfair	mix_50_50.jsonl	48	200	0	6.029853	33.168302
7	BAD_unfair	mix_50_50.jsonl	64	200	0	5.817353	34.379898

Final snapshot saved: /content/drive/MyDrive/vllm_sweeps/exp2_bad_unfair/summary_df_final.csv

Plot some meaningful charts from Sweep results

```
In [ ]: df = summary_df.copy()

num_cols = [
    "concurrency", "rps", "wall_time_s", "p50_s", "p90_s", "p99_s", "mean_s",
    "short_p50_s", "short_p90_s", "short_p99_s",
    "long_p50_s", "long_p90_s", "long_p99_s",
    "short_n", "long_n"
]
for c in num_cols:
    if c in df.columns:
        df[c] = pd.to_numeric(df[c], errors="coerce")

df = df.sort_values("concurrency").reset_index(drop=True)

# output directory
PLOT_DIR = "plots_bad_sweep"
os.makedirs(PLOT_DIR, exist_ok=True)

def savefig(name):
    path = os.path.join(PLOT_DIR, name)
    plt.tight_layout()
    plt.savefig(path, dpi=200, bbox_inches="tight")
    print("Saved:", path)

# Throughput vs concurrency
plt.figure()
plt.plot(df["concurrency"], df["rps"], marker="o")
plt.xlabel("Client concurrency")
plt.ylabel("Throughput (requests/sec)")
plt.title("BAD_unfair: Throughput vs Concurrency")
plt.grid(True, alpha=0.3)
savefig("bad_throughput_vs_concurrency.png")
plt.show()

# Overall tail latency vs concurrency (p50 / p90 / p99)
```

```

for metric, title in [
    ("p50_s", "BAD_unfair: Overall p50 vs Concurrency"),
    ("p90_s", "BAD_unfair: Overall p90 vs Concurrency"),
    ("p99_s", "BAD_unfair: Overall p99 vs Concurrency"),
]:
    plt.figure()
    plt.plot(df["concurrency"], df[metric], marker="o")
    plt.xlabel("Client concurrency")
    plt.ylabel(f"Latency (s) - {metric.replace('_s', '')}")
    plt.title(title)
    plt.grid(True, alpha=0.3)
    savefig(f"bad_{metric}_vs_concurrency.png")
    plt.show()

# Short vs Long p99 comparison (same plot)
plt.figure()
plt.plot(df["concurrency"], df["short_p99_s"], marker="o", label="short p99")
plt.plot(df["concurrency"], df["long_p99_s"], marker="o", label="long p99")
plt.xlabel("Client concurrency")
plt.ylabel("Latency (s)")
plt.title("BAD_unfair: Short vs Long p99 vs Concurrency")
plt.grid(True, alpha=0.3)
plt.legend()
savefig("bad_short_long_p99_vs_concurrency.png")
plt.show()

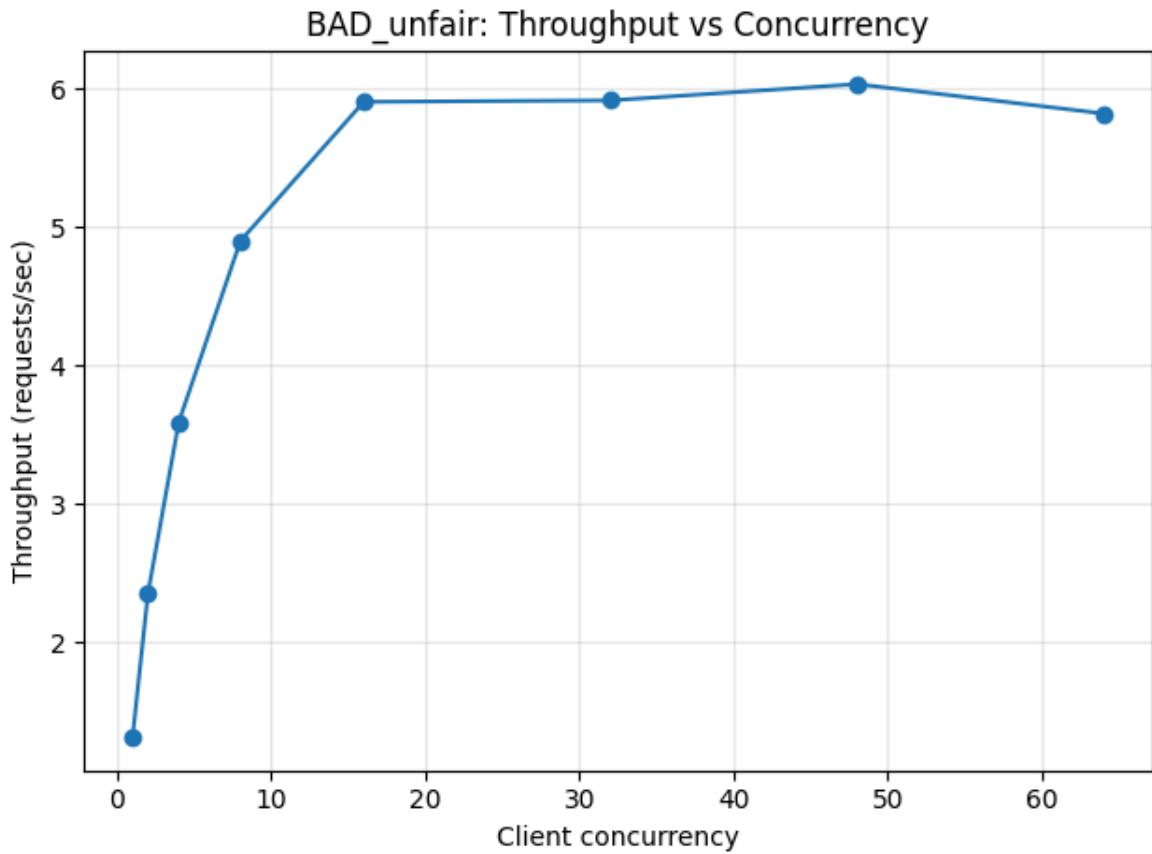
# "Knee" visualization: p99 vs RPS (tradeoff curve)
plt.figure()
plt.plot(df["rps"], df["p99_s"], marker="o")
for _, r in df.iterrows():
    plt.text(r["rps"], r["p99_s"], str(int(r["concurrency"])), fontsize=9)

plt.xlabel("Throughput (requests/sec)")
plt.ylabel("Overall p99 latency (s)")
plt.title("BAD_unfair: p99 vs Throughput (labels = concurrency)")
plt.grid(True, alpha=0.3)
savefig("bad_p99_vs_rps.png")
plt.show()

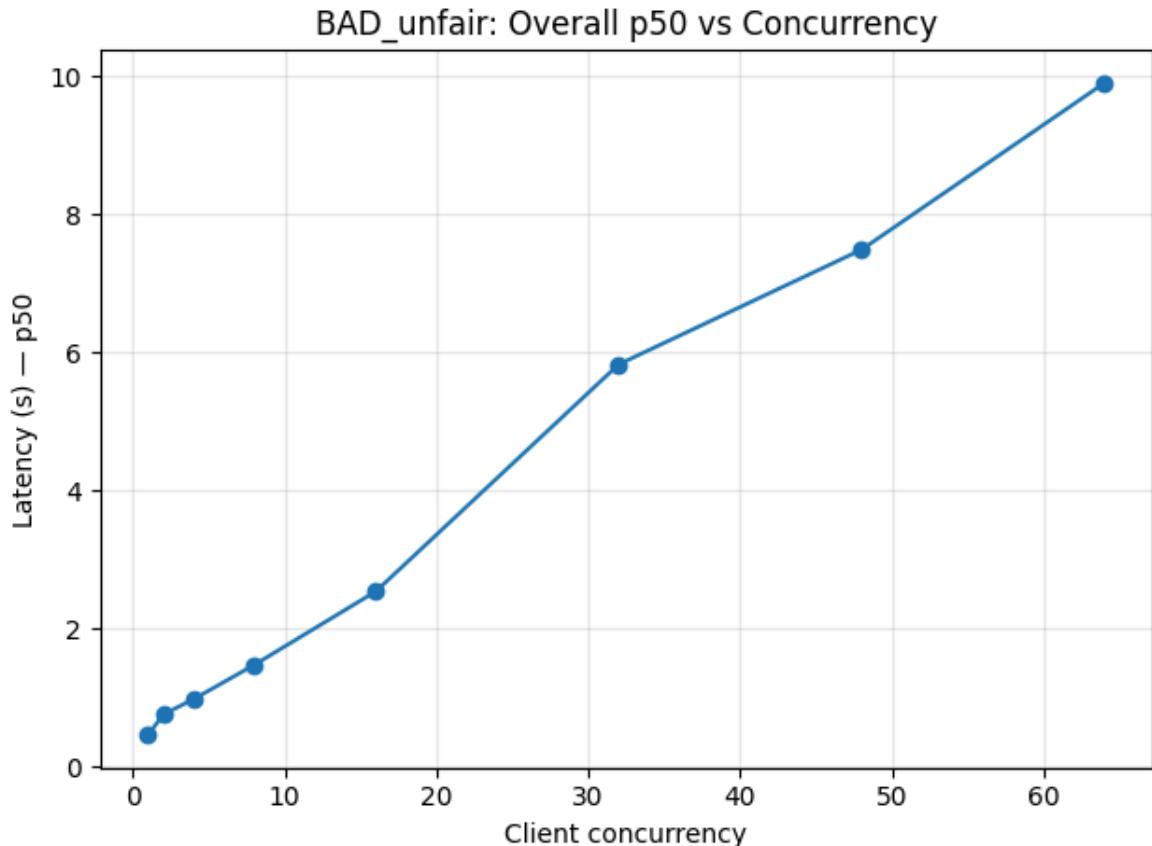
print("Done. Plot folder:", PLOT_DIR)
df[["concurrency", "rps", "p50_s", "p90_s", "p99_s", "short_p99_s", "long_p99_s"]]

```

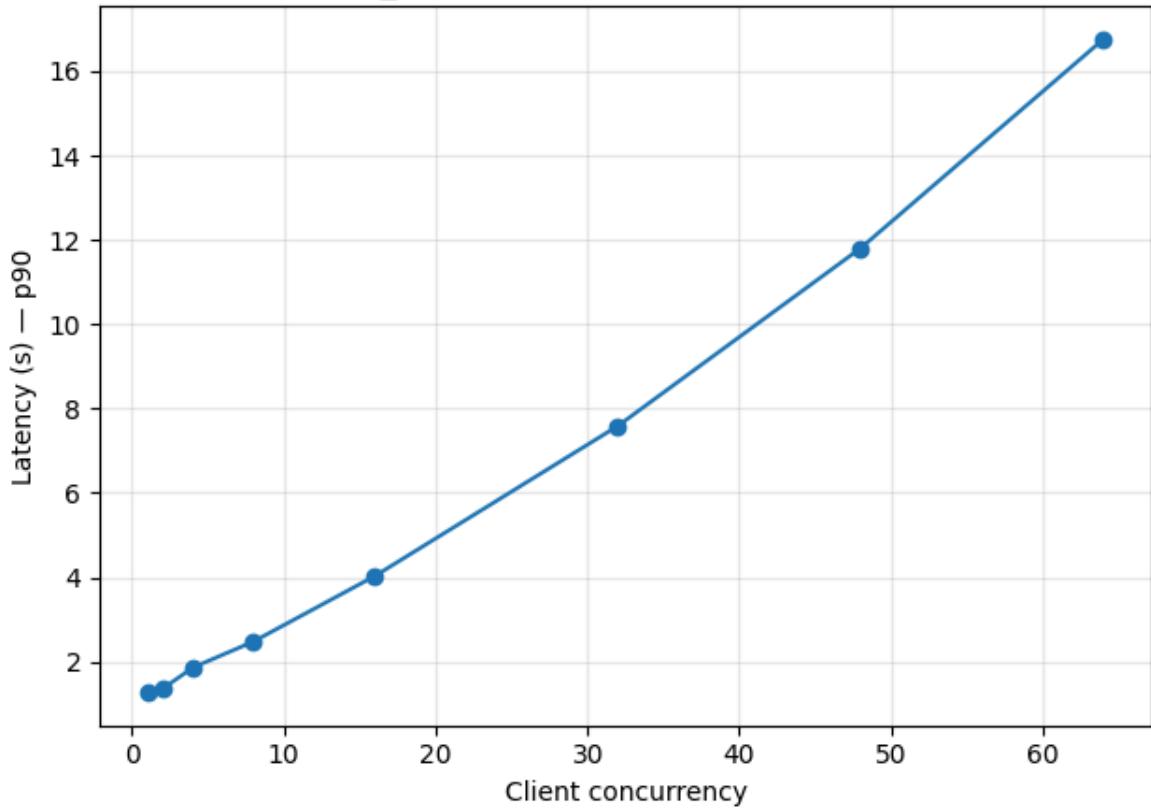
Saved: plots_bad_sweep/bad_throughput_vs_concurrency.png



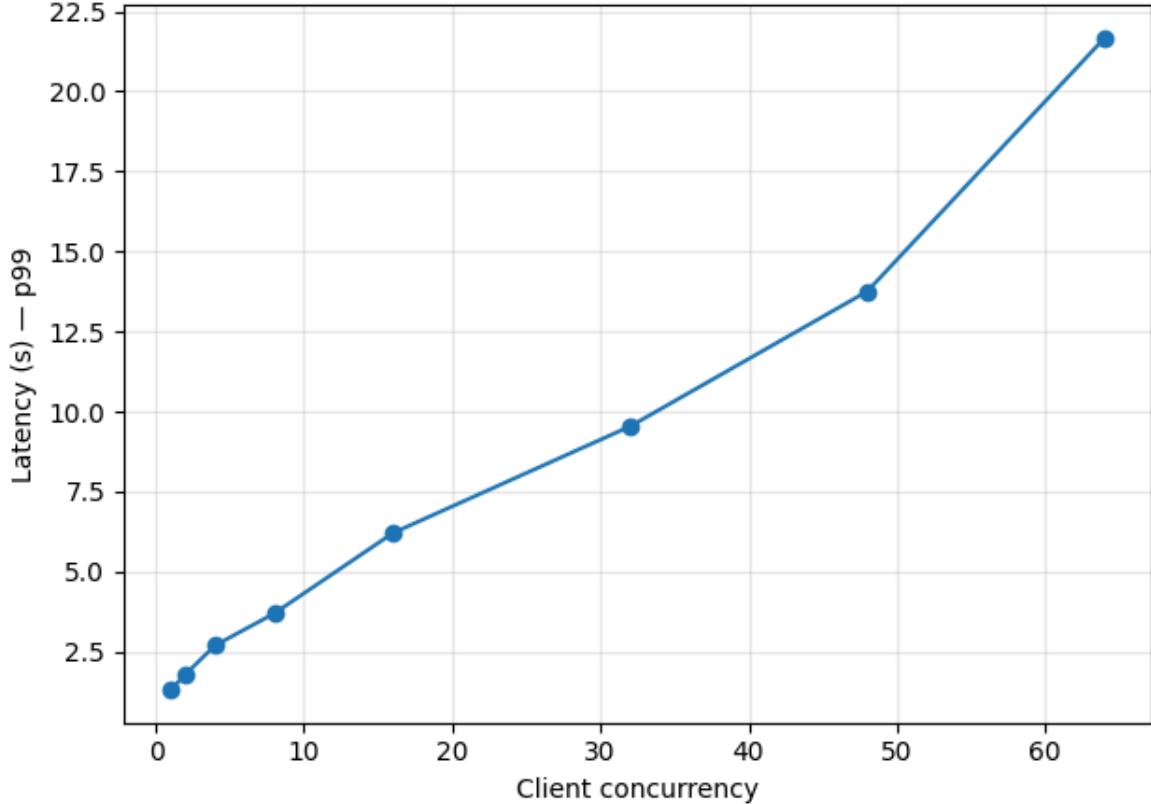
Saved: plots_bad_sweep/bad_p50_s_vs_concurrency.png



Saved: plots_bad_sweep/bad_p90_s_vs_concurrency.png

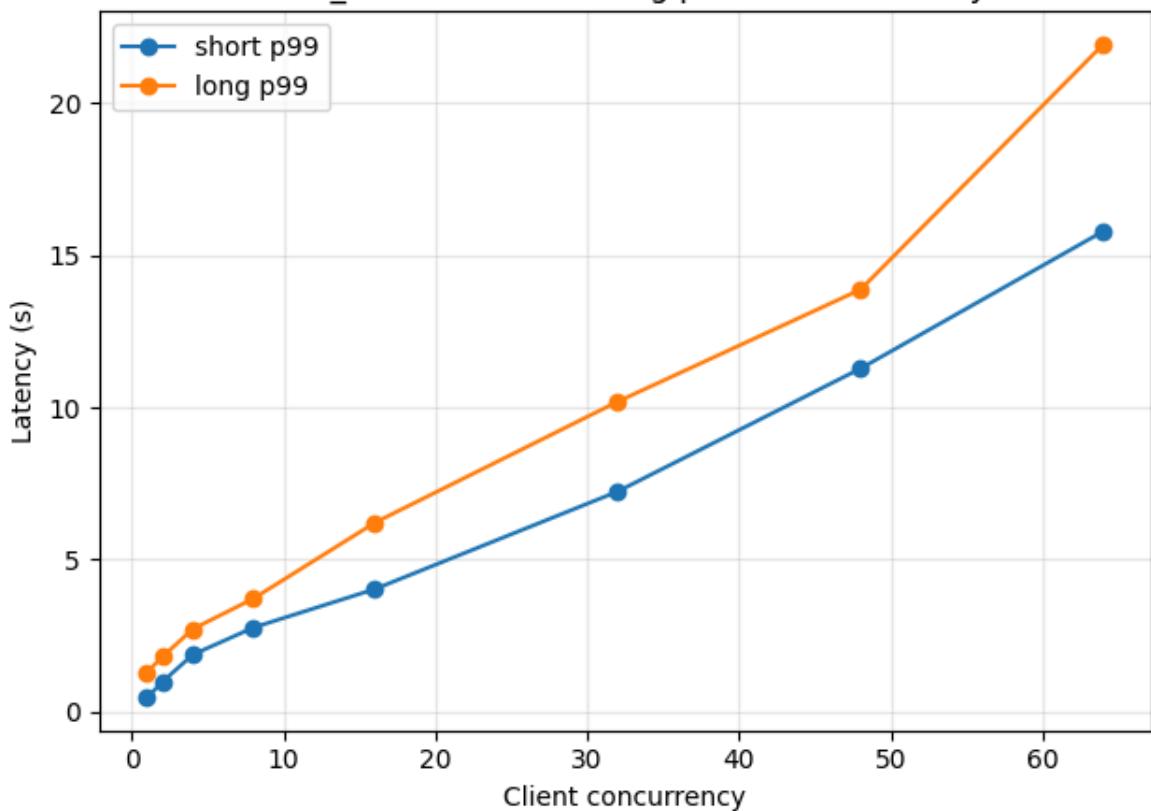
BAD_unfair: Overall p90 vs Concurrency

Saved: plots_bad_sweep/bad_p99_s_vs_concurrency.png

BAD_unfair: Overall p99 vs Concurrency

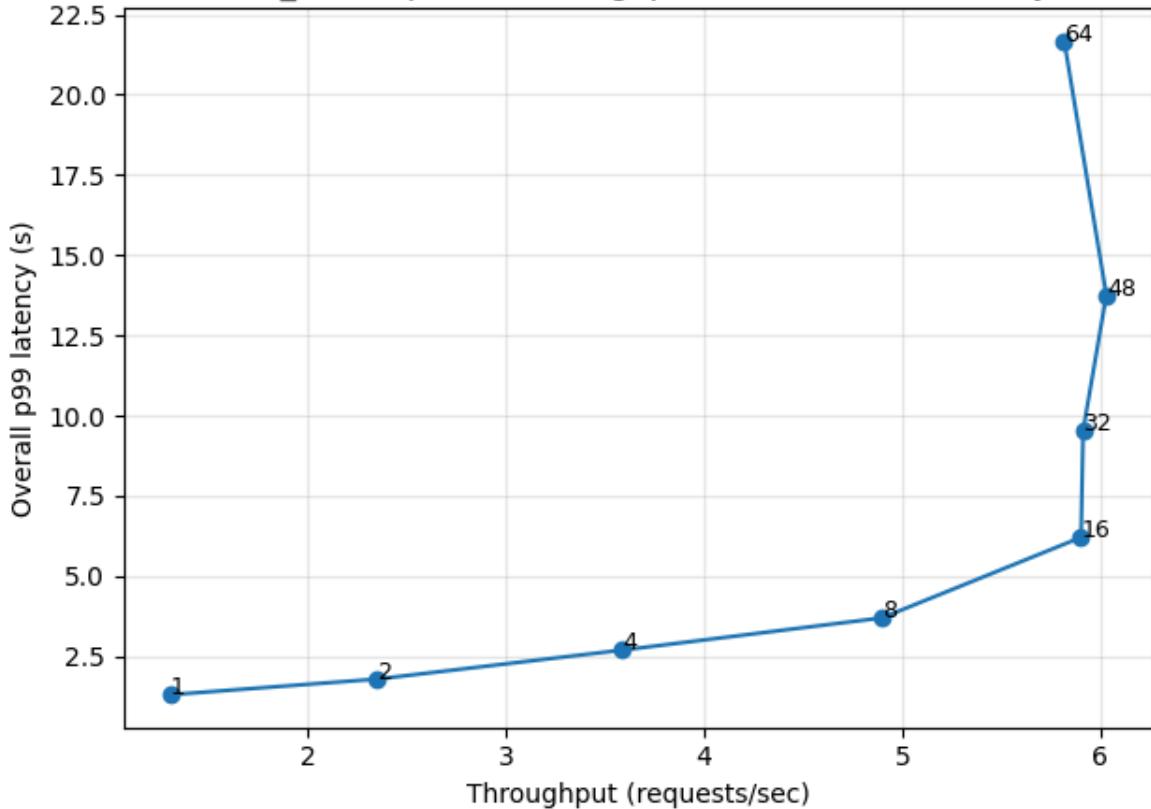
Saved: plots_bad_sweep/bad_short_long_p99_vs_concurrency.png

BAD_unfair: Short vs Long p99 vs Concurrency



Saved: plots_bad_sweep/bad_p99_vs_rps.png

BAD_unfair: p99 vs Throughput (labels = concurrency)



Done. Plot folder: plots_bad_sweep

Out []:

	concurrency	rps	p50_s	p90_s	p99_s	short_p99_s	long_p99_s
0	1	1.310728	0.455211	1.264692	1.302439	0.459093	1.3033
1	2	2.350512	0.753871	1.356037	1.788700	0.975944	1.7973
2	4	3.588961	0.974289	1.860058	2.692315	1.870385	2.7110
3	8	4.897939	1.472894	2.478667	3.694609	2.762490	3.7158
4	16	5.902047	2.537444	4.032396	6.204649	4.027770	6.2103
5	32	5.912537	5.829220	7.589303	9.536481	7.235541	10.1896
6	48	6.029853	7.492978	11.804981	13.755477	11.278168	13.8723
7	64	5.817353	9.910197	16.760385	21.667766	15.790760	21.937

Conclusion

- **Throughput scales then saturates:** RPS rises from ~1.31 (c=1) to ~5.90 (c=16) and then stays around ~5.8–6.0 RPS through c=64, showing the server hits a throughput ceiling.
- **Clear knee and tail blow-up:** overall p99 grows from ~1.30s (c=1) to ~6.20s (c=16), then jumps to ~9.54s (c=32) and reaches ~21.67s (c=64).
- **Short requests get dragged into the tail:** short p99 inflates from ~0.46s (c=1) to ~4.03s (c=16) and ~15.79s (c=64), which indicates strong queueing interference from long prompts.
- **At high concurrency, short and long tails converge:** by c=64, short p99 ~15.79s and long p99 ~21.94s, meaning even short prompts suffer large delays under this bad/unfair configuration.

Saving plots to Drive

```
In [ ]: LOCAL_PLOT_DIR = "plots_bad_sweep"
DRIVE_PLOT_DIR = "/content/drive/MyDrive/vllm_sweeps/exp2_bad_unfair/plot

if os.path.exists(DRIVE_PLOT_DIR):
    shutil.rmtree(DRIVE_PLOT_DIR)
shutil.copytree(LOCAL_PLOT_DIR, DRIVE_PLOT_DIR)

print("Copied plots to:", DRIVE_PLOT_DIR)
```

Copied plots to: /content/drive/MyDrive/vllm_sweeps/exp2_bad_unfair/plots_bad_sweep

Experiment 4: Sweep of concurrency and max-num-seqs

In this sweep, we stress-test a deliberately “bad” vLLM server configuration under a mixed workload (mix_50_50) to quantify how tail latency behaves as system load

increases. By sweeping both client concurrency and the server's admission capacity (max_num_seqs), we expose where throughput stops scaling, where queueing begins to dominate, and how long-prefill requests can degrade the latency of short requests in a shared serving setup. The goal is to establish a clear baseline of failure modes before introducing improved scheduling/configurations in the next steps.

```
In [ ]: BASE = "http://127.0.0.1:8000"
MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

CHECKPOINT_CSV = "bad_grid_checkpoint.csv"
MANIFEST_JSON = "bad_grid_manifest.json"

def kill_server():
    os.system('pkill -f "vllm.entrypoints.openai.api_server" || true')
    os.system('pkill -f "VLLM::EngineCore" || true')
    time.sleep(1.5)

def tail_log(log_path, n=60):
    if os.path.exists(log_path):
        print(f"\n--- tail {log_path} (last {n}) ---")
        os.system(f"tail -n {n} {log_path}")
    else:
        print(f"(log missing: {log_path})")

def wait_for_ready(base=BASE, log_path="vllm_bad.log", timeout_s=420):
    t0 = time.time()
    tries = 0
    while True:
        tries += 1
        try:
            r = requests.get(f"{base}/v1/models", timeout=2)
            if r.status_code == 200:
                data = r.json().get("data", [])
                ids = [m.get("id") for m in data]
                print(f" Server ready. Models: {ids}")
                return True
            else:
                print(f"[{tries}] /v1/models status={r.status_code}")
        except Exception as e:
            print(f"[{tries}] not ready yet: {type(e).__name__}")

    alive = os.system('ps -ef | grep -E "vllm.entrypoints.openai.api_'
    if not alive:
        print(" vLLM process is NOT running anymore.")
        tail_log(log_path, n=140)
        return False

    if time.time() - t0 > timeout_s:
        print(" Timed out waiting for server.")
        tail_log(log_path, n=140)
        return False

    if tries % 5 == 0:
        tail_log(log_path, n=30)

    time.sleep(2)
```

```

def warmup_request():
    payload = {
        "model": MODEL_ID,
        "messages": [{"role": "user", "content": "hi"}],
        "max_tokens": 8,
        "temperature": 0.0,
    }
    try:
        _ = requests.post(f"{BASE}/v1/chat/completions", json=payload, ti
    except Exception:
        pass

def start_bad_server(max_num_seqs, log_path="vllm_bad.log"):
    kill_server()
    try:
        os.remove(log_path)
    except FileNotFoundError:
        pass

    cmd = f"""
nohup python3 -m vllm.entrypoints.openai.api_server \
--model {MODEL_ID} \
--host 127.0.0.1 --port 8000 \
--dtype half \
--max-model-len 4096 \
--gpu-memory-utilization 0.70 \
--disable-log-stats \
--enforce-eager \
--max-num-batched-tokens 8192 \
--long-prefill-token-threshold 1024 \
--max-num-seqs {int(max_num_seqs)} \
> {log_path} 2>&1 </dev/null &
"""
    os.system(cmd)
    time.sleep(1.0)

def _safe_quantile(series, q):
    if series is None or len(series) == 0:
        return None
    return float(series.quantile(q))

def _safe_mean(series):
    if series is None or len(series) == 0:
        return None
    return float(series.mean())

def subset_metrics(df, short_cutoff_tokens=64, long_cutoff_tokens=2048):
    ok = df[df["ok"] == True].copy()
    lat = ok["latency_s"]

    out = {
        "overall_n": int(len(ok)),
        "overall_p50": _safe_quantile(lat, 0.50),
        "overall_p90": _safe_quantile(lat, 0.90),
        "overall_p99": _safe_quantile(lat, 0.99),
        "overall_mean": _safe_mean(lat),
    }

```

```

short = ok[ok["prompt_tokens"] <= short_cutoff_tokens]
out.update({
    "short_n": int(len(short)),
    "short_p50": _safe_quantile(short["latency_s"], 0.50),
    "short_p90": _safe_quantile(short["latency_s"], 0.90),
    "short_p99": _safe_quantile(short["latency_s"], 0.99),
    "short_mean": _safe_mean(short["latency_s"]),
})

long = ok[ok["prompt_tokens"] >= long_cutoff_tokens]
out.update({
    "long_n": int(len(long)),
    "long_p50": _safe_quantile(long["latency_s"], 0.50),
    "long_p90": _safe_quantile(long["latency_s"], 0.90),
    "long_p99": _safe_quantile(long["latency_s"], 0.99),
    "long_mean": _safe_mean(long["latency_s"]),
})
return out

def row_from_run(label, workload, concurrency, max_num_seqs, df, summ):
    m = subset_metrics(df, short_cutoff_tokens=64, long_cutoff_tokens=204)
    return {
        "label": label,
        "workload": workload,
        "concurrency": int(concurrency),
        "max_num_seqs": int(max_num_seqs),

        "status": "ok",
        "ok": int(summ.get("ok")) if summ.get("ok") is not None else int(
            "fail": int(summ.get("fail")) if summ.get("fail") is not None else
            "rps": float(summ.get("rps")) if summ.get("rps") is not None else
            "wall_time_s": float(summ.get("wall_time_s")) if summ.get("wall_t
            "summ_p50_s": float(summ.get("p50_s")) if summ.get("p50_s") is no
            "summ_p90_s": float(summ.get("p90_s")) if summ.get("p90_s") is no
            "summ_p99_s": float(summ.get("p99_s")) if summ.get("p99_s") is no
            "summ_mean_s": float(summ.get("mean_s")) if summ.get("mean_s") is
        )
    }

def load_checkpoint_df(path=CHECKPOINT_CSV):
    if os.path.exists(path):
        df = pd.read_csv(path)
        # normalize types
        if "max_num_seqs" in df.columns:
            df["max_num_seqs"] = df["max_num_seqs"].astype(int)
        if "concurrency" in df.columns:
            df["concurrency"] = df["concurrency"].astype(int)
        return df
    return pd.DataFrame()

def completed_pairs(df_ckpt):
    # Only treat status=="ok" as completed (you can also include server_
    if df_ckpt.empty:
        return set()
    done = df_ckpt[df_ckpt["status"] == "ok"][[ "max_num_seqs", "concurrent
    return set((int(r.max_num_seqs), int(r.concurrency))) for r in done.it

def append_row_checkpoint(row, path=CHECKPOINT_CSV):
    df_row = pd.DataFrame([row])

```

```

header = not os.path.exists(path)
df_row.to_csv(path, mode="a", header=header, index=False)

def write_manifest(payload, path=MANIFEST_JSON):
    # don't overwrite if already exists unless identical
    if os.path.exists(path):
        try:
            old = json.loads(open(path, "r").read())
            if old != payload:
                print(" Manifest differs from existing one. You may be re")
                print("Existing:", old)
                print("New:", payload)
        except Exception:
            print(" Could not read existing manifest. Proceeding.")
    return
with open(path, "w") as f:
    json.dump(payload, f, indent=2)

# sweep runner

WL_FILE = "mix_50_50.jsonl"
TOTAL_REQ = 200
DESIRED_MAX_TOKENS = 16

concurrency_list = [1, 2, 4, 8, 16, 32, 48, 64]
max_num_seqs_list = [16, 32, 48, 64, 80, 96]

# save manifest to know what this checkpoint corresponds to
write_manifest({
    "workload": WL_FILE,
    "total_requests": TOTAL_REQ,
    "desired_max_tokens": DESIRED_MAX_TOKENS,
    "concurrency_list": concurrency_list,
    "max_num_seqs_list": max_num_seqs_list,
    "model_id": MODEL_ID,
    "server_flags": {
        "max_num_batched_tokens": 8192,
        "long_prefill_token_threshold": 1024,
        "gpu_memory_utilization": 0.70,
        "max_model_len": 4096,
        "dtype": "half",
        "enforce_eager": True
    }
})

ckpt_df = load_checkpoint_df(CHECKPOINT_CSV)
done = completed_pairs(ckpt_df)
print(f" Loaded checkpoint rows: {len(ckpt_df)} | completed runs: {len(done)}")

for max_num_seqs in max_num_seqs_list:
    # If all concurrencies for this max_num_seqs are done, skip server re
    all_done_for_this = all((max_num_seqs, c) in done for c in concurrency_list)
    if all_done_for_this:
        print(f"\n Skipping max_num_seqs={max_num_seqs} (all concurrencie")
        continue

    print("\n=====")
    print(f"Starting BAD server with max-num-seqs={max_num_seqs}")

```

```

print("-----")
start_bad_server(max_num_seqs=max_num_seqs, log_path="vllm_bad.log")

ok = wait_for_ready(BASE, log_path="vllm_bad.log", timeout_s=420)
if not ok:
    # record server_failed for missing concurrencies only
    for c in concurrency_list:
        if (max_num_seqs, c) in done:
            continue
        row = {
            "label": "BAD_unfair",
            "workload": WL_FILE,
            "concurrency": int(c),
            "max_num_seqs": int(max_num_seqs),
            "status": "server_failed",
        }
        append_row_checkpoint(row, CHECKPOINT_CSV)
    # refresh checkpoint view
    ckpt_df = load_checkpoint_df(CHECKPOINT_CSV)
    done = completed_pairs(ckpt_df)
    continue

warmup_request()

for c in concurrency_list:
    if (max_num_seqs, c) in done:
        print(f" Skip (already done): max_seqs={max_num_seqs}, c={c}")
        continue

    print(f"\n--- RUN: BAD_unfair | {WL_FILE} | max_seqs={max_num_seqs}")

    try:
        df_bad, summ_bad = run_load(
            prompt_sets[WL_FILE],
            concurrency=c,
            total_requests=TOTAL_REQ,
            desired_max_tokens=DESIRED_MAX_TOKENS,
            desc=f"BAD_unfair | {WL_FILE} | max_seqs={max_num_seqs} |"
        )

        row = row_from_run(
            label="BAD_unfair",
            workload=WL_FILE,
            concurrency=c,
            max_num_seqs=max_num_seqs,
            df=df_bad,
            summ=summ_bad,
        )
        append_row_checkpoint(row, CHECKPOINT_CSV)

    except Exception as e:
        # record run_failed so you can retry later
        row = {
            "label": "BAD_unfair",
            "workload": WL_FILE,
            "concurrency": int(c),
            "max_num_seqs": int(max_num_seqs),
            "status": "run_failed",
            "error": type(e).__name__,
        }

```

```

append_row_checkpoint(row, CHECKPOINT_CSV)
print(" run_load failed:", type(e).__name__)
tail_log("vllm_bad.log", n=80)

# refresh done set so mid-run restarts still work cleanly
ckpt_df = load_checkpoint_df(CHECKPOINT_CSV)
done = completed_pairs(ckpt_df)

kill_server()

summary_df = load_checkpoint_df(CHECKPOINT_CSV)
print("\n BAD (concurrency x max-num-seqs) sweep (resume-capable) done.")
summary_df

```

✓ Loaded checkpoint rows: 0 | completed runs: 0

=====
Starting BAD server with max-num-seqs=16
=====

```
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError
```

```
--- tail vllm_bad.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError
```

```
--- tail vllm_bad.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError
```

```
--- tail vllm_bad.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError
```

```
--- tail vllm_bad.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
```

✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']

```
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=1 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=1: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=2 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=2: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=4 ---
```

```

BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=4: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=8 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=8: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=16 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=16: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=32 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=32: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=48 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=16 | c=48: 0%| 0/200
[00:00<?, ?it/s]
=====
Starting BAD server with max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
[25] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=1 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=1: 0%| 0/200
[00:00<?, ?it/s]

```

```

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=2 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=2: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=4 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=4: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=8 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=8: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=16 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=16: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=32 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=32: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=48 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=32 | c=48: 0%| 0/200
[00:00<?, ?it/s]
=====
Starting BAD server with max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=1 ---

```

```
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=1: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=2 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=2: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=4 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=4: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=8 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=8: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=16 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=16: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=32 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=32: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=48 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=48: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=64 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=48 | c=64: 0%| 0/200
[00:00<?, ?it/s]
```

```
=====
Starting BAD server with max-num-seqs=64
=====

[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=1 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=1: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=2 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=2: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=4 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=4: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=8 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=8: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=16 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=16: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=32 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=32: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=48 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=48: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=64 ---
```

```

BAD_unfair | mix_50_50.jsonl | max_seqs=64 | c=64: 0% | 0/200
[00:00<?, ?it/s]
=====
Starting BAD server with max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
[25] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=1 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=1: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=2 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=2: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=4 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=4: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=8 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=8: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=16 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=16: 0% | 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=32 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=32: 0% | 0/200
[00:00<?, ?it/s]

```

```

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=48 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=48: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=64 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=80 | c=64: 0%| 0/200
[00:00<?, ?it/s]
=====
Starting BAD server with max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_bad.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']

--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=1 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=1: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=2 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=2: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=4 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=4: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=8 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=8: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=16 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=16: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=32 ---

```

```
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=32: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=48 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=48: 0%| 0/200
[00:00<?, ?it/s]
--- RUN: BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=64 ---
BAD_unfair | mix_50_50.jsonl | max_seqs=96 | c=64: 0%| 0/200
[00:00<?, ?it/s]
 BAD (concurrency x max-num-seqs) sweep (resume-capable) done.
```

Out[]:

	label	workload	concurrency	max_num_seqs	status	ok	fail	
0	BAD_unfair	mix_50_50.jsonl	1	16	ok	200	0	
1	BAD_unfair	mix_50_50.jsonl	2	16	ok	200	0	2
2	BAD_unfair	mix_50_50.jsonl	4	16	ok	200	0	4
3	BAD_unfair	mix_50_50.jsonl	8	16	ok	200	0	7
4	BAD_unfair	mix_50_50.jsonl	16	16	ok	200	0	11
5	BAD_unfair	mix_50_50.jsonl	32	16	ok	200	0	13
6	BAD_unfair	mix_50_50.jsonl	48	16	ok	200	0	13
7	BAD_unfair	mix_50_50.jsonl	64	16	ok	200	0	13
8	BAD_unfair	mix_50_50.jsonl	1	32	ok	200	0	
9	BAD_unfair	mix_50_50.jsonl	2	32	ok	200	0	
10	BAD_unfair	mix_50_50.jsonl	4	32	ok	200	0	4
11	BAD_unfair	mix_50_50.jsonl	8	32	ok	200	0	7
12	BAD_unfair	mix_50_50.jsonl	16	32	ok	200	0	10
13	BAD_unfair	mix_50_50.jsonl	32	32	ok	200	0	15
14	BAD_unfair	mix_50_50.jsonl	48	32	ok	200	0	17
15	BAD_unfair	mix_50_50.jsonl	64	32	ok	200	0	15
16	BAD_unfair	mix_50_50.jsonl	1	48	ok	200	0	
17	BAD_unfair	mix_50_50.jsonl	2	48	ok	200	0	2
18	BAD_unfair	mix_50_50.jsonl	4	48	ok	200	0	4
19	BAD_unfair	mix_50_50.jsonl	8	48	ok	200	0	7
20	BAD_unfair	mix_50_50.jsonl	16	48	ok	200	0	11
21	BAD_unfair	mix_50_50.jsonl	32	48	ok	200	0	15
22	BAD_unfair	mix_50_50.jsonl	48	48	ok	200	0	17
23	BAD_unfair	mix_50_50.jsonl	64	48	ok	200	0	18
24	BAD_unfair	mix_50_50.jsonl	1	64	ok	200	0	
25	BAD_unfair	mix_50_50.jsonl	2	64	ok	200	0	2
26	BAD_unfair	mix_50_50.jsonl	4	64	ok	200	0	4
27	BAD_unfair	mix_50_50.jsonl	8	64	ok	200	0	7
28	BAD_unfair	mix_50_50.jsonl	16	64	ok	200	0	11
29	BAD_unfair	mix_50_50.jsonl	32	64	ok	200	0	15
30	BAD_unfair	mix_50_50.jsonl	48	64	ok	200	0	19
31	BAD_unfair	mix_50_50.jsonl	64	64	ok	200	0	20
32	BAD_unfair	mix_50_50.jsonl	1	80	ok	200	0	
33	BAD_unfair	mix_50_50.jsonl	2	80	ok	200	0	

	label	workload	concurrency	max_num_seqs	status	ok	fail
34	BAD_unfair	mix_50_50.jsonl	4	80	ok	200	0 4
35	BAD_unfair	mix_50_50.jsonl	8	80	ok	200	0 7
36	BAD_unfair	mix_50_50.jsonl	16	80	ok	200	0 10
37	BAD_unfair	mix_50_50.jsonl	32	80	ok	200	0 15
38	BAD_unfair	mix_50_50.jsonl	48	80	ok	200	0 17
39	BAD_unfair	mix_50_50.jsonl	64	80	ok	200	0 20
40	BAD_unfair	mix_50_50.jsonl	1	96	ok	200	0 1
41	BAD_unfair	mix_50_50.jsonl	2	96	ok	200	0 2
42	BAD_unfair	mix_50_50.jsonl	4	96	ok	200	0 3
43	BAD_unfair	mix_50_50.jsonl	8	96	ok	200	0 5
44	BAD_unfair	mix_50_50.jsonl	16	96	ok	200	0 8
45	BAD_unfair	mix_50_50.jsonl	32	96	ok	200	0 12
46	BAD_unfair	mix_50_50.jsonl	48	96	ok	200	0 13
47	BAD_unfair	mix_50_50.jsonl	64	96	ok	200	0 14

48 rows × 28 columns

Conclusion

- **Throughput scaling:** RPS increases with client concurrency, but the gains flatten after roughly `c=32` for most `max_num_seqs` values (diminishing returns region).
- **Tail latency inflation:** As concurrency rises, **p99 grows much faster than p50**, showing queueing and head-of-line effects dominating at high load.
- **Short vs long separation:** In the mixed workload, **short requests stay much faster than long requests**, but short p99 still inflates as concurrency increases because shorts get stuck behind long-prefill work.
- **Impact of max_num_seqs:** Larger `max_num_seqs` generally allows higher throughput at mid/high concurrency, but it can also **worsen tail latency** if the server admits too many concurrent sequences and amplifies contention/queueing under this "bad" configuration.

```
In [ ]: OUT_DIR = "/content/drive/MyDrive/vllm_sweeps/bad"
os.makedirs(OUT_DIR, exist_ok=True)

# Copy any existing on-disk artifacts to Drive (if they exist)
for f in ["bad_grid_checkpoint.csv", "bad_grid_manifest.json", "vllm_bad.json"]:
    if os.path.exists(f):
        shutil.copy(f, f"{OUT_DIR}/{f}")
```

```

        print("✅ Copied:", f"{OUT_DIR}/{f}")
    else:
        print("⚠ Missing locally:", f)

# save the final dataframe we have in memory (summary_df)
try:
    summary_df_path = f'{OUT_DIR}/bad_summary_df.csv'
    summary_df.to_csv(summary_df_path, index=False)
    print("✅ Saved summary_df:", summary_df_path)
except NameError:
    print("⚠ summary_df not found in memory. If your df has a different"
          "name, change it in the code above")

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

- ✓ Copied: /content/drive/MyDrive/vllm_sweeps/bad/bad_grid_checkpoint.csv
- ✓ Copied: /content/drive/MyDrive/vllm_sweeps/bad/bad_grid_manifest.json
- ✓ Copied: /content/drive/MyDrive/vllm_sweeps/bad/vllm_bad.log
- ✓ Saved summary_df: /content/drive/MyDrive/vllm_sweeps/bad/bad_summary_df.csv

Plots for Bidirectional Sweep in Bad Config

```

In [ ]: df = summary_df.copy()

# Keep only good rows
df = df[df["status"] == "ok"].copy()

# Ensure numeric
for col in ["concurrency", "max_num_seqs", "rps", "short_p99", "long_p99"]:
    if col in df.columns:
        df[col] = pd.to_numeric(df[col], errors="coerce")

# Sort for consistent plotting
df = df.sort_values(["max_num_seqs", "concurrency"]).reset_index(drop=True)

concurrency_list = sorted(df["concurrency"].unique().tolist())
max_num_seqs_list = sorted(df["max_num_seqs"].unique().tolist())

print("Concurrencies:", concurrency_list)
print("max_num_seqs:", max_num_seqs_list)
print("Rows:", len(df))

```

Concurrencies: [1, 2, 4, 8, 16, 32, 48, 64]
max_num_seqs: [16, 32, 48, 64, 80, 96]
Rows: 48

```

In [ ]: def plot_heatmap(pivot_df, title, xlabel="concurrency", ylabel="max_num_seqs",
                      pivot_df: index=max_num_seqs, columns=concurrency, values=metric,
                      data = pivot_df.values.astype(float))

        plt.figure(figsize=(10, 5.5))
        im = plt.imshow(data, aspect="auto", interpolation="nearest")
        plt.colorbar(im)

        plt.title(title)
        plt.xlabel(xlabel)
        plt.ylabel(ylabel)

```

```

plt.xticks(ticks=np.arange(pivot_df.shape[1]), labels=pivot_df.columns)
plt.yticks(ticks=np.arange(pivot_df.shape[0]), labels=pivot_df.index)

for i in range(pivot_df.shape[0]):
    for j in range(pivot_df.shape[1]):
        val = data[i, j]
        if np.isfinite(val):
            plt.text(j, i, fmt.format(val), ha="center", va="center",
                     color="white" if val > 0 else "black")

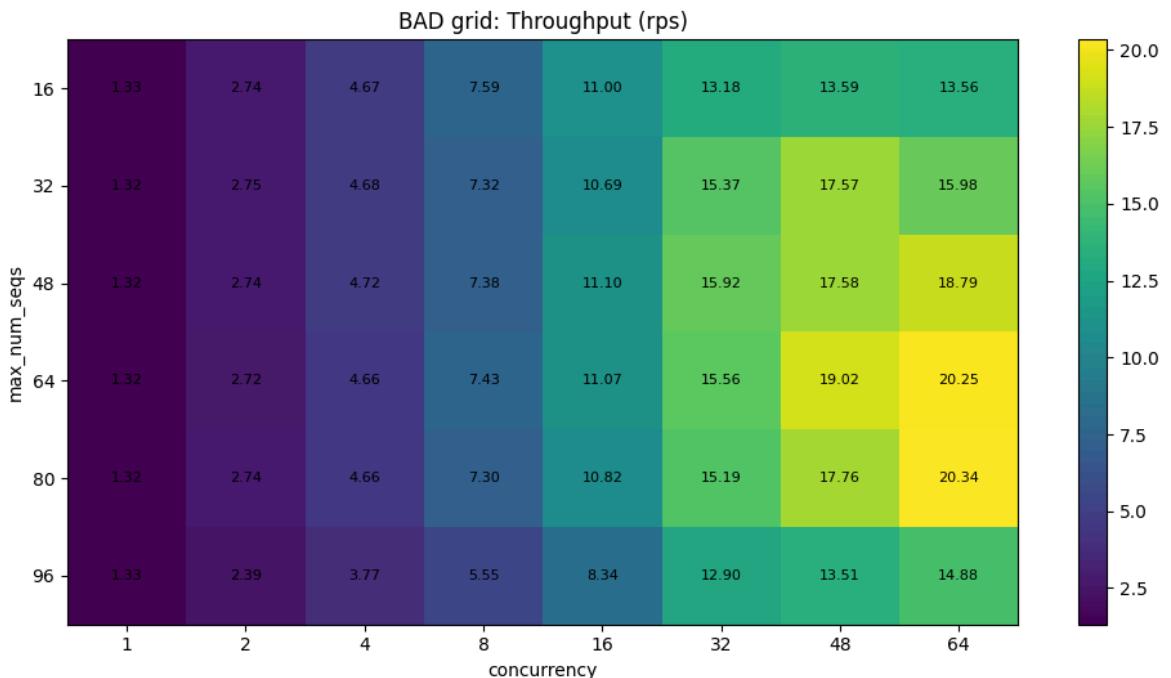
plt.tight_layout()
plt.show()

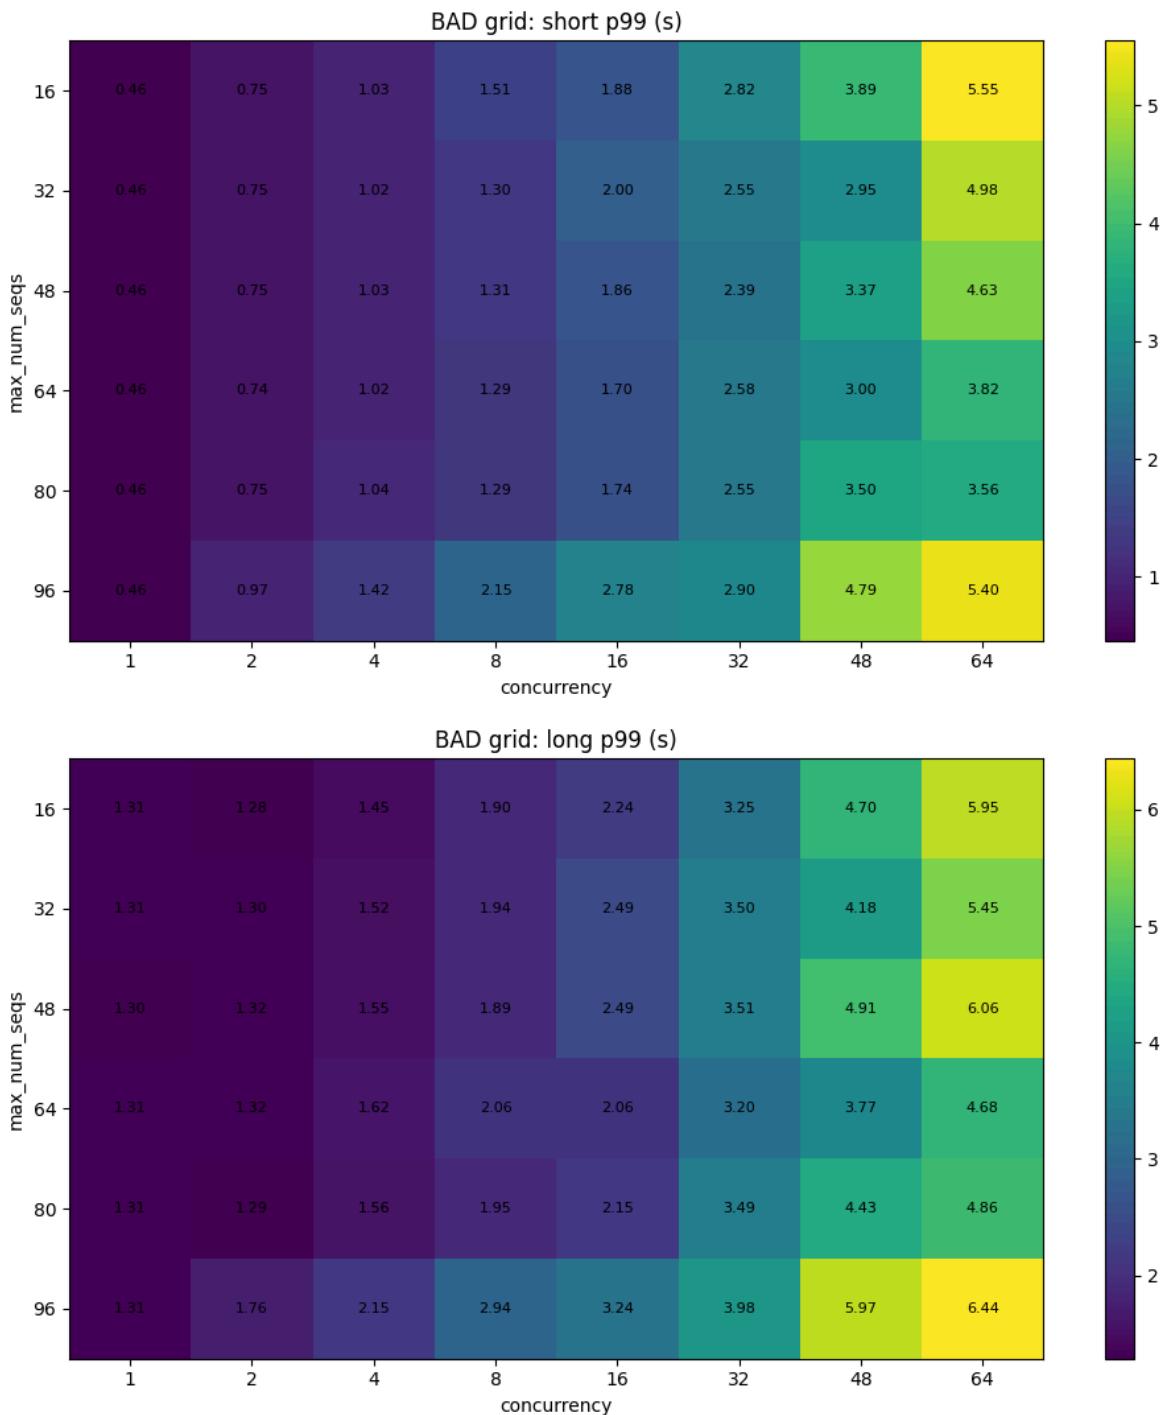
# Heatmap: rps
p_rps = df.pivot_table(index="max_num_seqs", columns="concurrency", value
plot_heatmap(p_rps, "BAD grid: Throughput (rps)")

# Heatmap: short_p99
p_sp99 = df.pivot_table(index="max_num_seqs", columns="concurrency", value
plot_heatmap(p_sp99, "BAD grid: short p99 (s)")

# Heatmap: long_p99
p_lp99 = df.pivot_table(index="max_num_seqs", columns="concurrency", value
plot_heatmap(p_lp99, "BAD grid: long p99 (s)")

```





Heatmap Interpretation

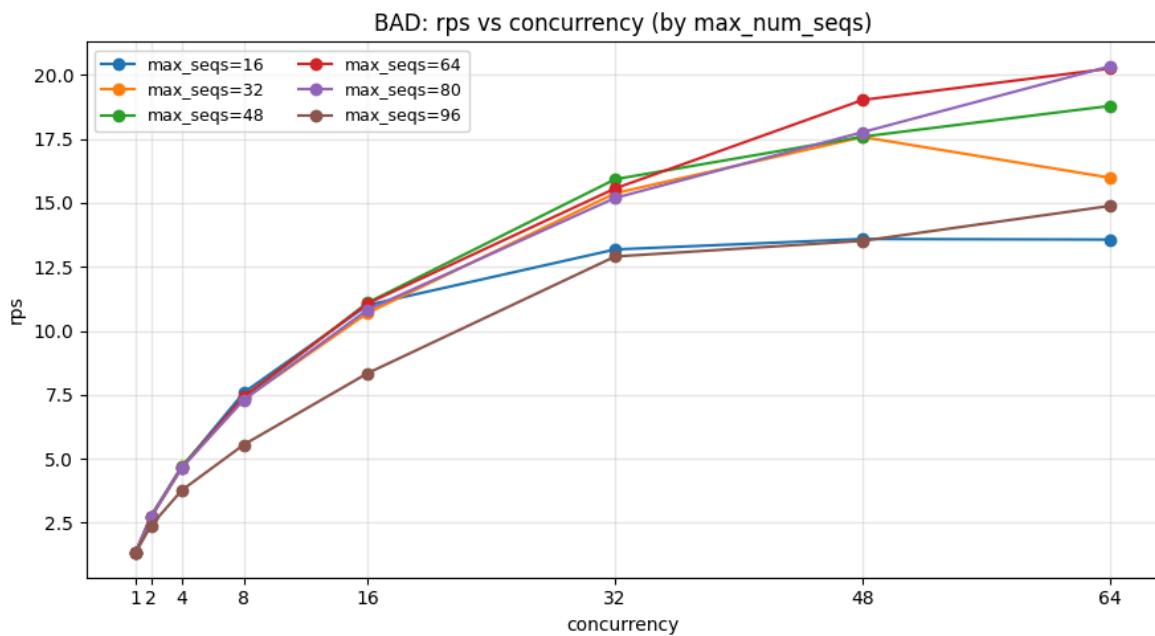
- **Throughput scales with concurrency until ~32–64**, but the gain after ~32 is smaller than the latency cost. You can see the plateau: rps improves from low concurrency to 32, then flattens for many max_num_seqs settings.
- **max_num_seqs acts like a capacity knob for concurrency**. At high concurrency (48–64), small max_num_seqs (like 16) cannot keep up, so both short and long p99 inflate. Mid-range values (around 64–80) sustain higher rps while keeping p99 lower.
- **Too large max_num_seqs can hurt**. The row at max_num_seqs = 96 shows lower throughput and much higher short and long p99 at high concurrency,

consistent with extra batching/queueing overhead and heavier contention when you allow too many in-flight sequences.

- **Short requests still suffer under mixed workloads when scheduling is unfair.**
Even though “short p99” is lower than “long p99”, it rises sharply with concurrency (especially at c=48–64), showing head-of-line blocking where long prefetches delay short requests.
- **Practical “knee” region:** max_num_seqs around **64–80** with concurrency around **16–32** gives a better tradeoff. It keeps rps high without the steep tail-latency blow-up seen at the extreme high concurrency settings.

```
In [ ]: plt.figure(figsize=(9, 5))
for m in max_num_seqs_list:
    sub = df[df["max_num_seqs"] == m].sort_values("concurrency")
    plt.plot(sub["concurrency"], sub["rps"], marker="o", label=f"max_seqs={m}")

plt.title("BAD: rps vs concurrency (by max_num_seqs)")
plt.xlabel("concurrency")
plt.ylabel("rps")
plt.xticks(concurrency_list)
plt.grid(True, alpha=0.3)
plt.legend(ncol=2, fontsize=9)
plt.tight_layout()
plt.show()
```



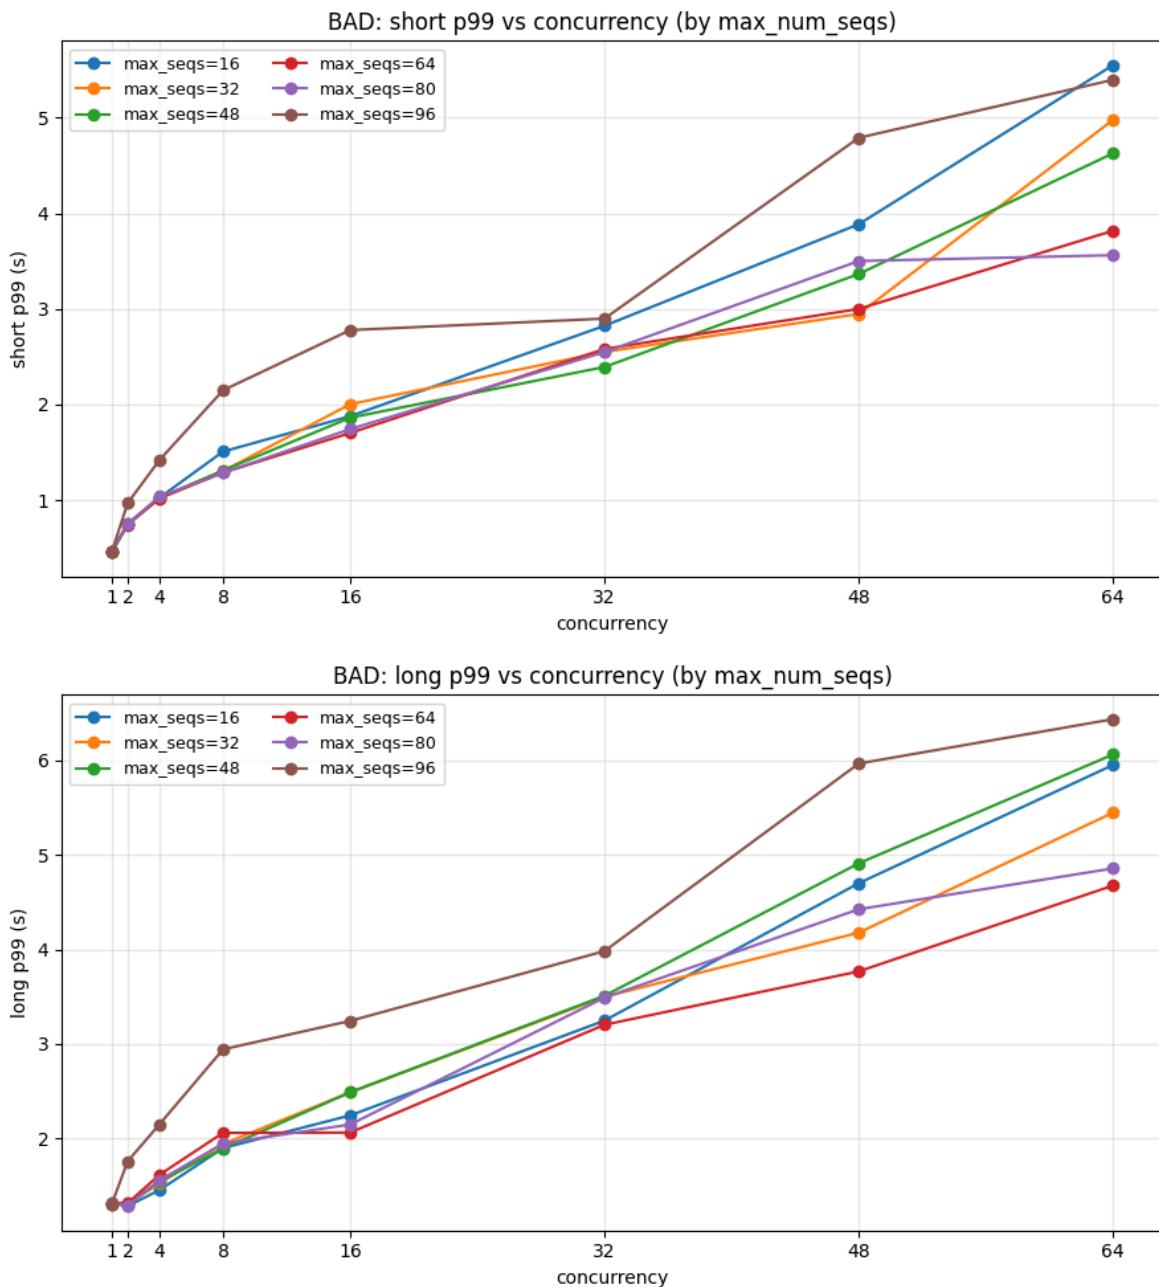
Interpretation: BAD configuration throughput (RPS) vs concurrency, grouped by max_num_seqs

- **RPS increases with concurrency for all settings up to ~16–32**, which is the “scaling region” where the GPU is getting better utilized as we add more in-flight work.
- **max_num_seqs controls how far you can scale before flattening.**

- With **max_num_seqs=16**, the curve **plateaus early** (around concurrency 32) and stays near ~13–14 rps, meaning the server cannot admit enough parallel sequences to benefit from more client concurrency.
- With **max_num_seqs=48/64/80**, the curves keep improving to higher concurrencies and reach the best throughput at the right edge (48–64).
- **The best throughput in this sweep occurs at high concurrency (48–64) with max_num_seqs around 64–80** (top curves near ~20 rps). That suggests this range is closer to the capacity point of this setup.
- **max_num_seqs=96 underperforms at low and mid concurrency and only partially catches up later.** This is consistent with “too many allowed in-flight sequences” adding overhead (more queueing, scheduling contention, memory pressure), so the system wastes efficiency instead of converting concurrency into throughput.
- **Diminishing returns are visible beyond ~48 concurrency**, even for the best max_num_seqs values. The slope flattens, meaning you are approaching saturation, and extra concurrency mostly increases queueing rather than throughput.

```
In [ ]: # short p99
plt.figure(figsize=(9, 5))
for m in max_num_seqs_list:
    sub = df[df["max_num_seqs"] == m].sort_values("concurrency")
    plt.plot(sub["concurrency"], sub["short_p99"], marker="o", label=f"max_{m}_seqs")
plt.title("BAD: short p99 vs concurrency (by max_num_seqs)")
plt.xlabel("concurrency")
plt.ylabel("short p99 (s)")
plt.xticks(concurrency_list)
plt.grid(True, alpha=0.3)
plt.legend(ncol=2, fontsize=9)
plt.tight_layout()
plt.show()

# long p99
plt.figure(figsize=(9, 5))
for m in max_num_seqs_list:
    sub = df[df["max_num_seqs"] == m].sort_values("concurrency")
    plt.plot(sub["concurrency"], sub["long_p99"], marker="o", label=f"max_{m}_seqs")
plt.title("BAD: long p99 vs concurrency (by max_num_seqs)")
plt.xlabel("concurrency")
plt.ylabel("long p99 (s)")
plt.xticks(concurrency_list)
plt.grid(True, alpha=0.3)
plt.legend(ncol=2, fontsize=9)
plt.tight_layout()
plt.show()
```



Interpretation: BAD short p99 vs concurrency and BAD long p99 vs concurrency

- For both short and long requests, **p99 increases monotonically with concurrency**. This is classic queueing behavior: higher offered load increases worst-case wait time.
- Short p99 inflates sharply at high concurrency** ($c=48, c=64$), showing that short requests lose their "fast path" when mixed with long prompts under unfair scheduling.
- Long p99 also climbs steadily**, and it typically climbs faster than short p99 because long prompts contribute more prefill work and hold resources longer.
- max_num_seqs=64 (and sometimes 80) provides a better tradeoff** than 16/32 (too restrictive) and 96 (too aggressive): higher throughput without blowing up p99 as badly as the overcommitted setting.

Experiment 5: Creating Fair Scheduling to improve the latency and RPS (or the trade-off)

Fair sanity check: lower max-num-batched-tokens (btok=768) at c=32

Before running the full FAIR sweeps, I ran a single controlled trial by reducing `max-num-batched-tokens` from the BAD baseline (8192) down to **768**, while keeping everything else identical and evaluating at the **same client concurrency (c=32)** used for the BAD runs.

- **Purpose:** verify that `max-num-batched-tokens` is a meaningful lever under mixed 50/50 traffic, and that lowering it can reduce the “long-prefill domination” effect that inflates tail latency.
- **Why 768:** it’s a mid-range value expected to encourage more frequent scheduling turns (smaller batches), which can improve responsiveness for short requests in mixed workloads.
- **Why c=32:** match the BAD configuration’s stress level so the comparison is apples-to-apples and any improvement is not just because the system is under-loaded.

```
In [ ]: !pkill -f "vllm.entrypoints.openai.api_server" || true
!pkill -f "VLLM::EngineCore" || true
!rm -f vllm_fair.log
!sleep 1
```

```
In [ ]: !nohup python3 -m vllm.entrypoints.openai.api_server \
--model Qwen/Qwen2.5-3B-Instruct \
--host 127.0.0.1 --port 8000 \
--dtype half \
--max-model-len 4096 \
--gpu-memory-utilization 0.70 \
--disable-log-stats \
--enforce-eager \
--max-num-batched-tokens 768 \
--long-prefill-token-threshold 1024 \
> vllm_fair.log 2>&1 </dev/null &
```

^C
^C

```
In [ ]: !tail -n 40 vllm_fair.log
```

```
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:38] Available routes are:  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /openapi.json, Methods: GET, HEAD  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /docs, Methods: GET, HEAD  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /docs/oauth2-redirect, Methods: GET, HEAD  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /redoc, Methods: GET, HEAD  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /scale_elastic_ep, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /is_scaling_elastic_ep, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /tokenize, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /detokenize, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /inference/v1/generate, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /pause, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /resume, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /is_paused, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /metrics, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /health, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /load, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/models, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /version, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/responses, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/responses/{response_id}, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/responses/{response_id}/cancel, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/messages, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/chat/completions, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/completions, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/audio/transcriptions, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/audio/translations, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /ping, Methods: GET  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /ping, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /invocations, Methods: POST  
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /classify, Methods: POST
```

```
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/embeddings, Methods: POST
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /score, Methods: POST
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/score, Methods: POST
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /rerank, Methods: POST
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v1/rerank, Methods: POST
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /v2/rerank, Methods: POST
(APIServer pid=22255) INFO 01-10 15:52:09 [launcher.py:46] Route: /pooling, Methods: POST
(APIServer pid=22255) INFO:      Started server process [22255]
(APIServer pid=22255) INFO:      Waiting for application startup.
(APIServer pid=22255) INFO:      Application startup complete.
```

```
In [ ]: df_fair_50_50, summ_fair_50_50 = run_load(
    prompt_sets["mix_50_50.jsonl"],
    concurrency=32,
    total_requests=200,
    desired_max_tokens=16,
    desc="FAIR_protected | mix_50_50 | c=16"
)

print("FAIR summary:", summ_fair_50_50)
df_fair_50_50.head()
```

FAIR_protected | mix_50_50 | c=16: 0% | 0/200 [00:00<?, ?it/s]
FAIR summary: {'concurrency': 32, 'total_requests': 200, 'ok': 200, 'fail': 0, 'wall_time_s': 41.268839522000235, 'rps': 4.84627148028674, 'p50_s': 6.04282219300012, 'p90_s': 9.18175588419981, 'p99_s': 11.05694531869025, 'mean_s': 6.266493499729964}

```
Out[ ]:   latency_s  prompt_tokens  completion_tokens  total_tokens  status_code  prompt
          0       3.036116            37                  16             53        200
          1       3.129861            37                  16             53        200
          2       3.074753            45                  16             61        200
          3       3.092510            37                  16             53        200
          4       3.141762            38                  16             54        200
```

```
In [ ]: def latency_stats(df, label, mask, subset_name):
    x = df.loc[mask & (df["ok"] == True), "latency_s"].dropna().values
    if len(x) == 0:
        return {"label": label, "subset": subset_name, "n": 0, "p50": None}
    return {
        "label": label,
        "subset": subset_name,
        "n": int(len(x)),
        "p50": float(np.quantile(x, 0.50)),
        "p90": float(np.quantile(x, 0.90)),
        "p99": float(np.quantile(x, 0.99)),
        "mean": float(np.mean(x)),
    }
```

```
def compare_short_long(df, label, short_cut=64, long_cut=2048):
    return [
        latency_stats(df, label, mask=df["prompt_tokens"] >= 0, subset_na=True),
        latency_stats(df, label, mask=df["prompt_tokens"] <= short_cut, subset_na=True),
        latency_stats(df, label, mask=df["prompt_tokens"] >= long_cut, subset_na=True)
    ]

bad_stats = compare_short_long(df_bad, "BAD_unfair", short_cut=64, long_cut=2048)
fair_stats = compare_short_long(df_fair_50_50, "FAIR_768", short_cut=64, long_cut=2048)

bad_stats, fair_stats
```

```
Out[ ]: ([{'label': 'BAD_unfair',
            'subset': 'overall',
            'n': 200,
            'p50': 5.534205468999517,
            'p90': 10.461404818999926,
            'p99': 11.045583048669595,
            'mean': 6.076663721505029},
           {'label': 'BAD_unfair',
            'subset': 'short<= 64',
            'n': 108,
            'p50': 4.6954104774999905,
            'p90': 8.427761869799998,
            'p99': 10.15342957642024,
            'mean': 5.170086654425977},
           {'label': 'BAD_unfair',
            'subset': 'long>= 2048',
            'n': 92,
            'p50': 7.038852010499795,
            'p90': 10.96015420079957,
            'p99': 11.058848744559482,
            'mean': 7.140906365467394}],
          [{'label': 'FAIR_768',
            'subset': 'overall',
            'n': 200,
            'p50': 6.04282219300012,
            'p90': 9.18175588419981,
            'p99': 11.05694531869025,
            'mean': 6.266493499729964},
           {'label': 'FAIR_768',
            'subset': 'short<= 64',
            'n': 108,
            'p50': 5.629765399999542,
            'p90': 8.433177871999579,
            'p99': 10.081671103819572,
            'mean': 5.578912106185119},
           {'label': 'FAIR_768',
            'subset': 'long>= 2048',
            'n': 92,
            'p50': 7.1285672229996635,
            'p90': 9.995110139300506,
            'p99': 11.105451201459774,
            'mean': 7.073654266065215}])
```

Conclusion: btok = 768 and c = 32

- **Overall:** FAIR_768 slightly **hurts median** (p50: 6.04s vs 5.53s) but **improves p90** (9.18s vs 10.46s). **p99 is basically unchanged** (~11.06s in both).
- **Short requests (≤ 64 tokens):** FAIR_768 **slows p50** (5.63s vs 4.70s) but gives a **small p99 improvement** (10.08s vs 10.15s). Net: better tail for short, worse typical latency.
- **Long requests (≥ 2048 tokens):** FAIR_768 **improves p90** (10.00s vs 10.96s) while **p99 stays ~the same** (11.11s vs 11.06s). This suggests btok=768 reduces “most long requests” queueing, but the worst-case is still dominated by contention.
- **Implication:** Lowering `max-num-batched-tokens` to 768 shifts the distribution (especially **p90**) but doesn’t meaningfully reduce the **worst-case p99** at this concurrency—so a sweep is needed to find the btok “sweet spot” that improves tails without sacrificing too much throughput.

max-num-batched-token Sweep @ c = 16

```
In [ ]: BASE = "http://127.0.0.1:8000"
MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

def kill_server():
    os.system('pkill -f "vllm.entrypoints.openai.api_server" || true')
    os.system('pkill -f "VLLM::EngineCore" || true')
    time.sleep(1.5)

def start_fair_server(max_btok, log_path="vllm_fair.log"):
    # clean + restart
    kill_server()
    try:
        os.remove(log_path)
    except FileNotFoundError:
        pass

    cmd = f"""
nohup python3 -m vllm.entrypoints.openai.api_server \
--model {MODEL_ID} \
--host 127.0.0.1 --port 8000 \
--dtype half \
--max-model-len 4096 \
--gpu-memory-utilization 0.70 \
--disable-log-stats \
--enforce-eager \
--max-num-batched-tokens {int(max_btok)} \
--long-prefill-token-threshold 1024 \
> {log_path} 2>&1 </dev/null &
"""
    os.system(cmd)
    time.sleep(1.0)

def tail_log(log_path, n=60):
    if os.path.exists(log_path):
        print(f"\n--- tail {log_path} (last {n}) ---")
        os.system(f"tail -n {n} {log_path}")
    else:
        print(f"(log missing: {log_path})")
```

```

def wait_for_ready(base=BASE, log_path="vllm_fair.log", timeout_s=300):
    t0 = time.time()
    tries = 0
    while True:
        tries += 1
        try:
            r = requests.get(f"{base}/v1/models", timeout=2)
            if r.status_code == 200:
                data = r.json().get("data", [])
                ids = [m.get("id") for m in data]
                print(f"✅ Server ready. Models: {ids}")
                return True
            else:
                print(f"[{tries}] /v1/models status={r.status_code}")
        except Exception as e:
            print(f"[{tries}] not ready yet: {type(e).__name__}")

        # if process died, fail fast
        alive = os.system('ps -ef | grep -E "vllm.entrypoints.openai.api_')
        if not alive:
            print("❌ vLLM process is NOT running anymore.")
            tail_log(log_path, n=120)
            return False

        if time.time() - t0 > timeout_s:
            print("❌ Timed out waiting for server.")
            tail_log(log_path, n=120)
            return False

        if tries % 5 == 0:
            tail_log(log_path, n=30)

        time.sleep(2)

def warmup_request():
    # forces one small request to populate caches
    payload = {
        "model": MODEL_ID,
        "messages": [{"role": "user", "content": "hi"}],
        "max_tokens": 8,
        "temperature": 0.0
    }
    try:
        _ = requests.post(f"{BASE}/v1/chat/completions", json=payload, ti
    except Exception:
        pass

# Sweep runner

WL_FILE = "mix_50_50.jsonl"
CONCURRENCY = 16
TOTAL_REQ = 400
DESIRED_MAX_TOKENS = 16

btok_list = [256, 384, 512, 768, 1024, 1536, 2048] # edit as you like
rows = []
fair_runs = []

```

```

for btok in btok_list:
    print(f"\n====")
    print(f"Starting FAIR server: max-num-batched-tokens={btok}")
    print(f"====")

    start_fair_server(btok, log_path="vllm_fair.log")

    ok = wait_for_ready(BASE, log_path="vllm_fair.log", timeout_s=420)
    if not ok:
        rows.append({
            "label": "FAIR",
            "workload": WL_FILE,
            "concurrency": CONCURRENCY,
            "max_num_batched_tokens": btok,
            "status": "server_failed"
        })
        continue

    warmup_request()

    df_fair, summ_fair = run_load(
        prompt_sets[WL_FILE],
        concurrency=CONCURRENCY,
        total_requests=TOTAL_REQ,
        desired_max_tokens=DESIRED_MAX_TOKENS,
        desc=f"FAIR | {WL_FILE} | c={CONCURRENCY} | btok={btok}"
    )

    fair_runs[btok] = (df_fair, summ_fair)

    m = subset_metrics(df_fair)

    rows.append({
        "label": "FAIR",
        "workload": WL_FILE,
        "concurrency": CONCURRENCY,
        "max_num_batched_tokens": btok,
        "status": "ok",
        "ok_count": int((df_fair["ok"] == True).sum()),
        "fail_count": int((df_fair["ok"] != True).sum()),
        "p50_s": float(df_fair[df_fair["ok"]==True]["latency_s"].quantile(0.5)),
        "p90_s": float(df_fair[df_fair["ok"]==True]["latency_s"].quantile(0.9)),
        "p99_s": float(df_fair[df_fair["ok"]==True]["latency_s"].quantile(0.99)),
        "mean_s": float(df_fair[df_fair["ok"]==True]["latency_s"].mean())
    }**m
    )

kill_server()

fair_summary_df = pd.DataFrame(rows)
print("\n✓ FAIR btok sweep done.")
fair_summary_df

```

```
=====
Starting FAIR server: max-num-batched-tokens=256
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=256: 0% | 0/400 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=384
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=384:    0%|          | 0/400 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=512
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=512: 0% | 0/400 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=768
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=768: 0% | 0/400 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=1024
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=1024: 0%| 0/400 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=1536
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=1536: 0%| 0/400 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=2048
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=16 | btok=2048:   0%|          | 0/400 [00:00
<?, ?it/s]
✓ FAIR btok sweep done.
```

Out[]:

	label	workload	concurrency	max_num_batched_tokens	status	ok_count
0	FAIR	mix_50_50.jsonl	16	256	ok	400
1	FAIR	mix_50_50.jsonl	16	384	ok	400
2	FAIR	mix_50_50.jsonl	16	512	ok	400
3	FAIR	mix_50_50.jsonl	16	768	ok	400
4	FAIR	mix_50_50.jsonl	16	1024	ok	400
5	FAIR	mix_50_50.jsonl	16	1536	ok	400
6	FAIR	mix_50_50.jsonl	16	2048	ok	400

7 rows × 24 columns

In []:

```
import os, time, requests
import pandas as pd
import numpy as np

BASE = "http://127.0.0.1:8000"
```

```

MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

def subset_metrics(df, short_cutoff=64, long_cutoff=3500):
    ok = df[df["ok"] == True].copy()

    def _q(s, p):
        if len(s) == 0:
            return None
        return float(s.quantile(p))

    overall = {
        "overall_n": int(len(ok)),
        "overall_p50": _q(ok["latency_s"], 0.50),
        "overall_p90": _q(ok["latency_s"], 0.90),
        "overall_p99": _q(ok["latency_s"], 0.99),
        "overall_mean": float(ok["latency_s"].mean()) if len(ok) else None
    }

    short = ok[ok["prompt_tokens"] <= short_cutoff]
    short_stats = {
        "short_n": int(len(short)),
        "short_p50": _q(short["latency_s"], 0.50),
        "short_p90": _q(short["latency_s"], 0.90),
        "short_p99": _q(short["latency_s"], 0.99),
        "short_mean": float(short["latency_s"].mean()) if len(short) else None
    }

    long_ = ok[ok["prompt_tokens"] >= long_cutoff]
    long_stats = {
        "long_n": int(len(long_)),
        "long_p50": _q(long_["latency_s"], 0.50),
        "long_p90": _q(long_["latency_s"], 0.90),
        "long_p99": _q(long_["latency_s"], 0.99),
        "long_mean": float(long_["latency_s"].mean()) if len(long_) else None
    }

    out = {}
    out.update(overall)
    out.update(short_stats)
    out.update(long_stats)
    return out

def kill_server():
    os.system('pkill -f "vllm.entrypoints.openai.api_server" || true')
    os.system('pkill -f "VLLM::EngineCore" || true')
    time.sleep(1.5)

def start_fair_server(max_btok, log_path="vllm_fair.log"):
    kill_server()
    try:
        os.remove(log_path)
    except FileNotFoundError:
        pass

    cmd = f"""
nohup python3 -m vllm.entrypoints.openai.api_server \
--model {MODEL_ID} \
"""

```

```

--host 127.0.0.1 --port 8000 \
--dtype half \
--max-model-len 4096 \
--gpu-memory-utilization 0.70 \
--disable-log-stats \
--enforce-eager \
--max-num-batched-tokens {int(max_btok)} \
--long-prefill-token-threshold 1024 \
> {log_path} 2>&1 </dev/null &
"""
os.system(cmd)
time.sleep(1.0)

def tail_log(log_path, n=60):
    if os.path.exists(log_path):
        print(f"\n--- tail {log_path} (last {n}) ---")
        os.system(f"tail -n {n} {log_path}")

def wait_for_ready(base=BASE, log_path="vllm_fair.log", timeout_s=420):
    t0 = time.time()
    tries = 0
    while True:
        tries += 1
        try:
            r = requests.get(f"{base}/v1/models", timeout=2)
            if r.status_code == 200:
                data = r.json().get("data", [])
                ids = [m.get("id") for m in data]
                print(f"✓ Server ready. Models: {ids}")
                return True
        except Exception as e:
            print(f"[{tries}] not ready yet: {type(e).__name__}")

    alive = os.system('ps -ef | grep -E "vllm.entrypoints.openai.api')
    if not alive:
        print("✗ vLLM process is NOT running anymore.")
        tail_log(log_path, n=120)
        return False

    if time.time() - t0 > timeout_s:
        print("✗ Timed out waiting for server.")
        tail_log(log_path, n=120)
        return False

    if tries % 5 == 0:
        tail_log(log_path, n=30)

    time.sleep(2)

def warmup_request():
    payload = {
        "model": MODEL_ID,
        "messages": [{"role": "user", "content": "hi"}],
        "max_tokens": 8,
        "temperature": 0.0
    }
    try:
        _ = requests.post(f"{BASE}/v1/chat/completions", json=payload, ti
    except Exception:
        pass

```

```

WL_FILE = "mix_50_50.jsonl"
CONCURRENCY = 32
TOTAL_REQ = 200
DESIRED_MAX_TOKENS = 16

btok_list = [256, 384, 512, 768, 1024, 1536, 2048, 3072, 4096, 6144, 8192

rows = []
fair_runs = {}

for btok in btok_list:
    print(f"\n====")
    print(f"Starting FAIR server: max-num-batched-tokens={btok}")
    print(f"====")

    start_fair_server(btok, log_path="vllm_fair.log")

    ok = wait_for_ready(BASE, log_path="vllm_fair.log", timeout_s=420)
    if not ok:
        rows.append({
            "label": "FAIR",
            "workload": WL_FILE,
            "concurrency": CONCURRENCY,
            "max_num_batched_tokens": btok,
            "status": "server_failed"
        })
        continue

    warmup_request()

    df_fair, summ_fair = run_load(
        prompt_sets[WL_FILE],
        concurrency=CONCURRENCY,
        total_requests=TOTAL_REQ,
        desired_max_tokens=DESIRED_MAX_TOKENS,
        desc=f"FAIR | {WL_FILE} | c={CONCURRENCY} | btok={btok}"
    )

    fair_runs[btok] = (df_fair, summ_fair)

    # subset metrics + safe latency stats
    m = subset_metrics(df_fair, short_cutoff=64, long_cutoff=3500)
    ok_df = df_fair[df_fair["ok"] == True]

    # rps/wall_time from summary (safe if missing)
    rps = summ_fair.get("rps", None)
    wall = summ_fair.get("wall_time_s", None)

    rows.append({
        "label": "FAIR",
        "workload": WL_FILE,
        "concurrency": CONCURRENCY,
        "max_num_batched_tokens": btok,
        "status": "ok",
    })

```

```

        "ok_count": int((df_fair["ok"] == True).sum()),
        "fail_count": int((df_fair["ok"] != True).sum()),

        # added throughput fields
        "rps": float(rps) if rps is not None else None,
        "wall_time_s": float(wall) if wall is not None else None,

        "p50_s": float(ok_df["latency_s"].quantile(0.50)) if len(ok_df) > 0 else None,
        "p90_s": float(ok_df["latency_s"].quantile(0.90)) if len(ok_df) > 0 else None,
        "p99_s": float(ok_df["latency_s"].quantile(0.99)) if len(ok_df) > 0 else None,
        "mean_s": float(ok_df["latency_s"].mean()) if len(ok_df) > 0 else None

    **m
}

kill_server()

fair_summary_df = pd.DataFrame(rows)
print("\n✓ FAIR btok sweep done.")
fair_summary_df
=====

Starting FAIR server: max-num-batched-tokens=256
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=256: 0% | 0/200 [00:00<?, ?it/s]

```

```
=====
Starting FAIR server: max-num-batched-tokens=384
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=384: 0% | 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=512
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
[25] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[26] not ready yet: ConnectionError
[27] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512: 0% | 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=768
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=768: 0% | 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=1024
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1024: 0%| 0/200 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=1536
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536: 0%| 0/200 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=2048
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048: 0%| 0/200 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=3072
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072: 0%| 0/200 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=4096
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096: 0%| 0/200 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=6144
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144: 0%| 0/200 [00:00
<?, ?it/s]
```

```
=====
Starting FAIR server: max-num-batched-tokens=8192
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192: 0%| 0/200 [00:00
<?, ?it/s]
 FAIR btok sweep done.
```

Out[]:	label	workload	concurrency	max_num_batched_tokens	status	ok_count
0	FAIR	mix_50_50.jsonl	32	256	ok	20
1	FAIR	mix_50_50.jsonl	32	384	ok	20
2	FAIR	mix_50_50.jsonl	32	512	ok	20
3	FAIR	mix_50_50.jsonl	32	768	ok	20
4	FAIR	mix_50_50.jsonl	32	1024	ok	20
5	FAIR	mix_50_50.jsonl	32	1536	ok	20
6	FAIR	mix_50_50.jsonl	32	2048	ok	20
7	FAIR	mix_50_50.jsonl	32	3072	ok	20
8	FAIR	mix_50_50.jsonl	32	4096	ok	20
9	FAIR	mix_50_50.jsonl	32	6144	ok	20
10	FAIR	mix_50_50.jsonl	32	8192	ok	20

11 rows × 28 columns

Conclusion: btok Sweep @ c= 16

- Even at lower concurrency (FAIR sweep at `c=16`), FAIR is dramatically better than the BAD baseline you saw (BAD at `c=32`).

Not an apples-to-apples comparison due to different `concurrency`, but it's still a strong signal that the FAIR knob is addressing the bottleneck.

 - Overall **p99: 11.05s (BAD) → 5.88s (FAIR @ btok=2048)** ($\approx 46.8\%$ lower)
 - Overall **p90: 10.46s (BAD) → 4.88s (FAIR @ btok=2048)** ($\approx 53.3\%$ lower)
 - Overall **p50: 5.53s (BAD) → 2.88s (FAIR @ btok=2048)** ($\approx 47.9\%$ lower)
 - Overall **mean: 6.08s (BAD) → 3.09s (FAIR @ btok=2048)** ($\approx 49.2\%$ lower)
- Short and long tails both improve vs BAD (still with the same cross-concurrency caveat).
 - Short **p99: 10.15s (BAD) → 5.05s (FAIR @ btok=2048)** ($\approx 50.2\%$ lower)
 - Long **p99: 11.06s (BAD) → 5.97s (FAIR @ btok=2048)** ($\approx 46.0\%$ lower)
- Within the FAIR sweep itself (`c=16`), tail latency improves as `max-num-batched-tokens` increases (`256 → 2048`).
 - Overall **p99: 7.49s → 5.88s** ($\approx 21.5\%$ lower)
 - Long **p99: 7.64s → 5.97s** ($\approx 21.9\%$ lower)
 - This strongly suggests the server was **prefill-throttled** at low btok.
- Not just tail — median improves too (FAIR `c=16`, `256 → 2048`).
 - Overall **p50: 3.92s → 2.88s** ($\approx 26.6\%$ lower)
 - Overall **p90: 6.00s → 4.88s** ($\approx 18.7\%$ lower)

Caveat: concurrency confounds the BAD vs FAIR comparison

- The BAD baseline numbers were measured at `concurrency=32`, while this FAIR btok sweep was run at `concurrency=16`.
- Lower concurrency reduces queueing pressure by itself, so part of the latency improvement could be explained by **less load**, not only by better batching/scheduling.
- To isolate the effect of server knobs, we need an apples-to-apples sweep at **the same concurrency**.

What to do next (to remove the confound)

- Re-run the btok sweep at `concurrency=32` (same workload and request count) to validate whether the btok trend holds at the real comparison point.
- Then run a **2D grid sweep** at `c=32` over:
 - `max-num-batched-tokens` (controls prefill batching / throughput)
 - `max-num-seqs` (controls how many sequences can be in-flight; too low underutilizes, too high overcommits)
- This will let us attribute changes correctly:
 - If FAIR still wins at `c=32`, the gains are genuinely from scheduling/batching.
 - If FAIR only wins at lower `c`, the improvement is mostly due to reduced load rather than config.

```
In [ ]: output_path = os.path.join(DRIVE_OUT_DIR, "fair_scheduling_sweep_summary")
fair_summary_df.to_csv(output_path, index=False)
print(f"✅ Saved fair_summary_df to Drive: {output_path}")

✅ Saved fair_summary_df to Drive: /content/drive/MyDrive/llm_loadtest_wor
kloads/fair_scheduling_sweep_summary.csv
```

```
In [ ]: import os, time, json, requests
import pandas as pd

BASE = "http://127.0.0.1:8000"
MODEL_ID = "Qwen/Qwen2.5-3B-Instruct"

CHECKPOINT_CSV = "fair_grid_checkpoint.csv"
MANIFEST_JSON = "fair_grid_manifest.json"

def kill_server():
    os.system('pkill -f "vllm.entrypoints.openai.api_server" || true')
    os.system('pkill -f "VLLM::EngineCore" || true')
    time.sleep(1.5)

def tail_log(log_path, n=80):
    if os.path.exists(log_path):
        print(f"\n--- tail {log_path} (last {n}) ---")
        os.system(f"tail -n {n} {log_path}")
    else:
```

```

        print(f"(log missing: {log_path})")

def wait_for_ready(base=BASE, log_path="vllm_fair.log", timeout_s=480):
    t0 = time.time()
    tries = 0
    while True:
        tries += 1
        try:
            r = requests.get(f"{base}/v1/models", timeout=2)
            if r.status_code == 200:
                data = r.json().get("data", [])
                ids = [m.get("id") for m in data]
                print(f"\u2713 Server ready. Models: {ids}")
                return True
            else:
                print(f"[{tries}] /v1/models status={r.status_code}")
        except Exception:
            print(f"[{tries}] not ready yet: ConnectionError")

    alive = os.system('ps -ef | grep -E "vllm.entrypoints.openai.api_')
    if not alive:
        print("X vLLM process is NOT running anymore.")
        tail_log(log_path, n=160)
        return False

    if time.time() - t0 > timeout_s:
        print("X Timed out waiting for server.")
        tail_log(log_path, n=160)
        return False

    if tries % 5 == 0:
        tail_log(log_path, n=30)

    time.sleep(2)

def warmup_request():
    payload = {
        "model": MODEL_ID,
        "messages": [{"role": "user", "content": "hi"}],
        "max_tokens": 8,
        "temperature": 0.0,
    }
    try:
        _ = requests.post(f"{BASE}/v1/chat/completions", json=payload, ti
    except Exception:
        pass

def start_fair_server(max_btok, max_num_seqs, log_path="vllm_fair.log"):
    kill_server()
    try:
        os.remove(log_path)
    except FileNotFoundError:
        pass

    cmd = f"""
nohup python3 -m vllm.entrypoints.openai.api_server \
--model {MODEL_ID} \
--host 127.0.0.1 --port 8000 \
--dtype half \
--max-model-len 4096 \
"""

```

```
--gpu-memory-utilization 0.70 \
--disable-log-stats \
--enforce-eager \
--max-num-batched-tokens {int(max_btok)} \
--max-num-seqs {int(max_num_seqs)} \
--long-prefill-token-threshold 1024 \
> {log_path} 2>&1 </dev/null &
"""
os.system(cmd)
time.sleep(1.0)

def _safe_quantile(series, q):
    if series is None or len(series) == 0:
        return None
    return float(series.quantile(q))

def _safe_mean(series):
    if series is None or len(series) == 0:
        return None
    return float(series.mean())

def subset_metrics(df, short_cutoff_tokens=64, long_cutoff_tokens=2048):
    ok = df[df["ok"] == True].copy()
    lat = ok["latency_s"]

    out = {
        "overall_n": int(len(ok)),
        "overall_p50": _safe_quantile(lat, 0.50),
        "overall_p90": _safe_quantile(lat, 0.90),
        "overall_p99": _safe_quantile(lat, 0.99),
        "overall_mean": _safe_mean(lat),
    }

    short = ok[ok["prompt_tokens"] <= short_cutoff_tokens]
    out.update({
        "short_n": int(len(short)),
        "short_p50": _safe_quantile(short["latency_s"], 0.50),
        "short_p90": _safe_quantile(short["latency_s"], 0.90),
        "short_p99": _safe_quantile(short["latency_s"], 0.99),
        "short_mean": _safe_mean(short["latency_s"]),
    })

    long = ok[ok["prompt_tokens"] >= long_cutoff_tokens]
    out.update({
        "long_n": int(len(long)),
        "long_p50": _safe_quantile(long["latency_s"], 0.50),
        "long_p90": _safe_quantile(long["latency_s"], 0.90),
        "long_p99": _safe_quantile(long["latency_s"], 0.99),
        "long_mean": _safe_mean(long["latency_s"]),
    })
    return out

def row_from_run(label, workload, concurrency, max_btok, max_num_seqs, df
m = subset_metrics(df, short_cutoff_tokens=64, long_cutoff_tokens=204
return {
    "label": label,
    "workload": workload,
    "concurrency": int(concurrency),
```

```

        "max_num_batched_tokens": int(max_btok),
        "max_num_seqs": int(max_num_seqs),

        "status": "ok",
        "ok": int(summ.get("ok")) if summ.get("ok") is not None else int(
        "fail": int(summ.get("fail")) if summ.get("fail") is not None else
        "rps": float(summ.get("rps")) if summ.get("rps") is not None else
        "wall_time_s": float(summ.get("wall_time_s")) if summ.get("wall_t
        "summ_p50_s": float(summ.get("p50_s")) if summ.get("p50_s") is no
        "summ_p90_s": float(summ.get("p90_s")) if summ.get("p90_s") is no
        "summ_p99_s": float(summ.get("p99_s")) if summ.get("p99_s") is no
        "summ_mean_s": float(summ.get("mean_s")) if summ.get("mean_s") is

        **m,
    }

}

def write_manifest(payload, path=MANIFEST_JSON):
    if os.path.exists(path):
        try:
            old = json.loads(open(path, "r").read())
            if old != payload:
                print("⚠️ Manifest differs from existing one. You may be
                print("Existing:", old)
                print("New:", payload)
        except Exception:
            print("⚠️ Could not read existing manifest. Proceeding.")
        return
    with open(path, "w") as f:
        json.dump(payload, f, indent=2)

def load_checkpoint_df(path=CHECKPOINT_CSV):
    if os.path.exists(path):
        df = pd.read_csv(path)
        for col in ["concurrency", "max_num_batched_tokens", "max_num_seq
            if col in df.columns:
                df[col] = df[col].astype(int)
        return df
    return pd.DataFrame()

def completed_pairs(df_ckpt):
    if df_ckpt.empty:
        return set()
    done = df_ckpt[df_ckpt["status"] == "ok"][[ "max_num_batched_tokens",
    return set((int(r.max_num_batched_tokens), int(r.max_num_seqs))) for r

def append_row_checkpoint(row, path=CHECKPOINT_CSV):
    df_row = pd.DataFrame([row])
    header = not os.path.exists(path)
    df_row.to_csv(path, mode="a", header=header, index=False)

# Sweep settings
WL_FILE = "mix_50_50.jsonl"
CONCURRENCY = 32
TOTAL_REQ = 200
DESIRED_MAX_TOKENS = 16

btok_list = [512, 768, 1024, 1536, 2048, 3072, 4096, 6144, 8192]

```

```

max_seqs_list = [32, 48, 64, 80, 96, 128]

write_manifest({
    "workload": WL_FILE,
    "concurrency": CONCURRENCY,
    "total_requests": TOTAL_REQ,
    "desired_max_tokens": DESIRED_MAX_TOKENS,
    "btok_list": btok_list,
    "max_seqs_list": max_seqs_list,
    "model_id": MODEL_ID,
    "server_flags_fixed": {
        "long_prefill_token_threshold": 1024,
        "gpu_memory_utilization": 0.70,
        "max_model_len": 4096,
        "dtype": "half",
        "enforce_eager": True
    }
})

ckpt_df = load_checkpoint_df(CHECKPOINT_CSV)
done = completed_pairs(ckpt_df)
print(f"✅ Loaded checkpoint rows: {len(ckpt_df)} | completed runs: {len(ckpt_df)}")

# RUN GRID (resume-capable)
for max_num_seqs in max_seqs_list:
    for btok in btok_list:
        if (btok, max_num_seqs) in done:
            print(f"⏭ Skip (already done): btok={btok}, max_seqs={max_num_seqs}")
            continue

        print("\n=====")
        print(f"Starting FAIR server | btok={btok} | max-num-seqs={max_num_seqs}")
        print("=====")

        start_fair_server(max_btok=btok, max_num_seqs=max_num_seqs, log_path="vllm_fair.log")

        ok = wait_for_ready(BASE, log_path="vllm_fair.log", timeout_s=480)
        if not ok:
            row = {
                "label": "FAIR",
                "workload": WL_FILE,
                "concurrency": CONCURRENCY,
                "max_num_batched_tokens": btok,
                "max_num_seqs": max_num_seqs,
                "status": "server_failed",
            }
            append_row_checkpoint(row, CHECKPOINT_CSV)
            ckpt_df = load_checkpoint_df(CHECKPOINT_CSV)
            done = completed_pairs(ckpt_df)
            continue

        warmup_request()

        try:
            df_fair, summ_fair = run_load(
                prompt_sets[WL_FILE],
                concurrency=CONCURRENCY,
                total_requests=TOTAL_REQ,
            )
        
```

```

        desired_max_tokens=DESIRED_MAX_TOKENS,
        desc=f"FAIR | {WL_FILE} | c={CONCURRENCY} | btok={btok} | "
    )

    row = row_from_run(
        label="FAIR",
        workload=WL_FILE,
        concurrency=CONCURRENCY,
        max_btok=btok,
        max_num_seqs=max_num_seqs,
        df=df_fair,
        summ=summ_fair,
    )
    append_row_checkpoint(row, CHECKPOINT_CSV)

except Exception as e:
    row = {
        "label": "FAIR",
        "workload": WL_FILE,
        "concurrency": CONCURRENCY,
        "max_num_batched_tokens": btok,
        "max_num_seqs": max_num_seqs,
        "status": "run_failed",
        "error": type(e).__name__,
    }
    append_row_checkpoint(row, CHECKPOINT_CSV)
    print("X run_load failed:", type(e).__name__)
    tail_log("vllm_fair.log", n=120)

ckpt_df = load_checkpoint_df(CHECKPOINT_CSV)
done = completed_pairs(ckpt_df)

kill_server()

fair_grid_df = load_checkpoint_df(CHECKPOINT_CSV)
print("\n✓ FAIR (btok x max-num-seqs) sweep done (resume-capable).")
fair_grid_df

```

⚠ Manifest differs from existing one. You may be resuming with different sweep params.

Existing: {'workload': 'mix_50_50.jsonl', 'concurrency': 32, 'total_requests': 200, 'desired_max_tokens': 16, 'btok_list': [768, 1024, 1536, 2048, 3072, 4096, 6144, 8192], 'max_seqs_list': [32, 48, 64, 80, 96, 128], 'model_id': 'Qwen/Qwen2.5-3B-Instruct', 'server_flags_fixed': {'long_prefill_token_threshold': 1024, 'gpu_memory_utilization': 0.7, 'max_model_len': 4096, 'dtype': 'half', 'enforce_eager': True}}

New: {'workload': 'mix_50_50.jsonl', 'concurrency': 32, 'total_requests': 200, 'desired_max_tokens': 16, 'btok_list': [512, 768, 1024, 1536, 2048, 3072, 4096, 6144, 8192], 'max_seqs_list': [32, 48, 64, 80, 96, 128], 'model_id': 'Qwen/Qwen2.5-3B-Instruct', 'server_flags_fixed': {'long_prefill_token_threshold': 1024, 'gpu_memory_utilization': 0.7, 'max_model_len': 4096, 'dtype': 'half', 'enforce_eager': True}}

✓ Loaded checkpoint rows: 0 | completed runs: 0

```
=====
Starting FAIR server | btok=512 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512 | max_seqs=32: 0% | 0/200 [00:00<?, ?it/s]
=====

Starting FAIR server | btok=768 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
```

```
FAIR | mix_50_50.jsonl | c=32 | btok=768 | max_seqs=32: 0%|  
0/200 [00:00<?, ?it/s]  
=====  
Starting FAIR server | btok=1024 | max-num-seqs=32  
=====  
[1] not ready yet: ConnectionError  
[2] not ready yet: ConnectionError  
[3] not ready yet: ConnectionError  
[4] not ready yet: ConnectionError  
[5] not ready yet: ConnectionError  
  
--- tail vllm_fair.log (last 30) ---  
[6] not ready yet: ConnectionError  
[7] not ready yet: ConnectionError  
[8] not ready yet: ConnectionError  
[9] not ready yet: ConnectionError  
[10] not ready yet: ConnectionError  
  
--- tail vllm_fair.log (last 30) ---  
[11] not ready yet: ConnectionError  
[12] not ready yet: ConnectionError  
[13] not ready yet: ConnectionError  
[14] not ready yet: ConnectionError  
[15] not ready yet: ConnectionError  
  
--- tail vllm_fair.log (last 30) ---  
[16] not ready yet: ConnectionError  
[17] not ready yet: ConnectionError  
[18] not ready yet: ConnectionError  
[19] not ready yet: ConnectionError  
[20] not ready yet: ConnectionError  
  
--- tail vllm_fair.log (last 30) ---  
[21] not ready yet: ConnectionError  
[22] not ready yet: ConnectionError  
[23] not ready yet: ConnectionError  
✓ Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']  
FAIR | mix_50_50.jsonl | c=32 | btok=1024 | max_seqs=32: 0%|  
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1536 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536 | max_seqs=32: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=2048 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048 | max_seqs=32: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=3072 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072 | max_seqs=32: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=4096 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096 | max_seqs=32: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=6144 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144 | max_seqs=32: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=8192 | max-num-seqs=32
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192 | max_seqs=32: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=512 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=768 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=768 | max_seqs=48: 0% |
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1024 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1024 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1536 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=2048 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=3072 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=4096 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=6144 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=8192 | max-num-seqs=48
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192 | max_seqs=48: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=512 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=768 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=768 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1024 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1024 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1536 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=2048 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=3072 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=4096 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=6144 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=8192 | max-num-seqs=64
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192 | max_seqs=64: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=512 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=768 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=768 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1024 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1024 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1536 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=2048 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=3072 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=4096 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=6144 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=8192 | max-num-seqs=80
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192 | max_seqs=80: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=512 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512 | max_seqs=96:    0% |
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=768 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=768 | max_seqs=96:    0% |
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1024 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1024 | max_seqs=96: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1536 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536 | max_seqs=96: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=2048 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048 | max_seqs=96: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=3072 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072 | max_seqs=96: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=4096 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096 | max_seqs=96: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=6144 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144 | max_seqs=96:    0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=8192 | max-num-seqs=96
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192 | max_seqs=96: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=512 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=512 | max_seqs=128: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=768 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=768 | max_seqs=128: 0%|
0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1024 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1024 | max_seqs=128: 0%| 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=1536 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=1536 | max_seqs=128: 0%| 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=2048 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=2048 | max_seqs=128: 0%| 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=3072 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=3072 | max_seqs=128: 0%| 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=4096 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=4096 | max_seqs=128: 0%| 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=6144 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=6144 | max_seqs=128:    0%| 0/200 [00:00<?, ?it/s]
```

```
=====
Starting FAIR server | btok=8192 | max-num-seqs=128
=====
[1] not ready yet: ConnectionError
[2] not ready yet: ConnectionError
[3] not ready yet: ConnectionError
[4] not ready yet: ConnectionError
[5] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[6] not ready yet: ConnectionError
[7] not ready yet: ConnectionError
[8] not ready yet: ConnectionError
[9] not ready yet: ConnectionError
[10] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[11] not ready yet: ConnectionError
[12] not ready yet: ConnectionError
[13] not ready yet: ConnectionError
[14] not ready yet: ConnectionError
[15] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[16] not ready yet: ConnectionError
[17] not ready yet: ConnectionError
[18] not ready yet: ConnectionError
[19] not ready yet: ConnectionError
[20] not ready yet: ConnectionError

--- tail vllm_fair.log (last 30) ---
[21] not ready yet: ConnectionError
[22] not ready yet: ConnectionError
[23] not ready yet: ConnectionError
[24] not ready yet: ConnectionError
 Server ready. Models: ['Qwen/Qwen2.5-3B-Instruct']
FAIR | mix_50_50.jsonl | c=32 | btok=8192 | max_seqs=128: 0% |
0/200 [00:00<?, ?it/s]
 FAIR (btok x max-num-seqs) sweep done (resume-capable).
```

Out[]:	label	workload	concurrency	max_num_batched_tokens	max_num_seqs
0	FAIR	mix_50_50.jsonl	32	512	32
1	FAIR	mix_50_50.jsonl	32	768	32
2	FAIR	mix_50_50.jsonl	32	1024	32
3	FAIR	mix_50_50.jsonl	32	1536	32
4	FAIR	mix_50_50.jsonl	32	2048	32
5	FAIR	mix_50_50.jsonl	32	3072	32
6	FAIR	mix_50_50.jsonl	32	4096	32
7	FAIR	mix_50_50.jsonl	32	6144	32
8	FAIR	mix_50_50.jsonl	32	8192	32
9	FAIR	mix_50_50.jsonl	32	512	48
10	FAIR	mix_50_50.jsonl	32	768	48
11	FAIR	mix_50_50.jsonl	32	1024	48
12	FAIR	mix_50_50.jsonl	32	1536	48
13	FAIR	mix_50_50.jsonl	32	2048	48
14	FAIR	mix_50_50.jsonl	32	3072	48
15	FAIR	mix_50_50.jsonl	32	4096	48
16	FAIR	mix_50_50.jsonl	32	6144	48
17	FAIR	mix_50_50.jsonl	32	8192	48
18	FAIR	mix_50_50.jsonl	32	512	64
19	FAIR	mix_50_50.jsonl	32	768	64
20	FAIR	mix_50_50.jsonl	32	1024	64
21	FAIR	mix_50_50.jsonl	32	1536	64
22	FAIR	mix_50_50.jsonl	32	2048	64
23	FAIR	mix_50_50.jsonl	32	3072	64
24	FAIR	mix_50_50.jsonl	32	4096	64
25	FAIR	mix_50_50.jsonl	32	6144	64
26	FAIR	mix_50_50.jsonl	32	8192	64
27	FAIR	mix_50_50.jsonl	32	512	80
28	FAIR	mix_50_50.jsonl	32	768	80
29	FAIR	mix_50_50.jsonl	32	1024	80
30	FAIR	mix_50_50.jsonl	32	1536	80
31	FAIR	mix_50_50.jsonl	32	2048	80
32	FAIR	mix_50_50.jsonl	32	3072	80
33	FAIR	mix_50_50.jsonl	32	4096	80

label	workload	concurrency	max_num_batched_tokens	max_num_seqs
34	FAIR	mix_50_50.jsonl	32	6144
35	FAIR	mix_50_50.jsonl	32	8192
36	FAIR	mix_50_50.jsonl	32	512
37	FAIR	mix_50_50.jsonl	32	768
38	FAIR	mix_50_50.jsonl	32	1024
39	FAIR	mix_50_50.jsonl	32	1536
40	FAIR	mix_50_50.jsonl	32	2048
41	FAIR	mix_50_50.jsonl	32	3072
42	FAIR	mix_50_50.jsonl	32	4096
43	FAIR	mix_50_50.jsonl	32	6144
44	FAIR	mix_50_50.jsonl	32	8192
45	FAIR	mix_50_50.jsonl	32	512
46	FAIR	mix_50_50.jsonl	32	768
47	FAIR	mix_50_50.jsonl	32	1024
48	FAIR	mix_50_50.jsonl	32	1536
49	FAIR	mix_50_50.jsonl	32	2048
50	FAIR	mix_50_50.jsonl	32	3072
51	FAIR	mix_50_50.jsonl	32	4096
52	FAIR	mix_50_50.jsonl	32	6144
53	FAIR	mix_50_50.jsonl	32	8192

54 rows × 29 columns

```
In [ ]: import os
OUT_DIR = "/content/drive/MyDrive/vllm_sweeps"
os.makedirs(OUT_DIR, exist_ok=True)

fair_grid_df.to_csv(f"{OUT_DIR}/fair_grid_df_snapshot.csv", index=False)
print("✅ Saved:", f"{OUT_DIR}/fair_grid_df_snapshot.csv")
```

✅ Saved: /content/drive/MyDrive/vllm_sweeps/fair_grid_df_snapshot.csv

```
In [ ]: import shutil, os

local_ckpt = "fair_grid_checkpoint.csv"
local_manifest = "fair_grid_manifest.json"

if os.path.exists(local_ckpt):
    shutil.copy(local_ckpt, f"{OUT_DIR}/{local_ckpt}")
    print("✅ Copied:", f"{OUT_DIR}/{local_ckpt}")

if os.path.exists(local_manifest):
```

```

        shutil.copy(local_manifest, f"{OUT_DIR}/{local_manifest}")
        print("✅ Copied:", f"{OUT_DIR}/{local_manifest}")

    ✓ Copied: /content/drive/MyDrive/vllm_sweeps/fair_grid_checkpoint.csv
    ✓ Copied: /content/drive/MyDrive/vllm_sweeps/fair_grid_manifest.json

In [ ]: if os.path.exists("vllm_fair.log"):
        shutil.copy("vllm_fair.log", f"{OUT_DIR}/vllm_fair_last_run.log")

In [ ]: FAIR_CSV = "fair_grid_checkpoint.csv"
BAD_CSV   = "bad_grid_checkpoint.csv"

fair_df  = pd.read_csv(FAIR_CSV) if os.path.exists(FAIR_CSV) else None
bad_df   = pd.read_csv(BAD_CSV)  if os.path.exists(BAD_CSV) else None

def _prep(df):
    if df is None or df.empty:
        return None
    df = df.copy()

    # Keep only successful rows
    if "status" in df.columns:
        df = df[df["status"] == "ok"].copy()

    # numeric coercions
    for col in ["concurrency", "max_num_batched_tokens", "max_num_seqs"]:
        if col in df.columns:
            df[col] = pd.to_numeric(df[col], errors="coerce").astype("Int")

    for col in ["rps", "short_p99", "short_p90", "short_p50", "overall_p99",
                "long_p99", "long_p90", "long_p50", "wall_time_s"]:

        if col in df.columns:
            df[col] = pd.to_numeric(df[col], errors="coerce")

    return df

fair_df = _prep(fair_df)
bad_df  = _prep(bad_df)

print("FAIR rows:", 0 if fair_df is None else len(fair_df))
print("BAD rows :", 0 if bad_df is None else len(bad_df))

# heatmap plotter

def plot_heatmap(df, index_col, col_col, value_col, title):
    """
    index_col becomes y-axis, col_col becomes x-axis.
    """
    pivot = df.pivot_table(index=index_col, columns=col_col, values=value_col)
    pivot = pivot.sort_index().sort_index(axis=1)

    plt.figure()

    mat = pivot.values.astype(float)

    masked = np.ma.masked_invalid(mat)

```

```

im = plt.imshow(masked, aspect="auto", origin="lower")
plt.colorbar(im, label=value_col)

plt.title(title)
plt.xlabel(col_col)
plt.ylabel(index_col)

# tick labels
plt.xticks(ticks=np.arange(pivot.shape[1]), labels=pivot.columns.tolist())
plt.yticks(ticks=np.arange(pivot.shape[0]), labels=pivot.index.tolist())
plt.tight_layout()
plt.show()

# FAIR plots (btok x max_num_seqs)

if fair_df is not None and {"max_num_batched_tokens", "max_num_seqs"}.issubset(fair_df.columns):
    # Heatmap: short_p99
    plot_heatmap(
        fair_df,
        index_col="max_num_seqs",
        col_col="max_num_batched_tokens",
        value_col="short_p99",
        title="FAIR grid: short_p99 (lower is better)"
    )

    # Heatmap: rps
    plot_heatmap(
        fair_df,
        index_col="max_num_seqs",
        col_col="max_num_batched_tokens",
        value_col="rps",
        title="FAIR grid: throughput rps (higher is better)"
    )

    # Pareto scatter: rps vs short_p99
    plt.figure()
    plt.scatter(fair_df["short_p99"], fair_df["rps"])
    plt.xlabel("short_p99 (s) ↓")
    plt.ylabel("rps ↑")
    plt.title("FAIR grid: Pareto view (rps vs short_p99)")
    plt.tight_layout()
    plt.show()

    # Slices: short_p99 vs btok for each max_num_seqs
    for seqs in sorted(fair_df["max_num_seqs"].dropna().unique()):
        sub = fair_df[fair_df["max_num_seqs"] == seqs].sort_values("max_num_batched_tokens")
        if len(sub) < 2:
            continue
        plt.figure()
        plt.plot(sub["max_num_batched_tokens"], sub["short_p99"], marker="x")
        plt.xlabel("max_num_batched_tokens")
        plt.ylabel("short_p99 (s)")
        plt.title(f"FAIR: short_p99 vs btok (max_num_seqs={seqs})")
        plt.tight_layout()
        plt.show()

    # Slices: short_p99 vs max_num_seqs for each btok
    for btok in sorted(fair_df["max_num_batched_tokens"].dropna().unique()):
        sub = fair_df[fair_df["max_num_batched_tokens"] == btok].sort_values("max_num_seqs")
        if len(sub) < 2:
            continue
        plt.figure()
        plt.plot(sub["max_num_seqs"], sub["short_p99"], marker="x")
        plt.xlabel("max_num_seqs")
        plt.ylabel("short_p99 (s)")
        plt.title(f"FAIR: short_p99 vs max_num_seqs (btok={btok})")
        plt.tight_layout()
        plt.show()

```

```

        sub = fair_df[fair_df["max_num_batched_tokens"] == btok].sort_values()
        if len(sub) < 2:
            continue
        plt.figure()
        plt.plot(sub["max_num_seqs"], sub["short_p99"], marker="o")
        plt.xlabel("max_num_seqs")
        plt.ylabel("short_p99 (s)")
        plt.title(f"FAIR: short_p99 vs max_num_seqs (btok={btok})")
        plt.tight_layout()
        plt.show()

# BAD plots (concurrency x max_num_seqs)

if bad_df is not None and {"concurrency", "max_num_seqs"}.issubset(bad_df):
    plot_heatmap(
        bad_df,
        index_col="max_num_seqs",
        col_col="concurrency",
        value_col="short_p99",
        title="BAD grid: short_p99 across (max_num_seqs x concurrency)"
    )

    plot_heatmap(
        bad_df,
        index_col="max_num_seqs",
        col_col="concurrency",
        value_col="rps",
        title="BAD grid: rps across (max_num_seqs x concurrency)"
    )

# Pareto scatter: rps vs short_p99 (BAD)
plt.figure()
plt.scatter(bad_df["short_p99"], bad_df["rps"])
plt.xlabel("short_p99 (s) ↑")
plt.ylabel("rps ↑")
plt.title("BAD grid: Pareto view (rps vs short_p99)")
plt.tight_layout()
plt.show()

# FAIR vs BAD baseline delta plot (optional but very meaningful)

# If BAD baseline at concurrency=32
BAD_BASELINE_SHORT_P99 = 10.618033962269847 # from (c=32) experiment
BAD_BASELINE_RPS = None

if fair_df is not None and "short_p99" in fair_df.columns:
    fair_df2 = fair_df.copy()
    fair_df2["delta_short_p99_vs_bad"] = fair_df2["short_p99"] - BAD_BASELINE_SHORT_P99

# show configs that beat BAD (negative delta)
winners = fair_df2.sort_values("delta_short_p99_vs_bad").head(15)
print("\nTop FAIR configs by short_p99 improvement vs BAD baseline (m")
cols = ["max_num_batched_tokens", "max_num_seqs", "rps", "short_p50", "sh"]
print(winners[cols].to_string(index=False))

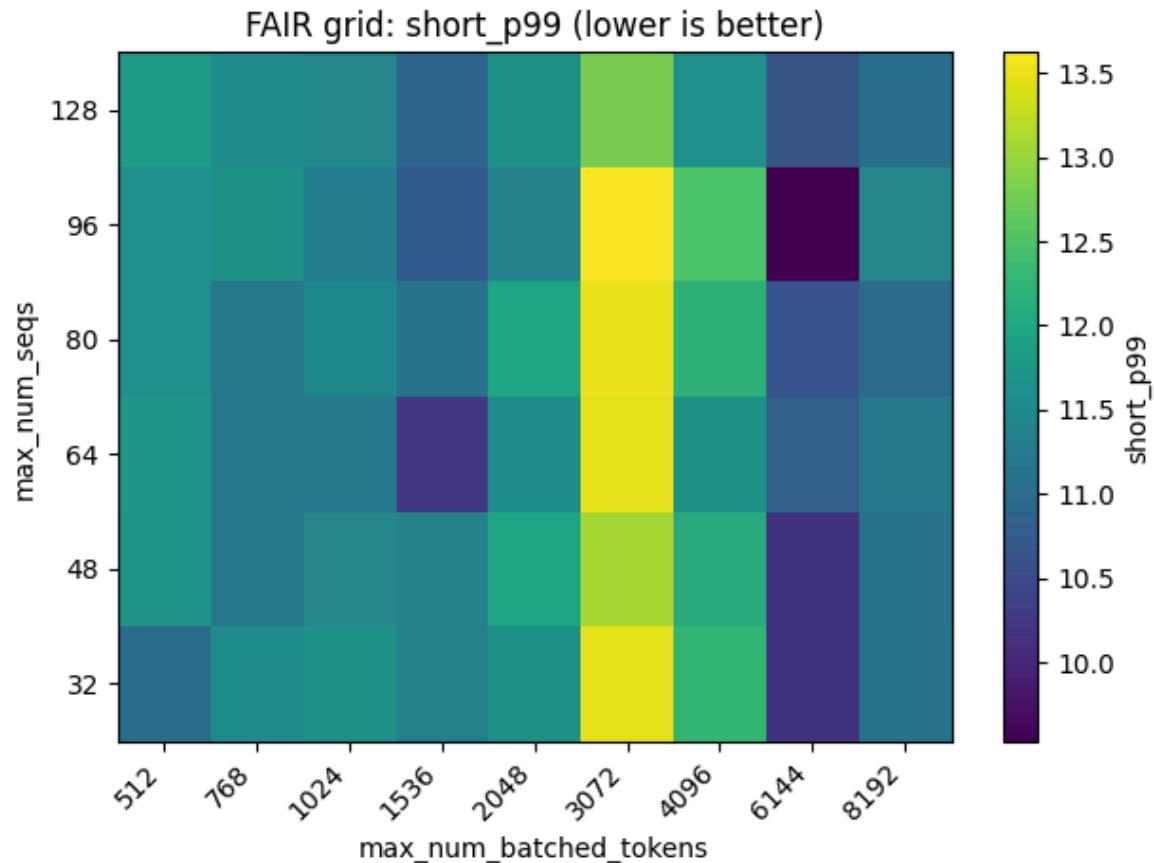
# Scatter: delta_short_p99 vs rps (winners are below 0)
plt.figure()
plt.scatter(fair_df2["delta_short_p99_vs_bad"], fair_df2["rps"])

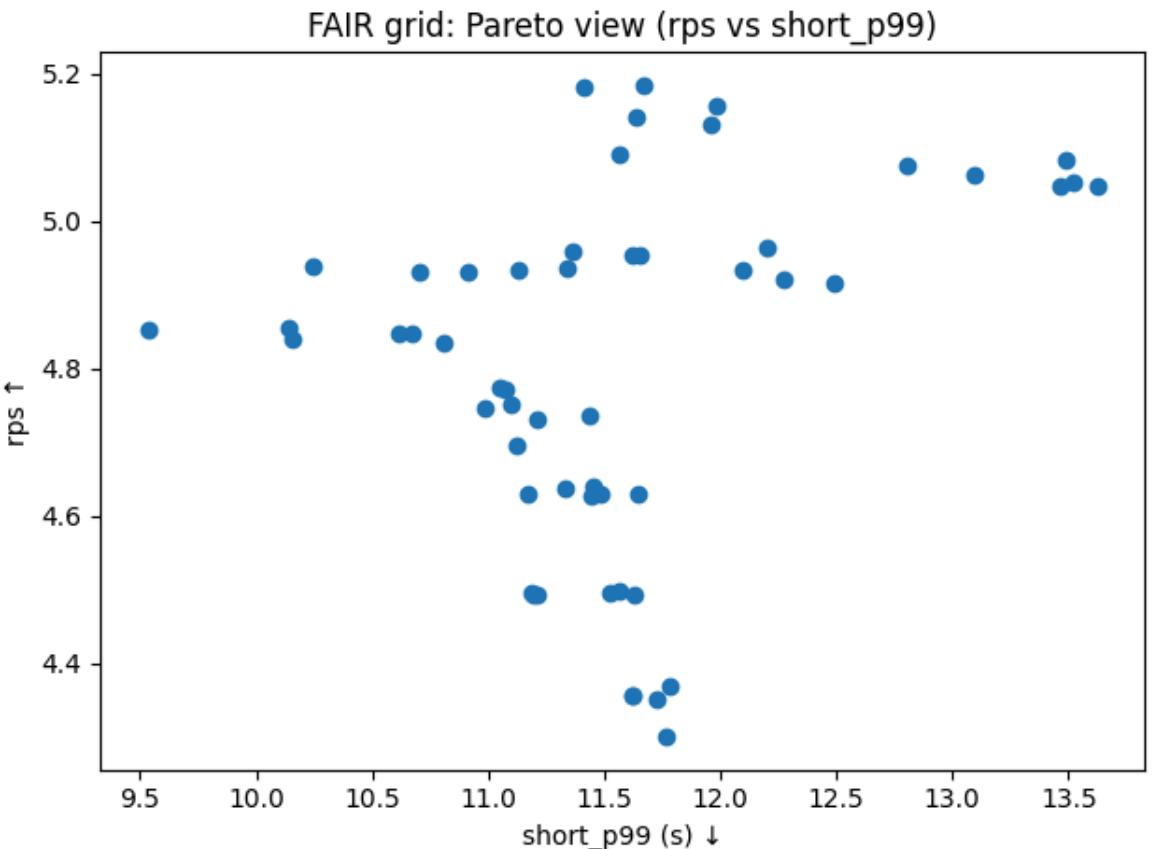
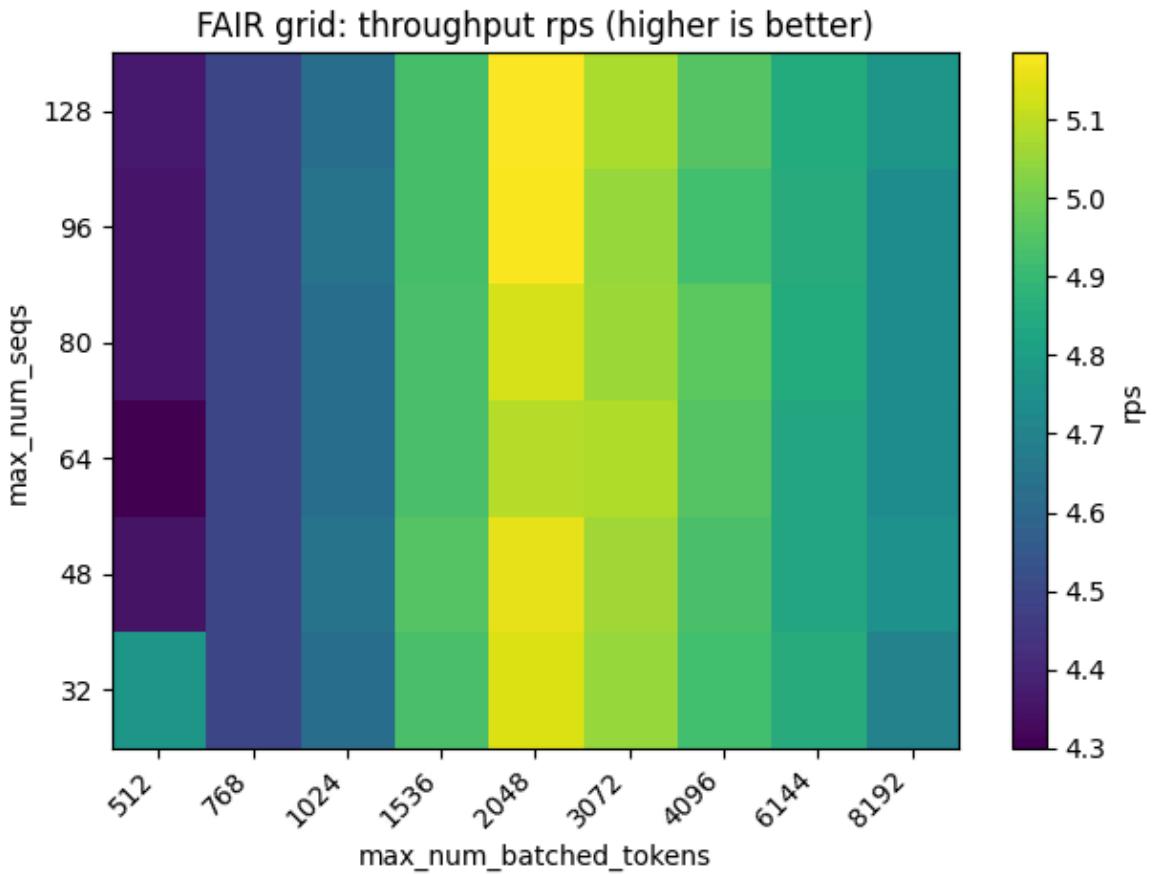
```

```
plt.axvline(0, linestyle="--")
plt.xlabel("Δ short_p99 vs BAD (s) (negative = FAIR better)")
plt.ylabel("rps")
plt.title("FAIR: improvement vs BAD baseline (rps vs Δshort_p99)")
plt.tight_layout()
plt.show()
```

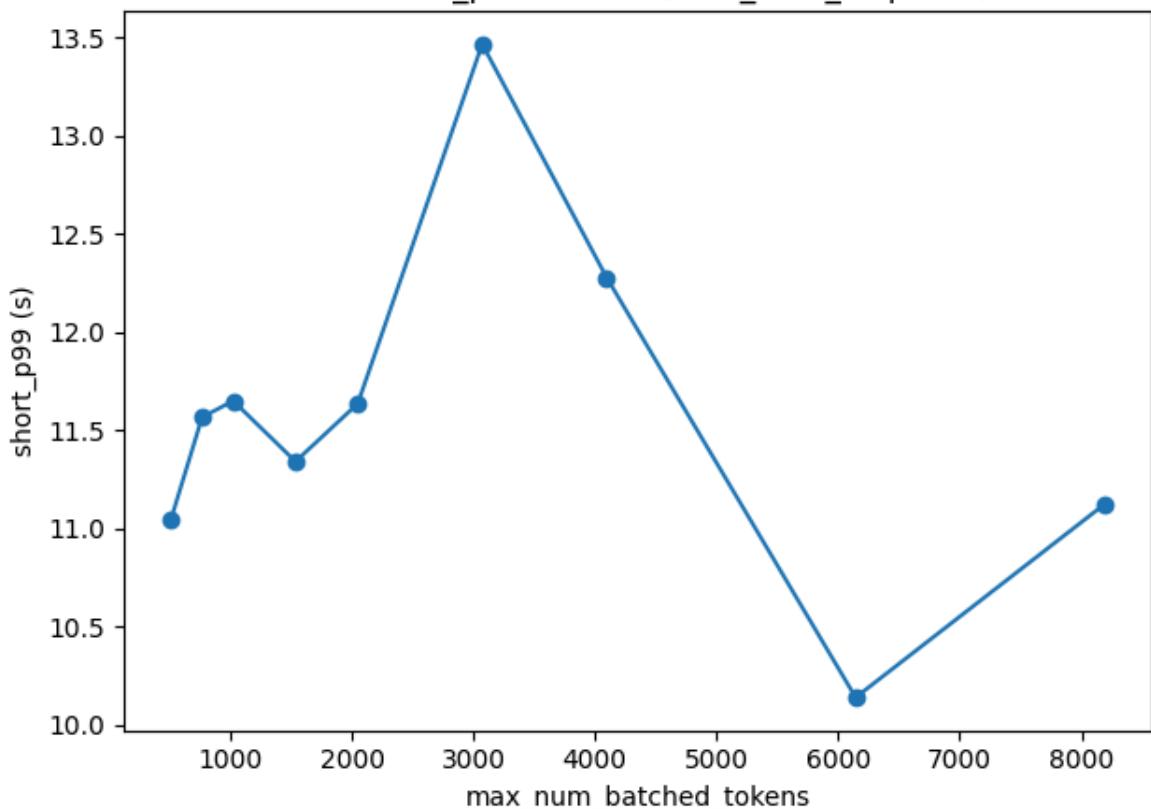
FAIR rows: 54

BAD rows : 48

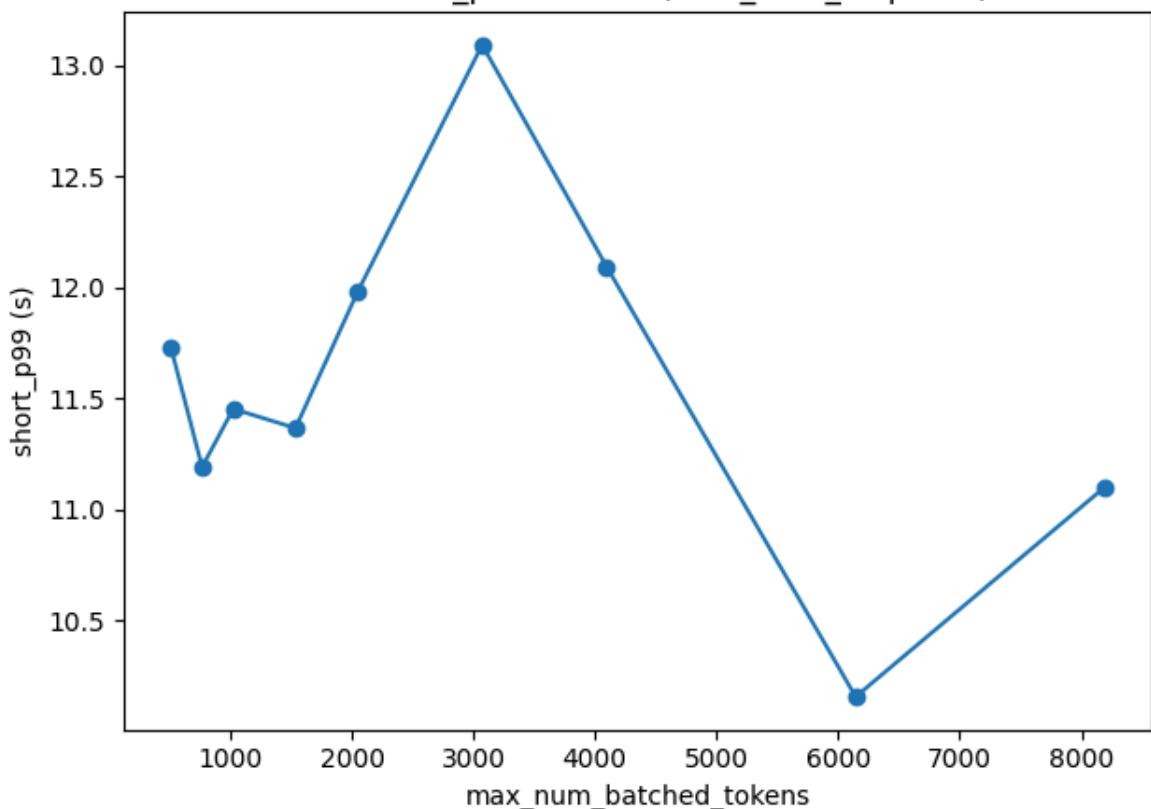




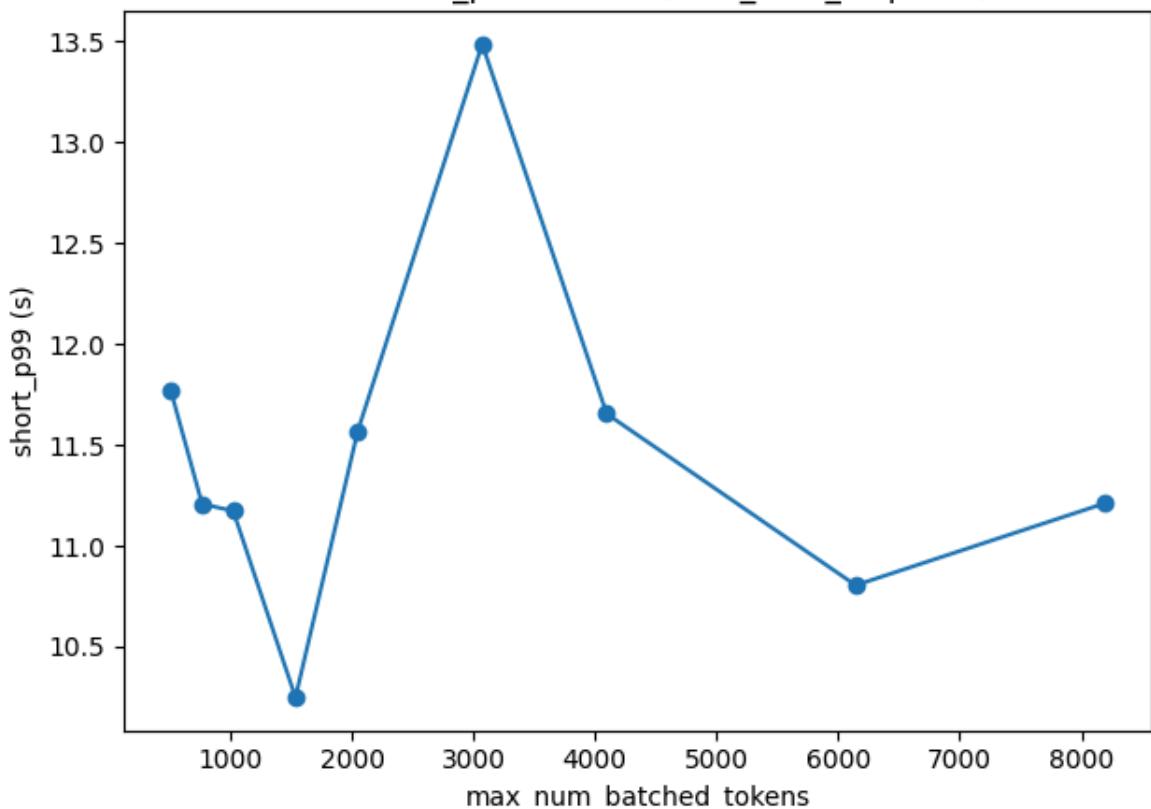
FAIR: short_p99 vs btok (max_num_seqs=32)



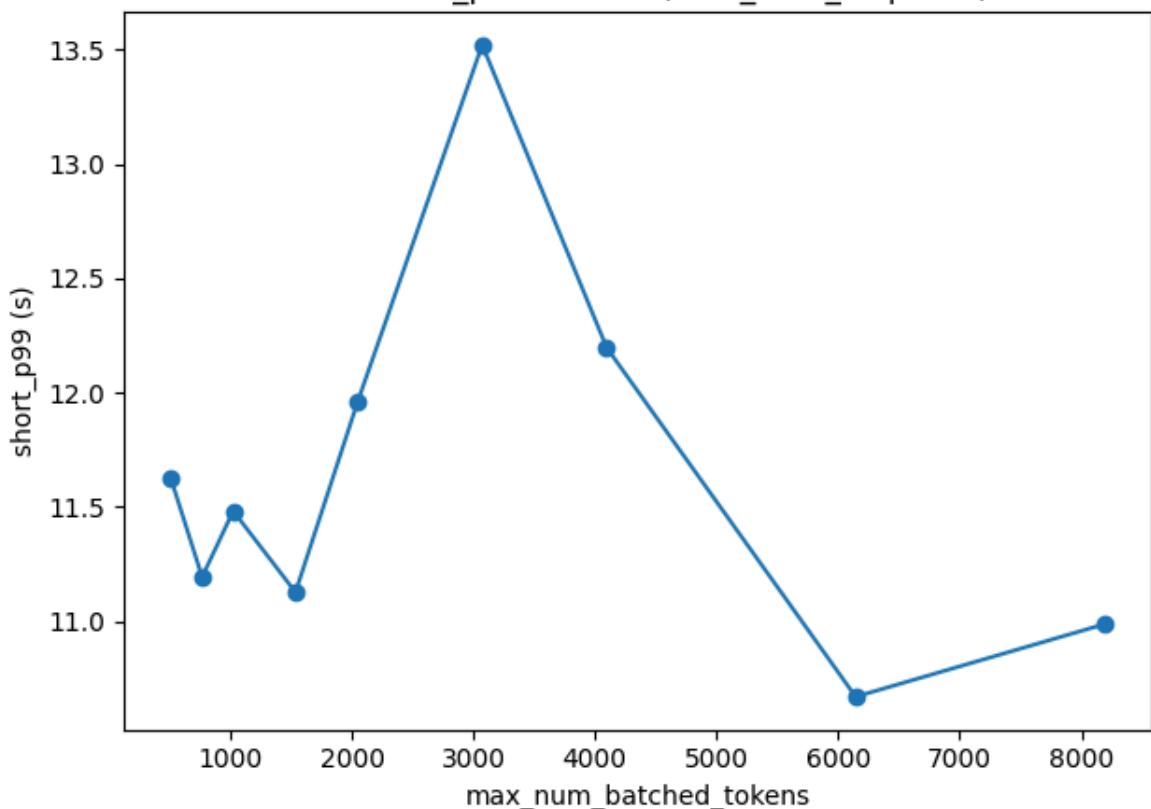
FAIR: short_p99 vs btok (max_num_seqs=48)



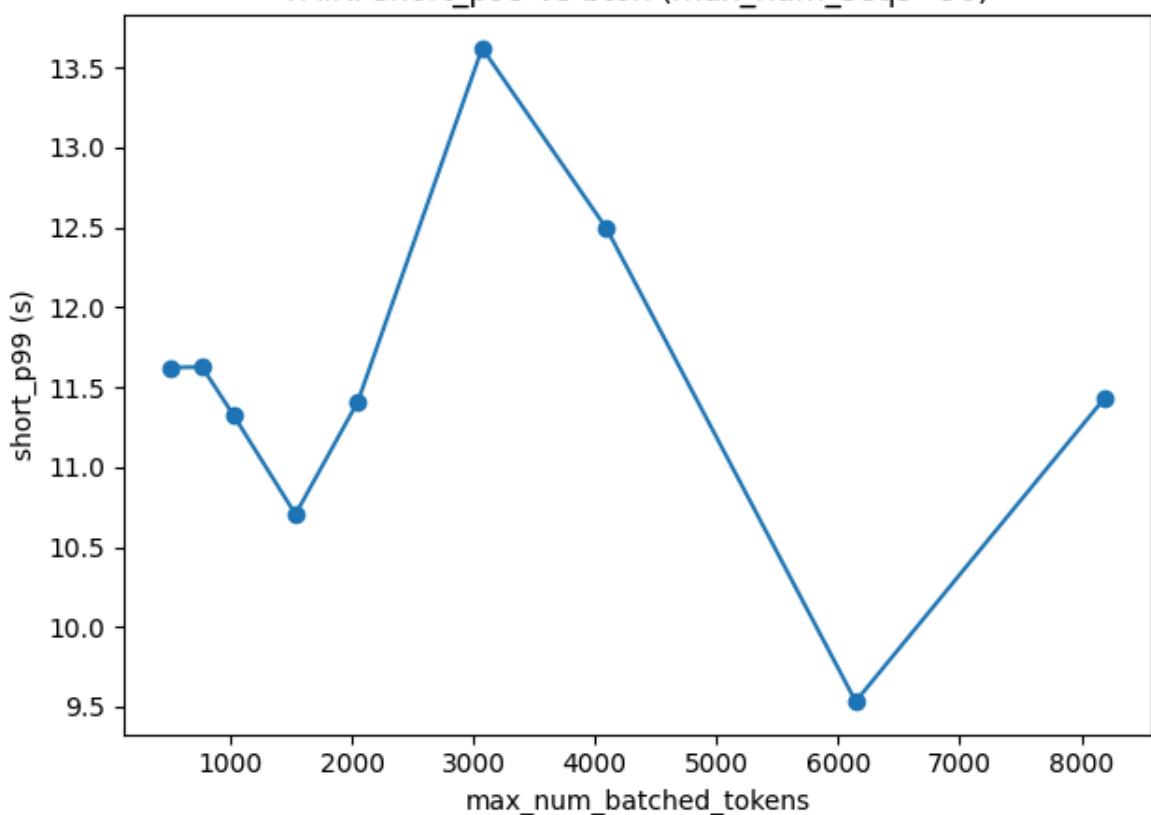
FAIR: short_p99 vs btok (max_num_seqs=64)



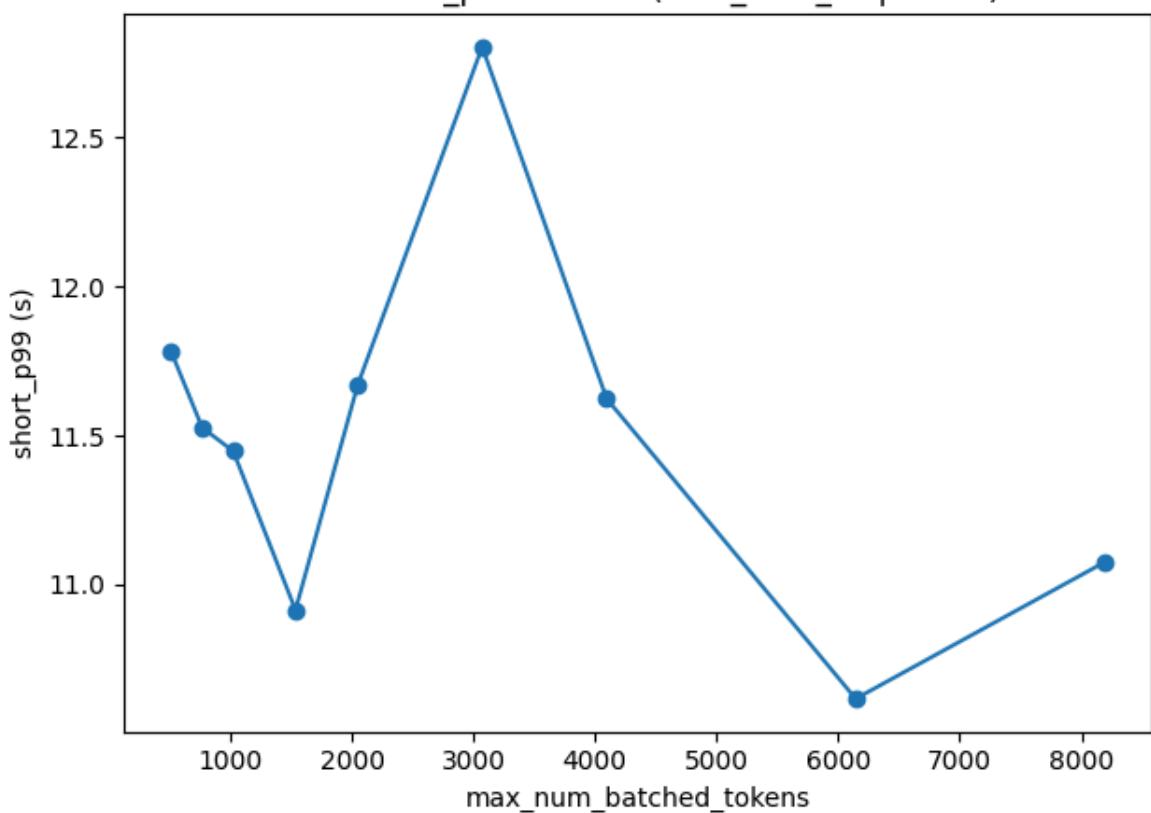
FAIR: short_p99 vs btok (max_num_seqs=80)

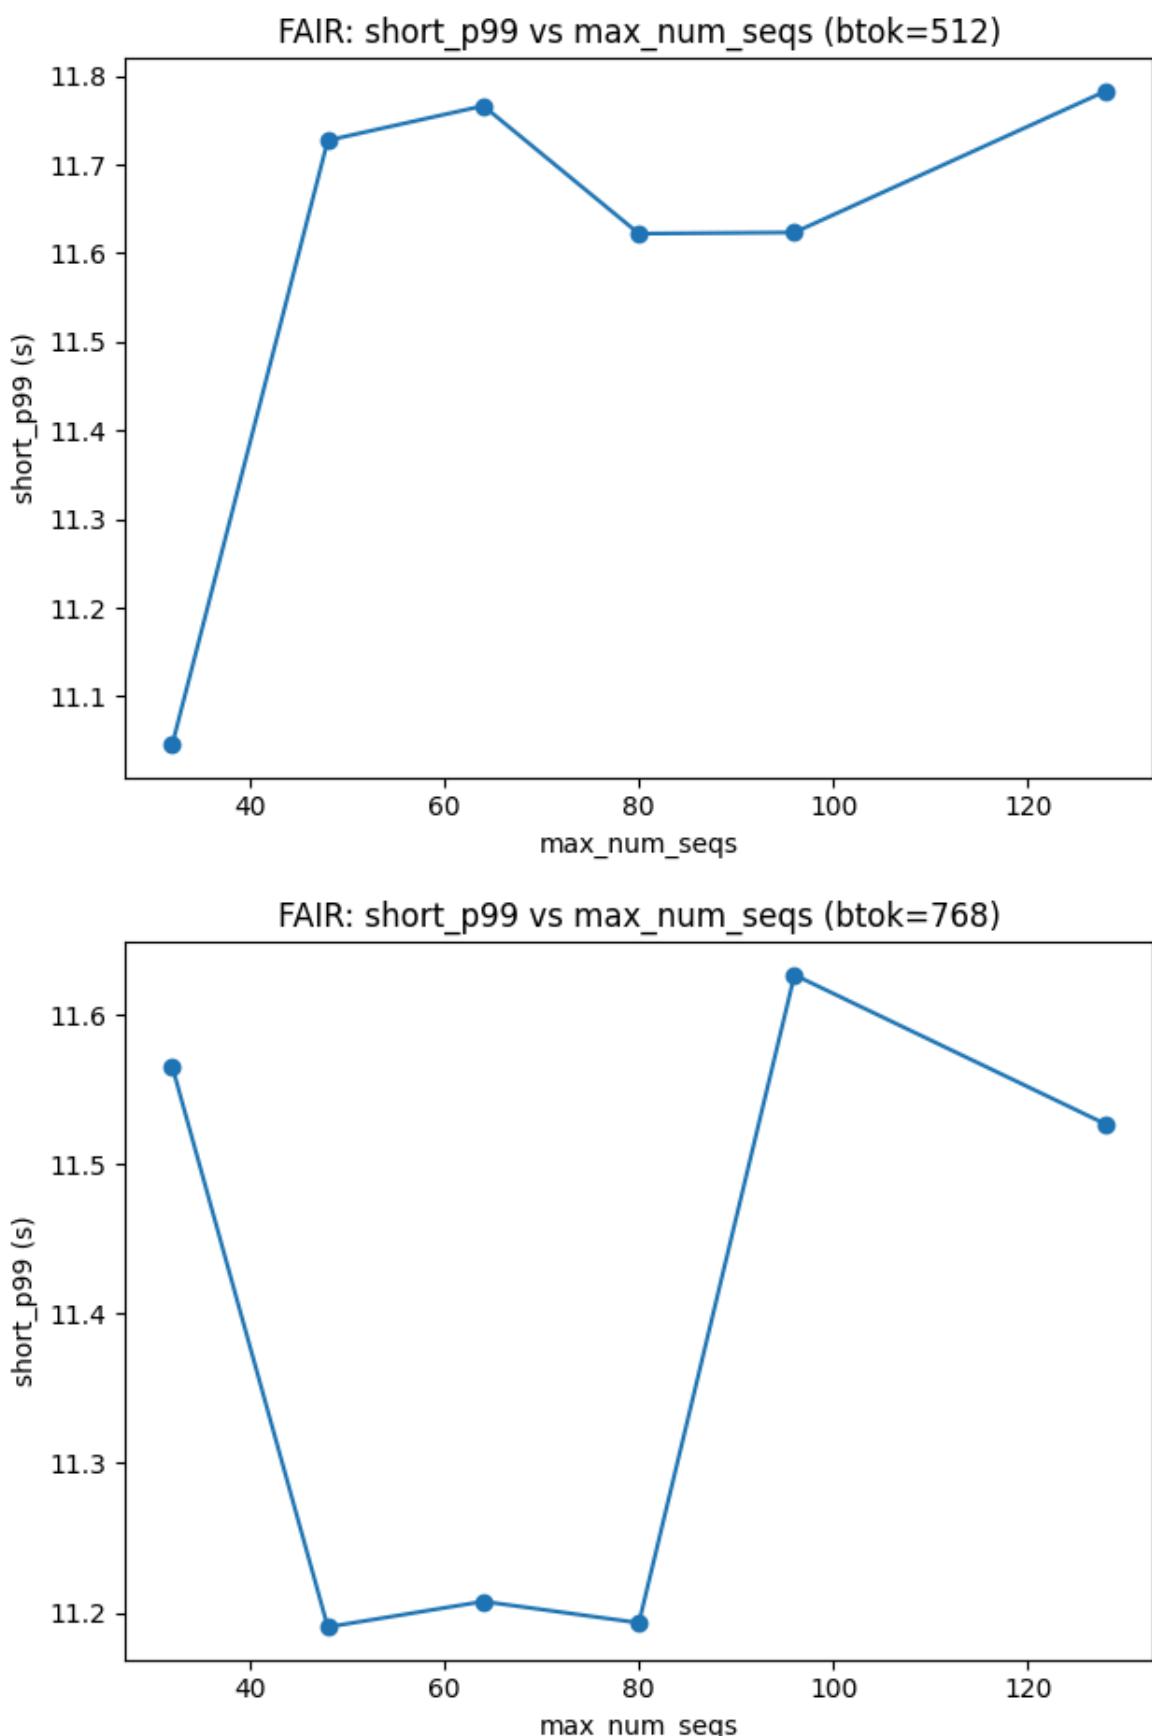


FAIR: short_p99 vs btok (max_num_seqs=96)

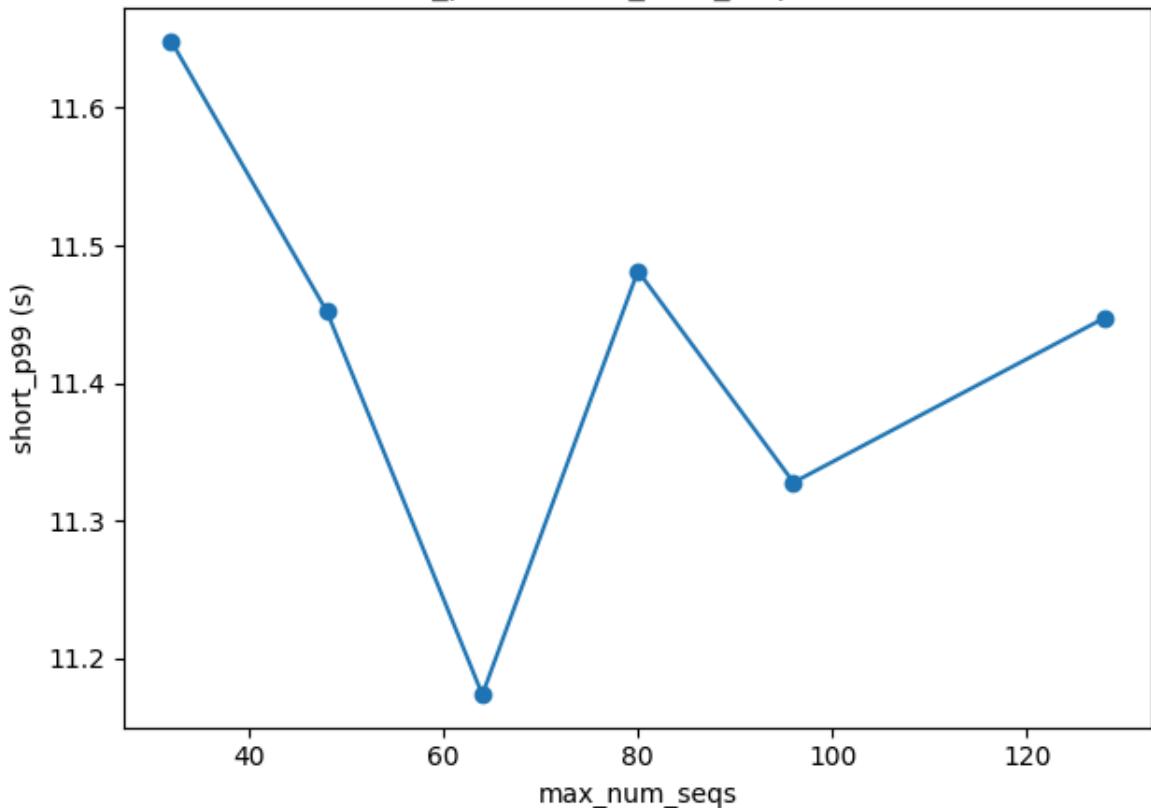


FAIR: short_p99 vs btok (max_num_seqs=128)

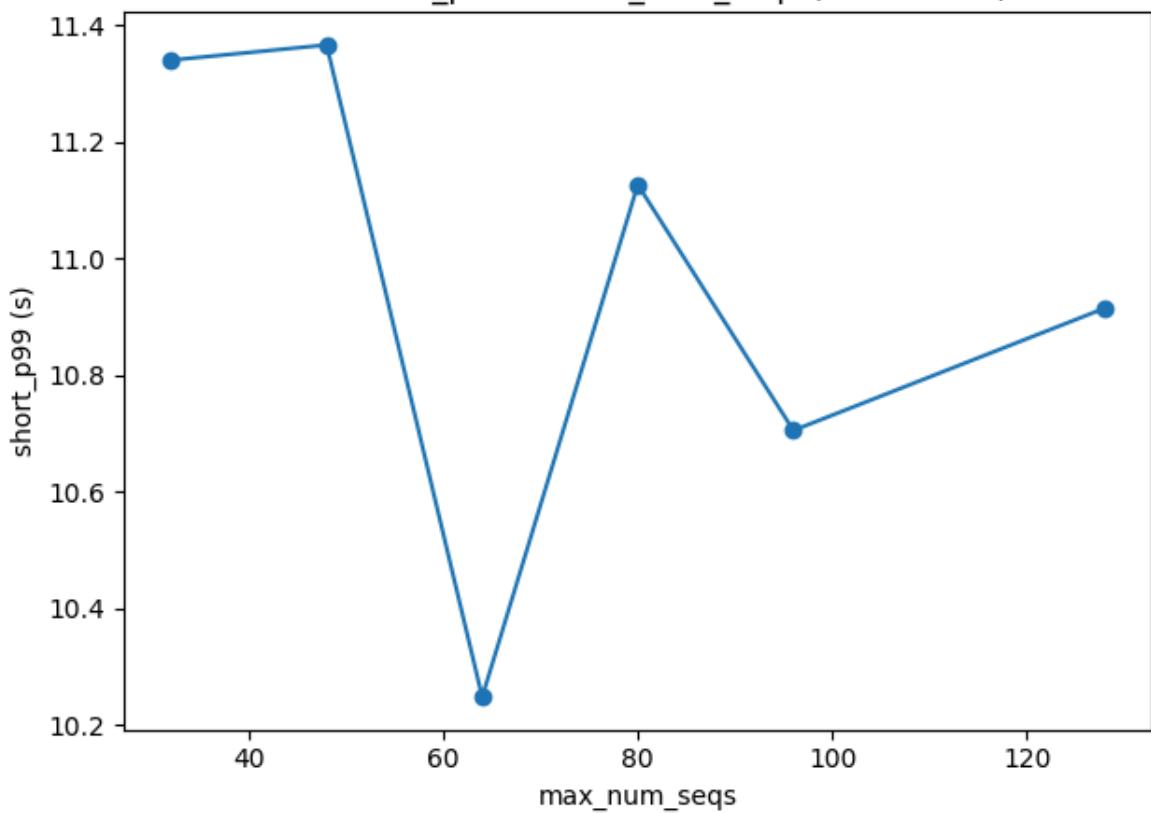


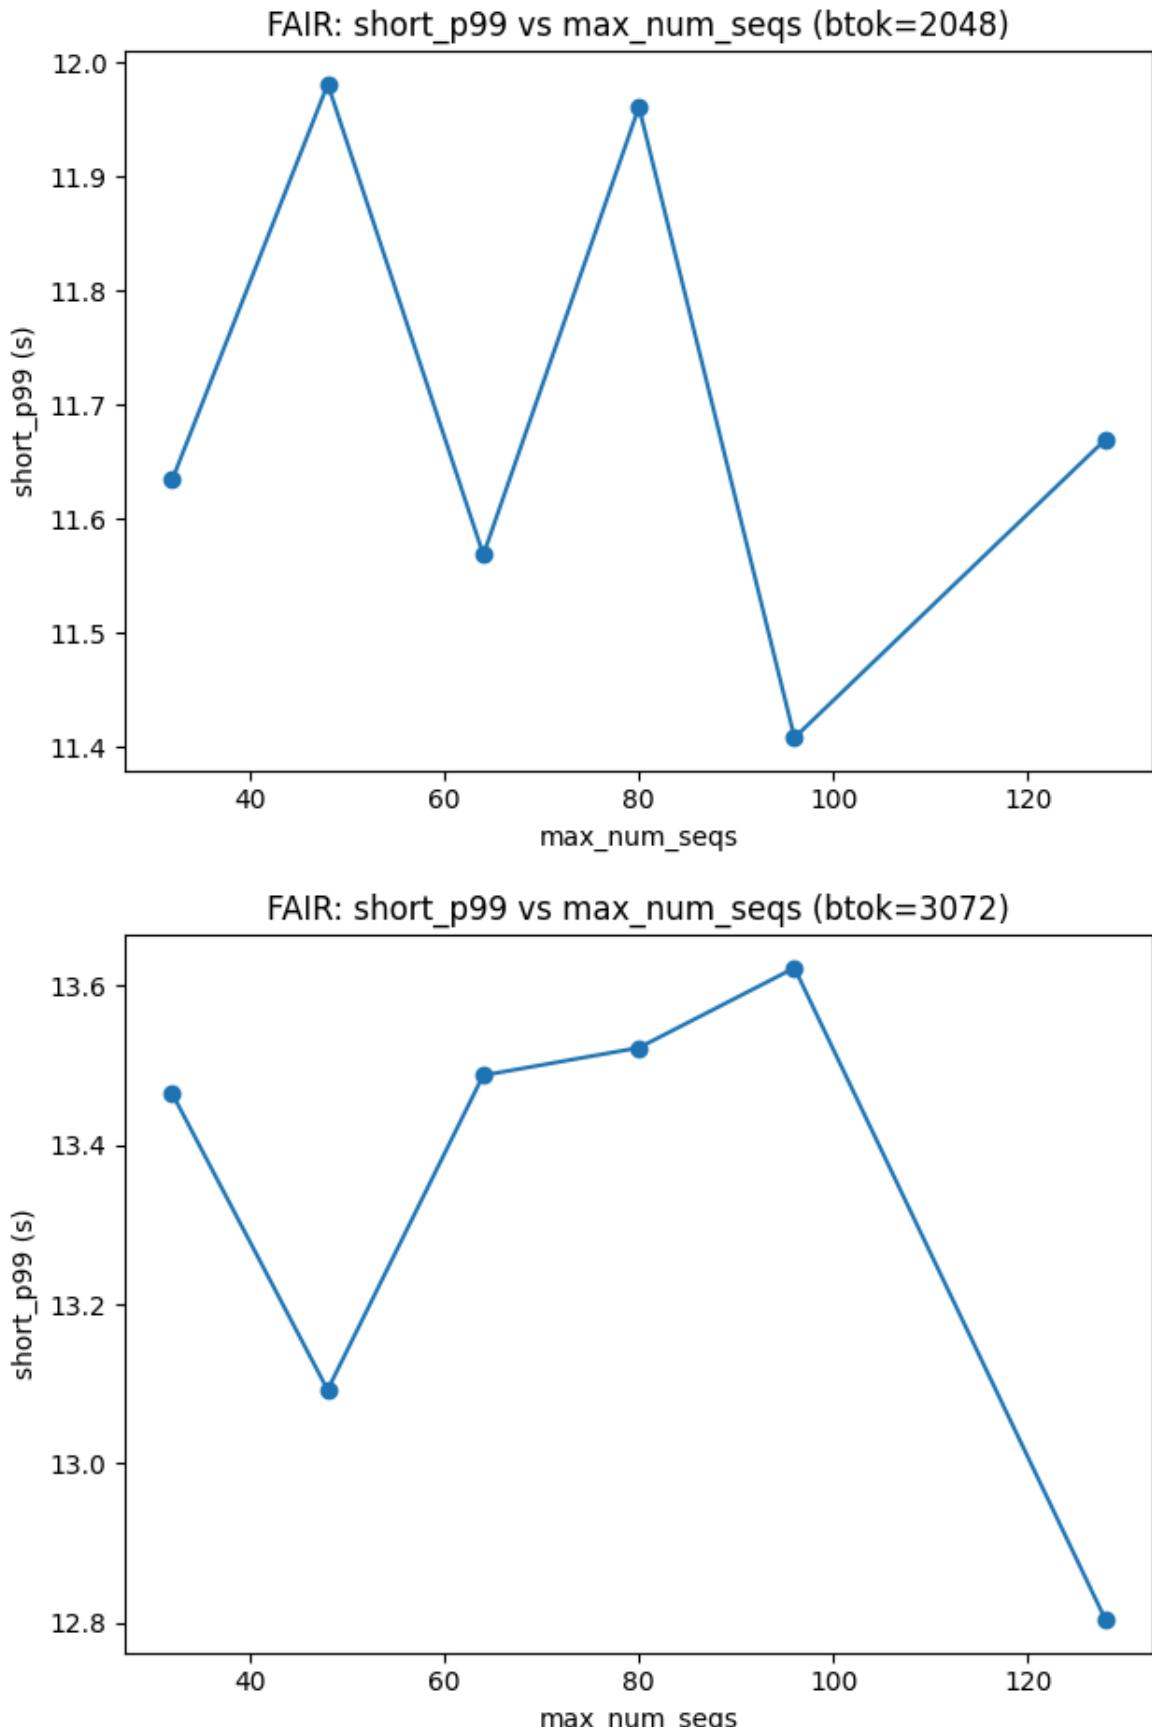


FAIR: short_p99 vs max_num_seqs (btok=1024)

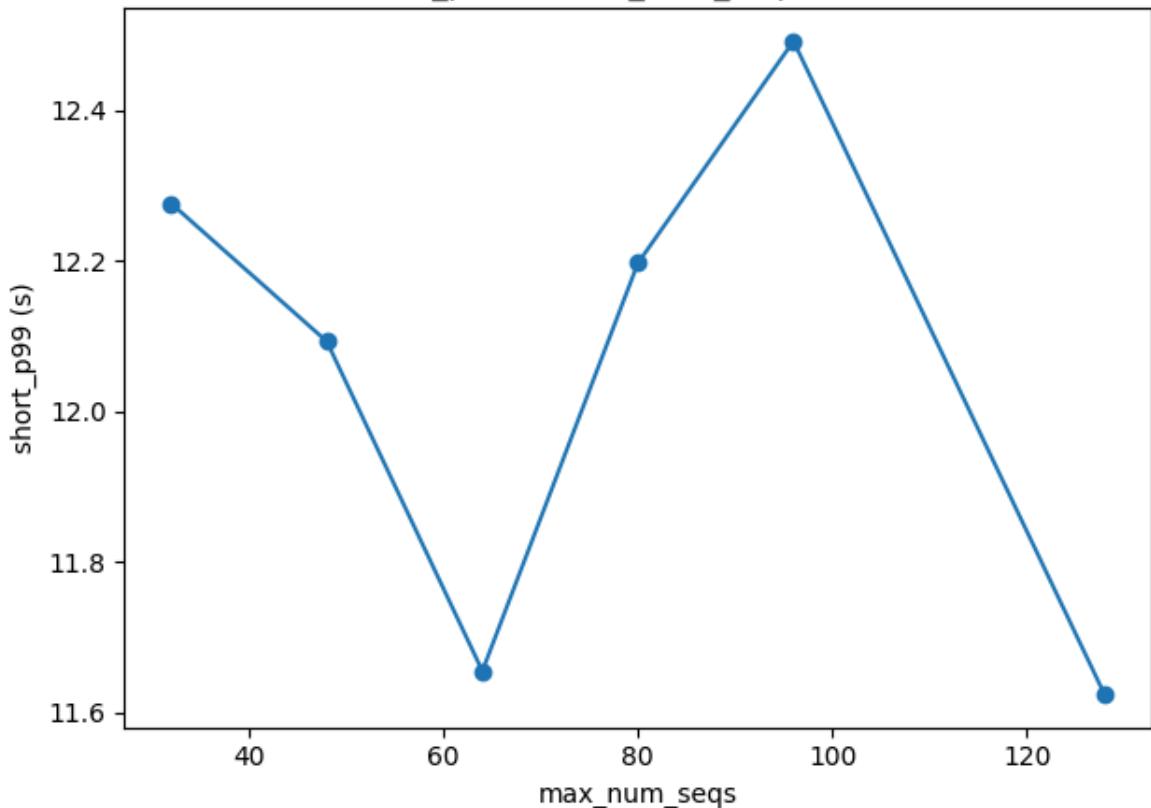


FAIR: short_p99 vs max_num_seqs (btok=1536)

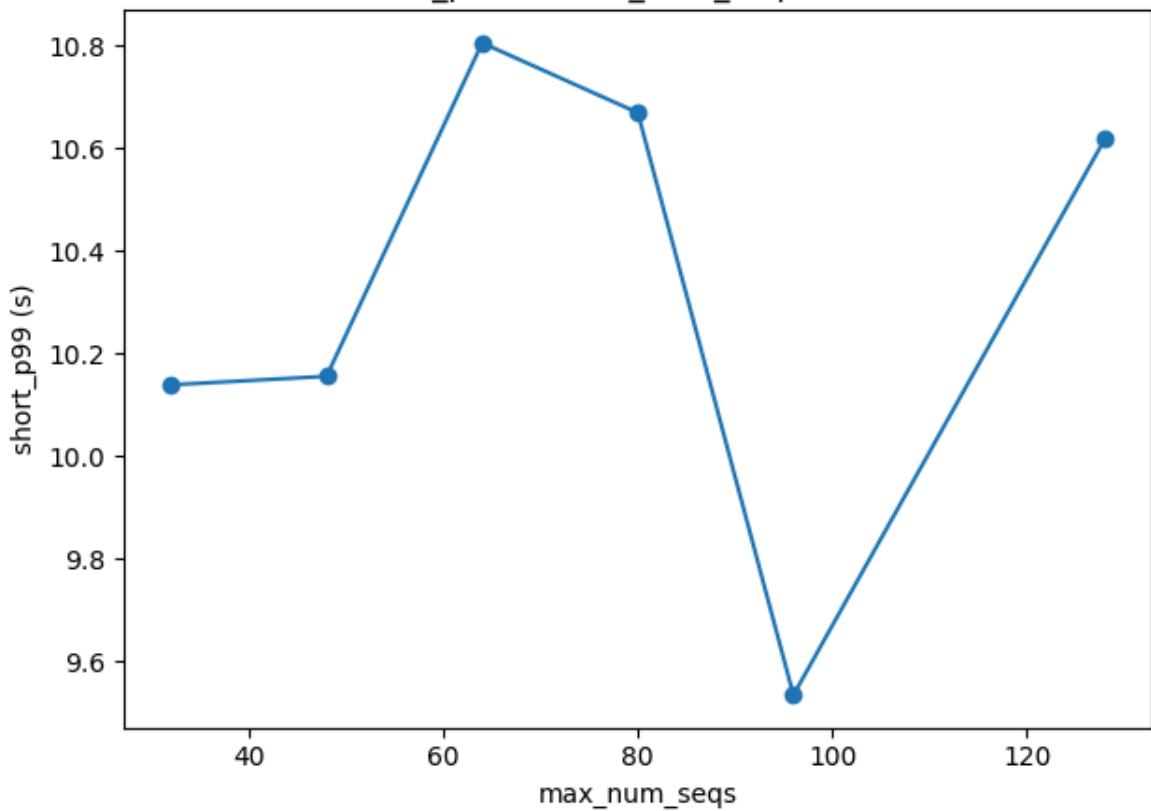




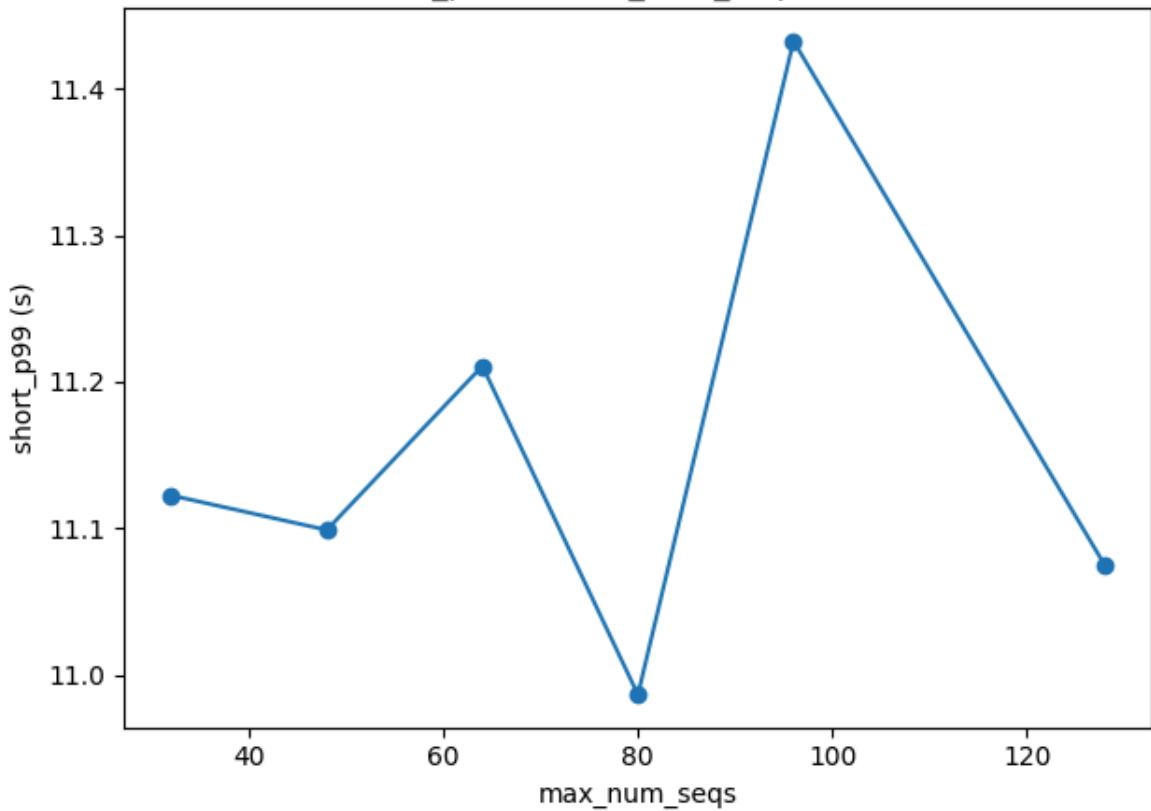
FAIR: short_p99 vs max_num_seqs (btok=4096)



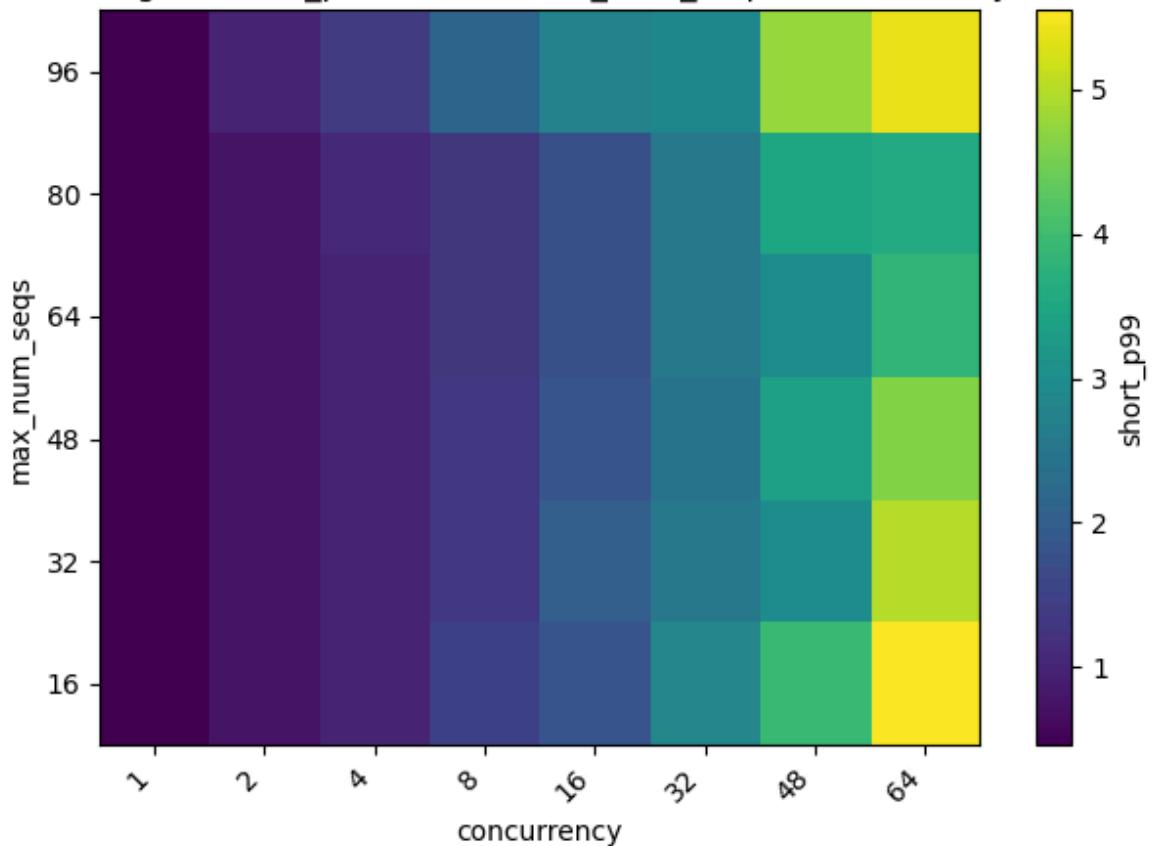
FAIR: short_p99 vs max_num_seqs (btok=6144)

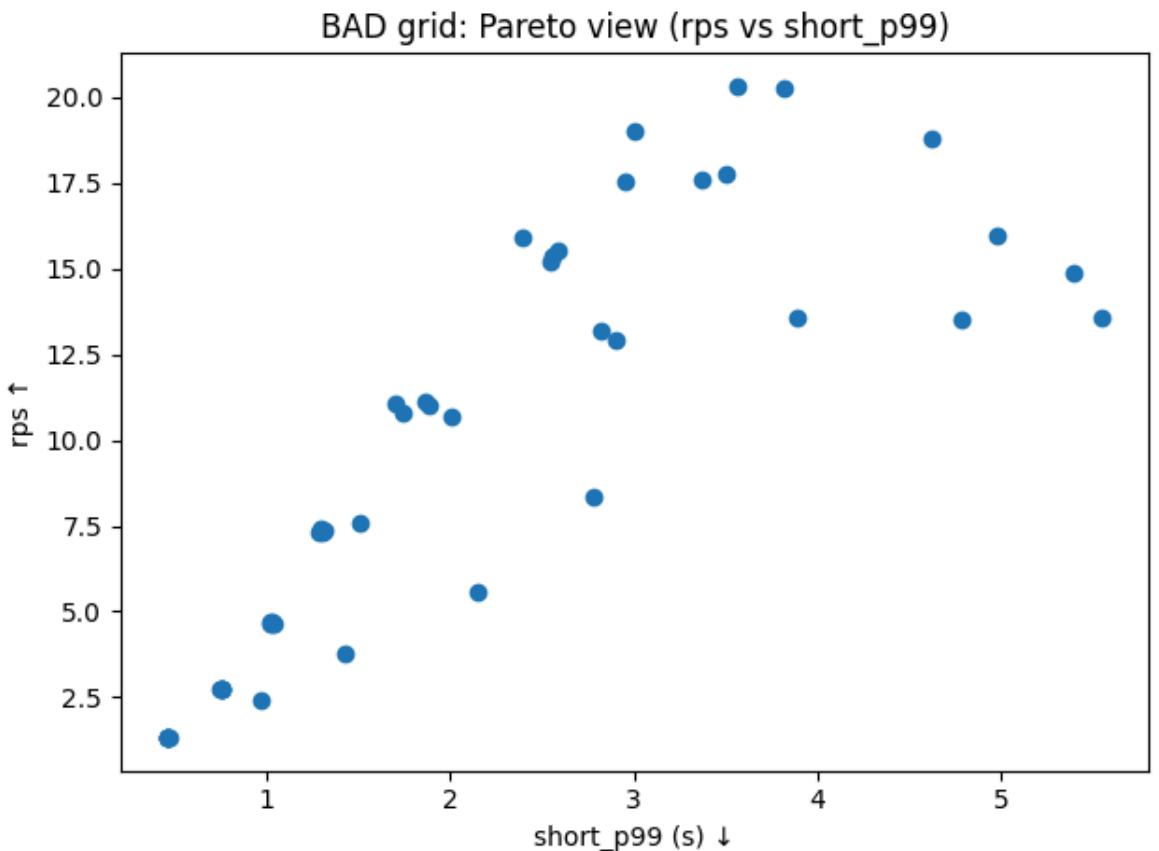
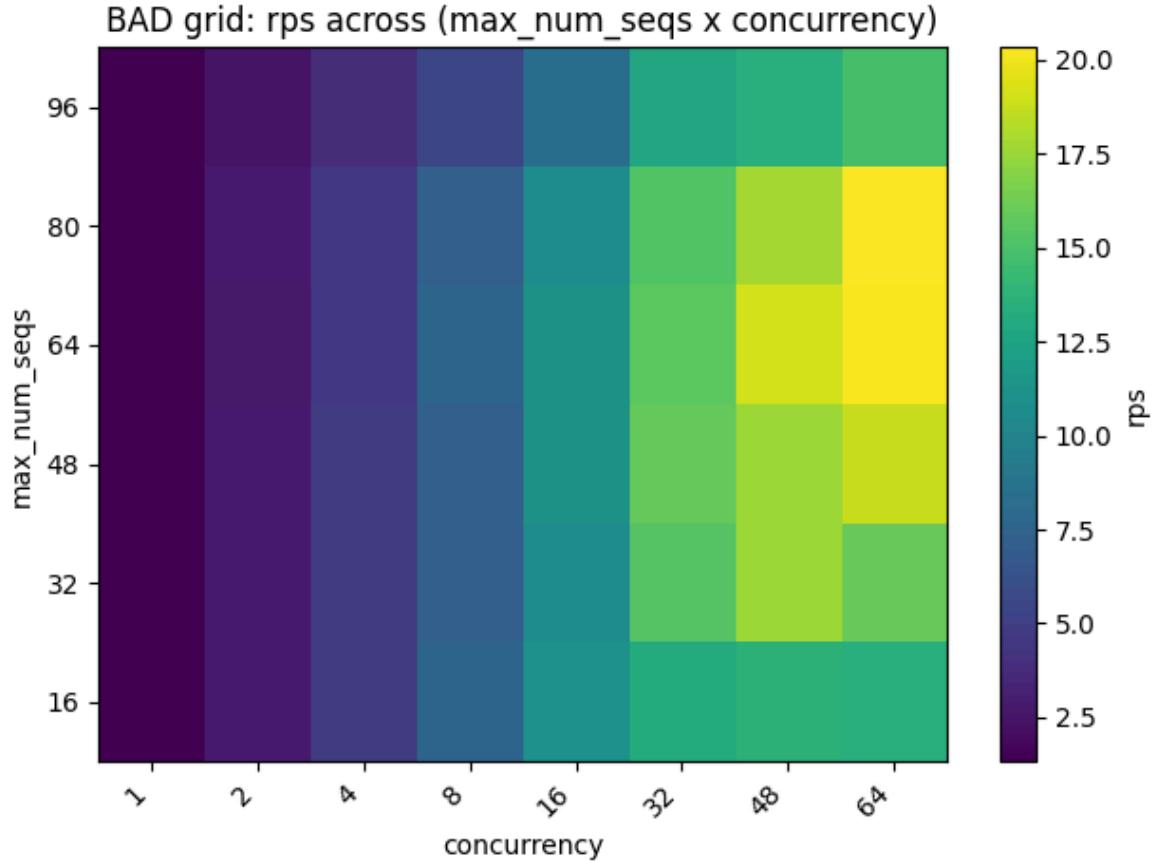


FAIR: short_p99 vs max_num_seqs (btok=8192)



BAD grid: short_p99 across (max_num_seqs x concurrency)

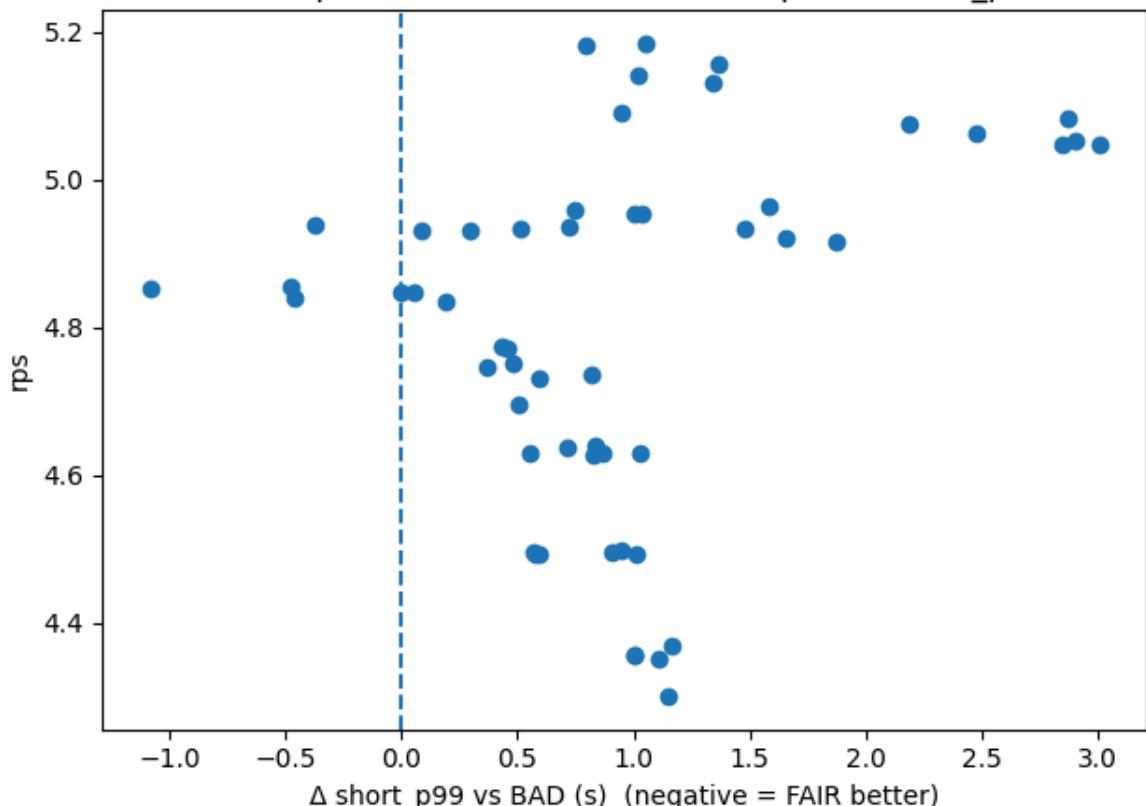




Top FAIR configs by short_p99 improvement vs BAD baseline (most negative is best):

	max_num_batched_tokens	max_num_seqs	rps	short_p50	short_p90	short_p99
	delta_short_p99_vs_bad					
33743	6144 -1.084291	96	4.853246	5.389539	9.325518	9.5
38313	6144 -0.479721	32	4.855452	4.769299	9.536749	10.1
54904	6144 -0.463130	48	4.838867	5.196080	9.575984	10.1
48558	1536 -0.369476	64	4.938646	6.086197	8.717633	10.2
16775	6144 -0.001259	128	4.847778	4.985013	10.048980	10.6
68710	6144 0.050676	80	4.848198	4.672418	9.520353	10.6
05149	1536 0.087115	96	4.931036	6.244107	9.136553	10.7
04741	6144 0.186707	64	4.834418	4.366397	10.301273	10.8
14505	1536 0.296471	128	4.930590	6.238226	8.916273	10.9
86636	8192 0.368602	80	4.745860	4.424287	9.958120	10.9
45613	512 0.427579	32	4.773907	5.700893	9.829510	11.0
75057	8192 0.457023	128	4.772023	4.867562	8.871036	11.0
98977	8192 0.480943	48	4.752496	4.614139	9.769913	11.0
22538	8192 0.504504	32	4.694382	5.120950	10.116705	11.1
26135	1536 0.508101	80	4.933366	6.336045	9.076619	11.1

FAIR: improvement vs BAD baseline (rps vs Δ short_p99)



Conclusion: FAIR grid ($\text{btok} \times \text{max-num-seqs}$ @ $c=32$): interpretation (throughput \leftrightarrow tail tradeoff)

- **Baseline (BAD_unfair @ c=32):** overall p99 $\approx 11.05\text{s}$, short p99 $\approx 10.15\text{s}$, long p99 $\approx 11.06\text{s}$.
This is the reference point for judging whether FAIR is “better” (and *which* metric we care about).
- **Within FAIR, throughput varies modestly, but tails move a lot.**
 - RPS stays roughly in $\sim 4.30 \rightarrow \sim 5.19 \text{ rps}$, so the main lever is *where* you land on the **RPS vs p99** curve.
- **Best FAIR config for *long-tail control* (closest to BAD long p99):**
 - **btok=512, max_num_seqs=32** \rightarrow long p99 $\approx 11.80\text{s}$ (best/lowest long p99 in the grid) with RPS ≈ 4.77 .
 - Compared to BAD long p99 (**11.06s**), FAIR here is **$\sim +6.7\%$ higher** long p99, but it’s the *tightest* long-tail option among FAIR settings.
- **Best FAIR config for *protecting short requests* (lowest short p99):**
 - **btok=6144, max_num_seqs=96** \rightarrow short p99 $\approx 9.53\text{s}$ with RPS ≈ 4.85 .
 - Compared to BAD short p99 (**10.15s**), this is **$\sim 6.1\%$ lower** (better), but it comes with a higher long p99 ($\approx 13.09\text{s}$), i.e., fairness helps short requests but can tax long tails if batching is too aggressive.
- **Peak-throughput region (if you’re chasing RPS) comes with a tail cost:**
 - Highest RPS is around **btok=2048, max_num_seqs=96–128** \rightarrow RPS $\approx 5.18\text{--}5.19$.
 - But long p99 is typically **$\sim 12.8\text{--}12.9\text{s}$** and short p99 **$\sim 11.4\text{--}11.7\text{s}$** , which is worse than BAD on p99 even though throughput is higher.
- **Takeaway:** there isn’t one “best” FAIR config — it depends on the SLO you want to optimize.
 - If your SLO is **short p99**, a high-btok / moderate-seqs region (e.g., **6144 × 96**) is attractive.
 - If your SLO is **long p99**, the most conservative FAIR corner (e.g., **512 × 32**) is safest.
 - For production, you’d pick the **knee**: *near-peak RPS without large p99 inflation*, then validate with a final concurrency sweep.

Code to convert notebook (.ipynb) to .pdf

In [7]: `import json`

```
in_path  = "/content/Copy_of_vllm_serving_project.ipynb"
out_path = "/content/Copy_of_vllm_serving_project_clean.ipynb"
```

```
with open(in_path, "r", encoding="utf-8") as f:
    nb = json.load(f)

nb.get("metadata", {}).pop("widgets", None)

for cell in nb.get("cells", []):
    if cell.get("cell_type") != "code":
        continue
    for out in cell.get("outputs", []):
        out.pop("metadata", {}).pop("widgets", None)

with open(out_path, "w", encoding="utf-8") as f:
    json.dump(nb, f, ensure_ascii=False, indent=1)

print("Wrote cleaned notebook:", out_path)
```

Wrote cleaned notebook: /content/Copy_of_vllm_serving_project_clean.ipynb

In [8]: !pip -q install nbconvert
!jupyter nbconvert --to html "/content/Copy_of_vllm_serving_project_clean.ipynb"

[NbConvertApp] Converting notebook /content/Copy_of_vllm_serving_project_clean.ipynb to html
[NbConvertApp] WARNING | Alternative text is missing on 39 image(s).
[NbConvertApp] Writing 3079569 bytes to /content/Copy_of_vllm_serving_project_clean.html